tiny-asm: an assembler for riscv

jacob navia

Contents

List of Tables

List of Figures

1 The RISC-V assembler

1.1 Introduction

The tiny assembler is a "digest" of the GNU `gas` assembler. I have extracted from the 1.3Gb of `binutils-gdb` source code[1] two files: `asm.c` and `asm.h`.

There are two goals here:

1. To produce a small and fast assembler to be used as a compiler back-end. The elimination of features proceeds according to this goal: assemble machine generated output, without consideration for any human user, since all input to the assembler is supposed to be machine generated.

2. To produce a minimal set of sources that is *easy to read and understand* so that people can hack away without a lengthy learning curve. This documentation also, contributes to this objective.

In this version of the tiny-assembler there isn't:

- No input pre-processing. No include files, nor any fancy macro processing.

- No fancy error messages, messages will be emitted only in english. If you want other language error output you are welcome to do it yourself. The rationale behind this is obviously that a high level language user, programming in C++ or C, will be completely unable to understand the assembler messages even if they are translated into his/her native language.

- This assembler is geared to the riscv CPU. All support for any other machine has been dropped, specially support for machines that have ceased to exist for more than 20 years: the Motorola 68000 family, the Sparc, the Z80, etc. I think that even `gas` could drop support for those machines also.

- The code has been cleaned up from all cruft like this:

```
/* The magic number BSD_FILL_SIZE_CROCK_4 is from BSD 4.2 VAX
 * flavoured AS.  The following bizarre behaviour is to be
 * compatible with above.  I guess they tried to take up to 8
 * bytes from a 4-byte expression and they forgot to sign
 * extend.  */
#define BSD_FILL_SIZE_CROCK_4 (4)
```

So, we are still in 2023 keeping bug compatibility with an assembler for a machine that ceased production in 2000?

---

[1] I have just done a `du -b ./binutils-gdb` Probably is a bit less since I didn't do an extensive search for only .c and .h files.

- All the indirection through macros that are expanded into members of function tables that makes the code impossible to follow are eliminated. Now, if you see code like `foo(bar);` it is highly likely that you are calling function `foo` with argument `bar`...

- All libraries are eliminated. Tiny-asm doesn't use `BFD` nor `libiberty` nor `libopcodes`. The only library used is `zlib`.

- There are only two files: `asm.c` and `asm.h`. No other include files are there, as far as I remember, excepting system includes like `stdio.h` of course.

I have avoided to put much code samples here. There only two source files, and if you want to see the exact code sequences you are free to look them up, it is not very difficult. I see no interest in filling pages with code.

1.1.1 Requirements

I have concentrated in explaining how things work, and that includes talking about specifications and the standards used. You should have:

1. Source code: If you want the official sources of the GNU assembler you should download the binutils-gdb package. It is available in many places, for instance in github:

   https://github.com/bminor/binutils-gdb.git.

   You can download the sources of the `tiny-asm` from:
   https://github.com/jacob-navia/tiny-asm.

2. Assembler user documentation in "Using as".
   https://sourceware.org/binutils/docs/as/ This is the official documentation for the Gnu Assembler. Tiny-asm has kept most of it, and the algorithms, names of functions and variables are almost always the same. Knowing what the user specifications are will help you understand what the different assembler directives are doing.

3. The RISC-V Instruction Set Manual Volume I: Unprivileged ISA. There are a lot of versions of this document in the internet. Please try the most recent that you can find, of course. The official sources of the documentation are in
   https://riscv.org/wp-content/uploads/2019/12/riscv-spec-20191213.pdf, but there are apparently more recent ones. There is a depository in github at
   https://github.com/riscv/riscv-isa-manual, but they are in a strange format called "adoc" that is difficult to find a translator for, in non-windows systems.

4. DWARF debug information standard, the most recent being DWARF5 (2017) at
   https://dwarfstd.org/doc/DWARF5.pdf. This will enable you to better understand the debug information (cfi) directives of the assembler.

5. The ELF (Executable and Link Format) standard has an official page in the linux foundation at
   https://refspecs.linuxfoundation.org/elf/elf.pdf.

   ELF is the object format standard followed by the assembler. This will help you understand the `write_object_file` better.

6. You should obviously have a riscv machine. If you don't use a simulator (slow) buy a cheap board that can run linux. The chinese propose several machines, like
   https://pine64.com/product-category/star64/. This is the machine I am using, for around 110 US$. You can buy similar ones directly from the chinese, for instance
   https://www.waveshare.com/visionfive2.htm, or buy it from amazon.com, there are several boards available there. The Sifive company sells riscv boards also, but they

are not interested in retail sales. Demands for price and availability go into the bit bucket unless you are a huge company with orders of several hundred boards probably. But you can always try at https://www.sifive.com.

1.2 Building tiny-asm

The build process runs as follows:

1. Download the software from github

2. Build it:

```
$ gcc -o asm -g asm.c -lz
```

That is it. There is no Makefile but you can write one. I wrote this one:

```
star64:~/tiny-asm$ cat Makefile
asm: asm.o
gcc -o asm asm.o  -g -lz

asm.o: asm.c asm.h
gcc  -W -Wall -Wstrict-prototypes -Wmissing-prototypes\
-Wshadow -Wwrite-strings -g -c asm.c
clean:
rm -f asm.o asm
```

The Makefile for `gas` is 2268 lines... an impressing piece of software. However I think that 9 lines is much easier to understand. The user wants to use an assembler, maybe modify it, so there is no point in making him/her try to modify a 2 thousand line Makefile.

1.3 Overview

Like all assemblers, this assembler has a **parser**, where the text of the input file is converted into logical units that represent either instructions for the machine, or for the assembler itself, called *pseudo instructions*, and an **encoder**, where the instruction and its arguments are encoded into a 32 or 16 bit instruction and added to the current fragment. And then there is the object file generation, where the instructions and associated information are packed into the ELF format.
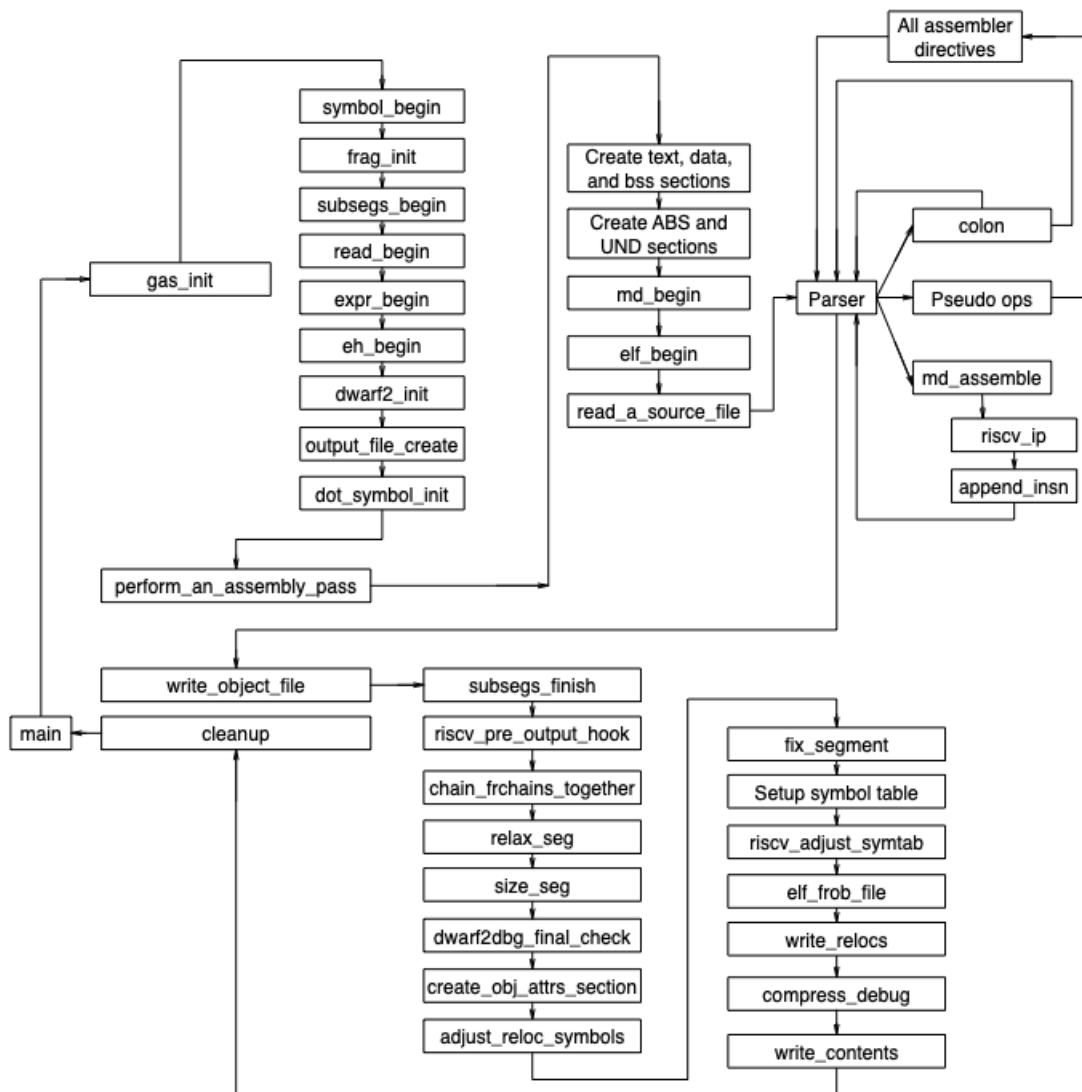
Figure 1.1: Overview of the assembler control flow

In figure 1.1 (page 12) we have these three main parts. Please keep in mind that this is a high level abstraction of the control flow. Obviously, if we would put each statement in the diagram we would have cram 40 000 lines into a diagram... too much.

We start with `main` that organizes all three parts [2]. It calls the initialization, `gas_init`, that initializes the symbols (`symbol_begin`,),the fragments initialization, the sub-segments, etc.

"Fragments" are understood in the assembler as pieces of code already assembled but that can grow, getting new instructions or other data. They are of variable length, and they

---

[2] Please be aware that in the diagram there is a direct link between, for instance, the function `dot_symbol_init` and `perform_an_assembly_pass`. This does NOT mean that the first calls the second directly. It means that the flow of the program returns to `gas_init` and then returns to the main function , and it is `main` function that calls `perform_an_assembly_pass`.

That would be quite complicated to draw, however. So, the diagram simplifies this.

will be strung together in a process called "relaxation" at the end of the assembly.

The initialization of the "sub-segments" means the text, data, and bss sections are created. Are "sub-segments" just plain object file sections? Not quite. There are "sections" like the "ABS" (absolute) section or the "UND" (undefined) sections that will never be written out in the object file.

There are other initializations that give us the opportunity of explaining some concepts that will be important later on. The `eh_begin` function, for instance, initializes the "exception handling" stuff. This is a complicated system that allows languages like C++ to walk the stack at run time, searching for a handler that will accept handling the exception that has just occurred.

This process involves an impressive machinery that contains a set of tables that associate addresses in the code to descriptions of the stack contents that allow a debugger or a runtime interpreter to see what functions have in terms of local variables and the space that each stack frame uses in the stack. And even if you are programming in C and you do not have any need for exceptions you will get them anyway since your C code could be called from a C++ program.

Other initializations concerns the start of the `dwarf2` debug information generation. Yes, the assembler can emit debug information for the program it is assembling. This way, the assembly programmer can follow the program line by line. `tiny-asm` has kept this even if it is highly unlikely that the compiler, that emits its own and much richer debug information, will need this.

The initialization of the "dot symbol" needs also some explaining. The current location when assembling a program is called "dot", i.e. a point. This symbol is always associated with the current address following a long assembler tradition that goes back to the start of the micro-computer age.

Eventually we come to the `perform_an_assembly_pass` function. This one continues the initialization process by creating the standard sections of the object file:

- The text section. This is a misnomer since there isn't anything textual inside. It contains the binary codes that will be interpreted by the integrated circuit. This is the most important output of the whole assembly process.

- The data section. This contains the tables, constants, structures and everything that the programmer has defined as static data that will be loaded at the start of the program by the program loader.

- The BSS section that contains nothing. It is just a reserved memory space that will be allocated by the program loader when it loads the program and contains always zeroes at the start.

- There are many other sections in an ELF format file. Let's stop here.

Then, we finish the setup process by calling `md_begin` and `elf_begin` functions.

The `md_begin` function reads all the static tables and builds hash tables from the for fast access. The opcodes are stored in hash tables, together with other data like the register names, the Control and Status Registers (CSRs) and what have you.

The `elf_begin` function builds symbols for each section in the object file. This allows to emit relocations or symbol addresses as an offset from the start of the section.

The setup phase behind us, we start the real work of the assembler: the well named `read_a_source_file`. This function does the parsing and the encoding of the instructions and directives.

In the diagram below, the functions aren't shown with their actual names but with their functional description. The GAS developers took (as you can see) a lot of effort to choose clear names that describe quite well what each function is doing. Still, I thought that here

we will use functional boxes instead of function names, since some of the functions described here do not exist as a separated subroutine but they are just pieces of `read_a_source_file`.

Figure 1.2: A more detailed view of the parser

We assume that the assembler input is a single file containing instructions, data, and assembler directives. In this version of the assembler, parsing is reduced to a bare minimum since we assume that we are assembling compiler output, and all the sophistication that is needed for an assembler adapted to human use is not needed for an assembler that is used to parse machine output.

We start with the function `read_a_source_file` that organizes the parsing and the instruction generation[3].

1. Setup. Here, we setup the input file name, in variable `physical_input_file` and we care about writing a file name record if we are emitting debug information.

---

[3]Actually, the initialization phase is executed before, but we will abstract that away for the time being
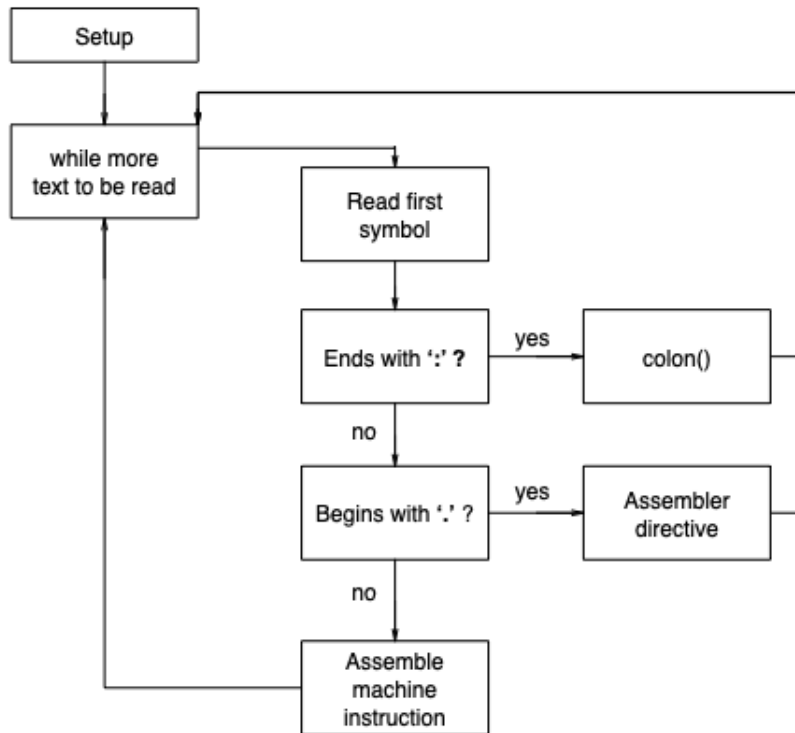
Figure 1.3: `read_a_source_file function`

2. We read a chunk of the input file. Currently, `BUFFER_SIZE` is set to 256 * 1024, and can be changed just by editing the corresponding line in asm.h

3. We start parsing lines. The first thing we read should be a symbol. If it ends with a colon, it is a label definition. We call the corresponding function `colon()` and continue parsing. If it is not finished by a colon, we see if the first letter is a point. If it is, it is an assembler directive. We call the corresponding function stored in the pseudo-ops structure (called `pseudo_typeS`) and we go fishing for the next line. If it is not a pseudo-operation, it *must* be a machine instruction. We call the `md_assemble` function.

The `md_assemble` function does basically following things:

1. Test if the instruction is valid using the current set of RISCV specifications. There are instructions that can be issued only with 64 or even 128 bits, or floating point instructions that depend on floating point being implemented in hardware, etc. RISCV machines can have a number of extensions implemented, since the basic ISA (Instruction Set Architecture) doesn't even have multiplication or division!

   Each "extension" has a letter that characterizes it. For instance, in the machine I am using we have in /proc/cpuinfo a line with:

   `isa : rv64imafdc`

   This means that the machine is a risc V 64 bits machine (rv64), with the integer (i), multiplication (m), a (Atomic), f (single precision floating point), d (double precision floating point) and c (Compressed instructions in 16 bits) extensions. The assembler

should test if any instruction is legal in the current subset, and reject those that do not comply.

Since we are an assembler for reading compiler output, we just assume the compiler doesn't emit wrong instructions and skip this test.

2. We call the `riscv_ip` function to encode the instruction. Basically it uses the `args` character string to know what arguments are expected. It verifies that those are correct, and inserts all necessary bits at the required positions. We will see later how these formats are defined.

3. If assembly succeeded the new instruction is added to the current fragment.

The `riscv_ip` function is basically a huge switch statement. The function will go through each one of the characters present in the `args` string of the opcode and add the necessary bits to the instruction.

In the tables below you will find the description of the different formats defined for each of the instructions in a riscv machine. This tables will help you understand how `riscv_ip` works.

1.4   General concepts and data structures

1.4.1   Binary File Descriptor (bfd)

This structure is at the heart of the BFD library. In the context of an assembler, there is only one of these beasts around, called `stdoutput`.

The main things stored here are:

- The file name.
- A table of function pointers that dispatches to the back end for doing most of the work.
- The input output stream
- A pointer to the back-end private data.

Of course there are a lot of other fields that you can study by reading the definition of this structure in asm.h.

It is important to underscore here that the table of function pointers has been completely eliminated in `tiny-asm`. There are no more indirection through the `xvec` field, since tiny-asm will only assemble riscv instructions. The front end and the back end have been merged into a monolithic whole. Still, this design is essential for understanding `gas`.

1.4.2   Symbols

There are several types of different structures that together represent a symbol. They will be described below, but in general they reflect the need by the bfd library to make a back-end independent structure that holds some general information, a high level abstraction of a symbol. Back-ends can differ in the object format, and in the cpu used,

so the information that is common to all those very different contexts is rather minimal.
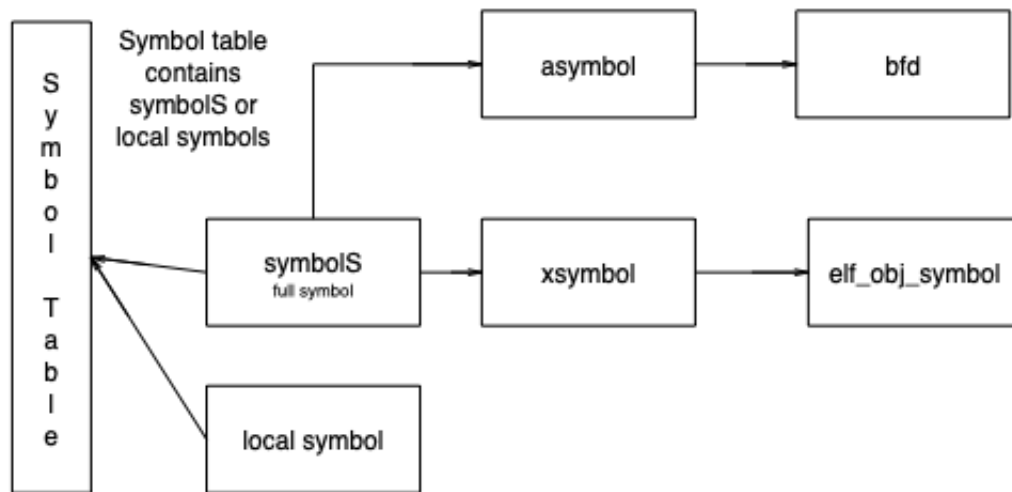
asymbol

This is the bfd-internal format, holding the following things:

- `the_bfd`. This points to the bfd that this symbol refers to. Since under tiny-asm there is only one bfd, called stdoutput, this is redundant. In other contexts, for instance in the linker where there are a lot of binary files opened for reading and one for writing, this makes much more sense.
- Name The name of the symbol.
- Value. Here there is either a pointer to some other symbol, or a numeric value.
- Flags. A long list of different flags. Some of them aren't used in tiny-asm, but their definition is still there since they are used in the linker.
- A pointer to the section.
- A pointer to special data used by the back end. It is a union of a generic pointer and an address.

To access the fields of an `asymbol` inline functions with rather lengthy names are provided. These functions look like this:

```
1  static inline asection * bfd_asymbol_section (const asymbol *sy)
2  {
3      return sy→section;
4  }
```

This function accesses the `section` field[4]

`asymbols` are global, and used for data that concerns a whole section. The constructor for these objects is `elf_make_empty_symbol`

symbolS

This structure is used for all kinds of symbols (labels, functions) that the assembler extracts from the source code. In the symbol table, we find pointers to this one and to the

---

[4]Is this really necessary? What are the real advantages of using `bfd_asymbol_section(foo)` instead of `foo→section`?

Yes, I know. It is called "information hiding". But the problem is that information hiding **hides** information, precisely, and if you are trying to understand what the code is doing, you do not know beforehand if that is a call to a lengthy function or just a field access.

local_symbol structures. The size of this should be the same or less than struct local_symbol, and fields that do not fit in that size go into an overflow structure called xsymbol for exension of symbol.

Listing 1.1: symbolS, xsymbol, elf_obj_sy

```
1  typedef struct symbol {
2      struct symbol_flags flags; /* Symbol flags. */
3      hashval_t hash;            /* Hash value calculated from name. */
4      const char *name;          /* The symbol name. */
5      fragS *frag;    /* Pointer to the frag of this symbol,if any. Otherwise NULL.*/
6      asymbol       *bsym;       /* BFD symbol */
7      struct xsymbol *x;         /* Extra symbol fields that won't fit. */
8  } symbolS;
9
10 /* Extra fields to make up a full symbol. */
11 struct xsymbol {
12     expressionS value;  /* Symbol value. Note that this is NOT a pointer */
13     /* Forwards and backwards chain pointers. */
14     struct symbol *next;
15     struct symbol *previous;
16     struct elf_obj_sy obj; /* Yet another symbol structure (YASS!) */
17 };
18
19 /* Additional information we keep for each symbol. */
20 struct elf_obj_sy {
21     unsigned int local : 1;   /* Whether the symbol has been marked as local. */
22     unsigned int rename : 1;  /* Whether the symbol has been marked for rename with
           @@@. */
23     unsigned int bad_version:1; /* Whether the symbol has a bad version name. */
24     /* Whether visibility of the symbol should be changed. */
25     ENUM_BITFIELD (elf_visibility) visibility : 2;
26     /* Keep track of .size expressions that involve yet unresolved differences */
27     expressionS *size;
28     /* The list of names specified by the .symver directive. */
29     struct elf_versioned_name_list *versioned_name;
30 };
```

The constructors for symbolS are:

- symbol_make. This constructor is simple, code below:

```
1      static symbolS *symbol_make(const char *name)
2      {
3          /* Let the machine description default it,e.g. for register names. */
4          symbolS *symbolP = md_undefined_symbol((char *)name);
5          if (!symbolP) symbolP = symbol_new(name,undefined_section,&
               zero_address_frag,0);
6          return (symbolP);
7      }
```

  So, if you read this you would think that first, as the commentary says, is calling to some function... Actually, for the riscv backend, we have a #define in asm.h:
  #define␣md_undefined_symbol(name)␣(0)
  symbol_make is just an alias for symbol_new.
- symbol_create. This function allocates space for the new symbol, sets some default fields, and then calls symbol_init that will finish the construction of the new symbol.
- symbol_new. This is a small function that calls symbol_create and then links the new symbol into the global list of symbols using the function symbol_append.

- `symbol_find_or_make(const char *name)`. This function searches for a symbol and if not found creates an undefined symbol, returning a pointer to it. When creating a symbol, it checks if it is a local symbol. Then either calls the constructor for a local or a true symbol.

In the symbol table, full fledged symbols or local symbols appear. The distinction between them is that for many symbols like labels, or similar, all the huge amount of information described above make no sense. A shorter and smaller structure is used, what makes considerable gains in memory space.

The `symbol_flags` structure

Flags that annotate an object are normally defined as `#defines` for instance

```
1 #define FLAG_FOO 1
2 #define FLAGG _BAR 2
3 ... etc
```

But there are better ways, for instance by defining a structure that can be accessed in a much clearer way.

Listing 1.2: The symbol_flags structure

```
struct symbol_flags {
    unsigned int  local_symbol:1; /* Whether the symbol is a local_symbol or not. */
    unsigned int  written:1;    /* Whether symbol has been written. */
    /* Whether symbol value has been completely resolved (used during final pass
     * over symbol table). */
    unsigned int  resolved:1;
    /* Whether the symbol value is currently being resolved (used to detect loops
     * in symbol dependencies). */
    unsigned int  resolving:1;
    /* Whether the symbol value is used in a reloc. This is used to ensure that
     * symbols used in relocs are written out,even if they are local and would
     * otherwise not be. */
    unsigned int  used_in_reloc:1;

    /* Whether the symbol is used as an operand or in an expression. */
    unsigned int  used:1;

    unsigned int  volatil:1;  /* Whether the symbol can be re-defined. */

    /* * Whether the symbol is a forward reference,and whether such has been
     * determined. */
    unsigned int  forward_ref:1;
    unsigned int  forward_resolved:1;

    /* This is set if the symbol is set with a .weakref directive. */
    unsigned int  weakrefr:1;

    /* This is set when the symbol is referenced as part of a .weakref
     * directive, but only if the symbol was not in the symbol table
     * before. It is cleared as soon as any direct reference to the symbol
     * is present. */
    unsigned int  weakrefd:1;

    /* Whether the symbol has been marked to be removed by a .symver
     * directive. */
    unsigned int  removed:1;
```

```
    /* Set when a warning about the symbol containing multibyte characters
     * is generated. */
    unsigned int  multibyte_warned:1;
};
```

Many of these fields are never used directly. The field `used_in_reloc`, for instance, is never accessed directly. The reason is clear when you look at the code for its access functions:

```
1  /* Mark a symbol as having been used in a reloc. */
2  static void symbol_mark_used_in_reloc(symbolS * s)
3  {
4      if (s→flags.local_symbol)
5          s = local_symbol_convert(s);
6      s→flags.used_in_reloc = 1;
7  }
8  /* Return if a symbol has been used in a relocation */
9  static int symbol_used_in_reloc_p(symbolS * s)/* Return whether a symbol has been
       used in a reloc. */
10 {
11     if (s→flags.local_symbol)
12         return 0; /* Local symbols can't be used in relocations */
13     return s→flags.used_in_reloc;
14 }
```

local_symbol

This structure can represent only symbols with offsets within the same `frag`.

Listing 1.3: local symbol

```
1  struct local_symbol {
2      struct symbol_flags flags; /* Flags: Only local_symbol and resolved relevant.*/
3      hashval_t  hash;           /* Hash value calculated from name. */
4      const char    *name;       /* The symbol name. */
5      fragS         *frag;       /* The symbol frag. */
6      asection      *section;    /* The symbol section. */
7      valueT     value;          /* The value of the symbol. */
8  };
```

Constructor for the `local_symbol` structure is the function `local_symbol_make`.

Symbol table

Listing 1.4: union symbol_entry_t

```
1  /* This structure makes up the entries of the symbol table */
2  typedef union symbol_entry {
3      struct local_symbol lsy;
4      struct symbol  sy;
5  } symbol_entry_t;
```

The symbol table is a hash table called `sy_hash`, created at initialization in the function `symbol_begin` called from `gas_init`.

Adding symbols into the symbol table is done with `symbol_table_insert`, the function `symbol_find` searches for a given symbol.

1.4.3 Fragments

As the assembler reads the source file, it stores the information it parses into `frags`. These structures contain two parts: a fixed part, i.e. the bytes that are known to be emitted, and a variable part, whose length can change. Memory is allocated with a worst case philosophy: the maximum size of the variable part is known, and will be used when reserving memory for each fragment.

When the source file has been fully parsed the relaxation process will go through all these fragments and build the linear structures needed in the object file.

Listing 1.5: struct frag

```
1  struct frag {
2      /* Object file address. This is zero until relaxation writes it. */
3      addressT fr_address;
4      /* When relaxing multiple times, remember the address the frag had in the last
5         relax pass. */
6      addressT last_fr_address;
7
8      /* (Fixed) number of octets we know we have. May be 0. */
9      valueT fr_fix;
10     /* May be used for (Variable) number of octets after above.
11     The generic frag handling code no longer makes any use of fr_var. */
12     offsetT fr_var;
13     /* For variable-length tail. */
14     offsetT fr_offset;
15     /* For variable-length tail. Points to the symbol to use according to fr_type */
16     symbolS *fr_symbol;
17     /* Points to opcode low addr byte, for relaxation. */
18     char *fr_opcode;
19
20     /* Chain forward; ascending address order. Rooted in frch_root. */
21     struct frag *fr_next;
22
23     /* Where the frag was created, or where it became a variant frag. */
24     const char *fr_file;
25     unsigned int fr_line;
26
27     /* A serial number for a sequence of frags having at most one alignment
28        or org frag, and that at the tail of the sequence. */
29     unsigned int region:16;
30
31     /* Flipped each relax pass so we can easily determine whether fr_address
32        has been adjusted. */
33     unsigned int relax_marker:1;
34
35     /* Used to ensure that all insns are emitted on proper address boundaries. */
36     unsigned int has_code:1;
37     unsigned int insn_addr:6;
38
39     /* This field indicates the interpretation of fr_offset, fr_symbol and the
40      * variable-length tail of the frag, as well as the treatment it gets in various
41      * phases of processing. It does not affect the initial fr_fix characters;
42      * they are always supposed to be output verbatim (fixups aside). */
43     relax_stateT fr_type;
44     /* This is used as a size for this frag and is updated during relaxation */
45     relax_substateT fr_subtype;
46
```

```
47      /* A simple structure (see below) that is used when creating a mapping symbol and
           when
48         checking if the mapping symbol is still useful. Mapiing symbols are create
           when
49         transitioning from a given mapping state to another */
50      struct riscv_frag_type tc_frag_data;
51
52      /* Data begins here. */
53      char fr_literal[1];
54 };
55
56 struct riscv_frag_type { symbolS *first_map_symbol, *last_map_symbol; };
```

1.4.4 Fixups

In many situations, the assembler can't finish a calculation because all data needed for it isn't available. For instance a symbol is yet unresolved, or the exact location for some instruction component is absent.

In those situations the assembler emits a `fixup`. This is nothing else than an instruction on how to patch the output later, when all the data is known.

Fixups are described in a structure called `fixS` that holds mainly following kinds of information:

- `next` The `fixS` structures are linked in a list.
- `fx_frag` The fragment where the fix should be applied.
- `fx_where` The position within that fragment where the fix should be applied.
- The quantity to be added or subtracted. If it is a symbol, a pointer to that symbol will be stored in the fields `fx_addsy` or `fx_subsy`. Otherwise, if it is just a number it will stored in the field `fx_offset`.
- `fx_size`. The size (in bytes) of the fixup, i.e. how many bytes should be written at the given location.
- There are many other fields that you can look up in the definition in asm.h. They are described in the comments surrounding their definition.

Constructors

Two functions build a fixup: `fix_new` and `fix_new_exp`. The second one is for a fixup referring to an expression, the first is for a symbol with an optional offset. They differ only in that `fix_new_exp` determines the symbol to add or subtract from the given expression. Both call `fix_new_internal` to do the actual fix.

Applying a fixup

A fixup is resolved by the function `md_apply_fix`. It uses the type of fixup to determine the sequence of actions to be performed: to fix the high 20 bits of a 32 bit address, or the lower 12, or add to a 64 bit address an addend, etc. The code consists (yes, you guessed it!) of a big `switch` statement with all the handled types of relocation existing for riscv machines, and it is not very difficult to follow.

1.4.5 Relocations

Sometimes a fixup can't be resolved. For instance this C code:

```
1 #include <stdio.h>
2 int main(void) {
```

```
3     printf("hello\n");
4 }
```

gcc translates this to:

```
1     .section    .rodata
2 .LC0:
3     .string "hello"
4     .text
5 main:
6     /* irrelevant stuff ellided */
7     lla a0,.LC0
8     call    puts@plt
9     /* further stuff ellided */
```

The address of the puts procedure can't be established by the assembler, nor the linker, only by the program loader that will know at load time the address of the shared library `libc6.so`. The assembler makes the same thing as when establishing a fixup. It makes a new fixup, this time for the linker, that will tell it where the address of the puts function needs to be stored.

This kind of fixup is called a *relocation.*

The linker can't resolve the address either, so it will make a relocation for the program loader, that will patch the code accordingly when the program starts[5].

Relocations, contrary to simple fixups have a standard format prescribed in the object file format, in our case ELF.

Listing 1.6: Elf relocation structure

```
1 typedef struct {
2     unsigned char r_offset[8]; /* Location at which to apply the action */
3     unsigned char r_info[8]; /* index and type of relocation */
4     unsigned char r_addend[8]; /* Constant addend used to compute value */
5 } Elf64_External_Rela;
```

This format doesn't exactly correspond to the internal one used by the assembler. The function `bfd_elf64_swap_reloca_out` converts from the bfd format to the ELF one.

1.4.6   Sections and subsections

Assembled data falls into four sections: opcodes, initialized data, uninitialized data and debug information. You may have separate groups of data in those sections that you want to end up near to each other in the object file, even though they are not contiguous in the assembler source.

Tiny-asm allows you to use subsections for this purpose. Within each section, there can be numbered subsections with values from 0 to 8191 [6].

Objects assembled into the same subsection go into the object file together with other objects in the same subsection. For example, a compiler might want to store constants in the text section, but might not want to have them interspersed with the program being assembled. In this case, the compiler could issue a `.text 0` before each section of code being output, and a `.text 1` before each group of constants being output.

Subsections are optional. If you do not use subsections, everything goes in subsection number zero.

Each subsection is zero-padded up to a multiple of four bytes.

---

[5]The process is obviously much more complicated. Here we leave all the details out, to take a high level view.

[6]This limit is mentioned in the GAS documentation. In the software, actually, there isn't a single test to enforce this limit, so you can write any number between 1 and MAX_INT.

Subsections appear in your object file in numeric order, lowest numbered to highest. The object file contains no representation of subsections; ld, objdump and other programs that manipulate object files see no trace of them. They just see all your text subsections as a single text section, and all your data subsections as a data section.

To specify which subsection you want subsequent statements assembled into, use a numeric argument to specify it, in a `.text <number>` or a `.data <number>` statement. If you just say `.text` then an implicit zero is assumed. Likewise `.data` means `.data 0`.

In the source code, sometimes subsections are called "subsegments".

1.5 Instruction formats and encoding

Yes, there are several parts in an assembler, but there is a fundamental part that makes the purpose of the whole program: **encoding instructions**. The essential part is here: transforming ASCII text representing instructions into a series of 16 or 32 bit sequences that encode each operation that the machine can do, including operations that are seldom, if ever, used.

To understand how the assembler works, it is important to keep in mind how the machine works, the names of its parts, and the intricacies of instruction encoding. Yes, yes, that looks awfully dry and uninteresting. But (to me) it is interesting, and if you do not like to *understand* how things work, please go to tik-tok and play some games...

There are several types of instruction encoding, named **R**, **I**, **S**, **B**, **U**, **J**.

- All are 32 bits, like the ARM.
- The first 7 bits are reserved for the opcode (bits 0 to 6).
- The same operand, for instance the source register 1 (sr1) is at the same position, bits 15 to 19.
- All instructions have at least one register operand.
- Since we have 32 registers, all register encoding take 5 bits.

The risc v introduces a more functional naming schema, where registers are assigned usage names, instead of the register numbers. Here is a correspondence table between them:

Table 1.1: RISCV symbolic register names

Register name	ABI name	Description	Register name	ABI name	Description
		Integer registers			
x0	zero	Hard-wired zero	x16	a6	Seventh argument
x1	ra	Return Address	x17	a7	Eighth argument
x2	sp	Stack pointer	x18	s2	Saved 2
x3	gp	Global pointer	x19	s3	Saved 3
x4	tp	Thread Pointer	x20	s4	Saved 4
x5	t0	Temporary/Alternate link register	x21	s5	Saved 5
x6	t1	Temporary	x22	s6	Saved 6
x7	t2	Temporary	x23	s7	Saved 7
x8	fp/s0	Frame pointer	x24	s8	Saved 8
x9	s1	Saved 1	x25	s9	Saved 9
x10	a0	First argument / Return value	x26	s10	Saved 10
x11	a1	Second Argument / Return value	x27	s11	Saved 11

Table 1.1: RISCV symbolic register names

x12-x15	a2-a5	Argument 3-5	x28-x31	t3-t6	Temporary registers
		Floating point registers			
f0-f7	ft0-ft7	Fp temps	f2-f7	fa2-fa7	function arguments
f8-f9	fs0-fs1	Fp saved registers	f18-f27	fs2-fs11	saved registers
f10-f11	fa0-fa1	Fp arguments/return value	f28-f31	ft8-ft11	Temporary registers

The difference between the ABI names and the actual register numbers is due to the fact that the ranges of registers are not contiguous. For instance the range of saved registers has two of them as x8 and x9, then the rest is x18 to x27.

1.6 The instruction formats

Each format is designed to be used by similar type of instructions.

- **R** Register to register ALU instructions.
- **I** Immediate and load.
- **S** Store and comparisons.
- **B** Branch.
- **U J** Jump and jump with link (call) instructions.

The RISC-V manual comments these formats like this

> The RISC-V ISA keeps the source (rs1 and rs2) and destination (rd) registers at the same position in all formats to simplify decoding. Except for the 5-bit immediates used in CSR instructions, immediates are always sign-extended, and are generally packed towards the leftmost available bits in the instruction and have been allocated to reduce hardware complexity. In particular, the sign bit for all immediates is always in bit 31 of the instruction to speed sign-extension circuitry. [7]

All instructions are 32 bits. This requirement, that riscv shares with ARM and other machines, is necessary to make possible parallel decoding of instructions. In the x86, for instance, each instructions is of variable length, what makes parallel decoding of instructions an incredibly difficult undertaking.

The `C` extension compresses 32 bit instructions into 16 bits, what makes for more compact code. They expand into 32 bits instructions.

Table 1.2: The different instruction formats

---

[7] RISC-V User level ISA V 2.2 §2.2. They add further down: Decoding register specifiers is usually on the critical paths in implementations, and so the instruction format was chosen to keep all register specifiers at the same position in all formats at the expense of having to move immediate bits across formats

Table 1.2: The different instruction formats

``"R" format:
struct rFormat {
    unsigned extra7:7;
    unsigned src2:5;
    unsigned src1:5;
    unsigned extra3:3;
    unsigned dst:5;
    unsigned opcode:7;
};```
	``"I" format
	struct iFormat {
	    unsigned imm12:12;
	    unsigned src1:5;
	    unsigned extra3:3;
	    unsigned dst:3;
	    unsigned opcode:7;
	};```
		``"U" format
		struct uFormat {
		    unsigned imm20:20;
		    unsigned dst:5;
		    unsigned opcode:7;
		};```

``"S" format
struct sFormat {
    unsigned imm12_2:7;
    unsigned src2:5;
    unsigned src1:5,
    unsigned extra3:3;
    unsigned imm12_1:5;
    unsigned opcode:7;
};```
	``"B" format
	struct bFormat {
	    unsigned imm12_sign:1;
	    unsigned imm12_10_5:6;
	    unsigned src2:5;
	    unsigned src1:5;
	    unsigned extra3:3;
	    unsigned imm12_1_4:4;
	    unsigned imm12_11:1;
	    unsigned opcode:7;
	};```
		``"J" format
		struct jFormat {
		    unsigned imm12_sign:1;
		    unsigned imm12_1_10:10;
		    unsigned imm12_11:1;
		    unsigned imm12_12_19:7;
		    unsigned dst:5;
		    unsigned opcode:7;
		};```

1.6.1   The "R" format

This format features 3 registers (destination, source 1 and source 2) and has two fields of 3 and seven bits available for use to customize the opcodes. In C we could describe that as: We use a 32 bit addition as an example of this format: `addw a0,a0,a1`. The addition using ABI

Figure 1.4: **R** Instruction layout

names is `addw a0,a0,a1` but using actual register numbers we have `addw x10,x10,x11`. For this instruction the 10 bits of `extra-3` and `extra-7` are empty.

We have then:

- Opcode: 0 1 1 1 0 1 1 → 0x3b (59 decimal).

- Destination register: 0 1 0 1 0 → 0xA (10 decimal). Register 10 is `a0`, that contains the first argument and is loaded with the result.

- Source 1: 0 1 0 1 0 → 0xA (10 decimal). Register 10 (a0) is the first source.

- Source 2: 0 1 0 1 1 → 0xB (11 decimal). Register 11 (a1) is the second source.

Software handling

We have an instruction with the `args` format of "Cs,Cw,Ct" that expects source and destination to be identical (`s` and `w`) followed by a target register in the expected range for compressed registers. All of that is true, and we succeed with a compressed 16 bit instruction.

Obviously this is not what we wanted. We wanted a 32 bit 'R' instruction. To be able to do that, we add the following instruction at the top of our assembler file

```
.option arch -c
```

I.e. we disable all compressed instructions.

We see here that the *order* in the layout of the opcode table is very important. The instructions that are **more** constrained should come first, and the general formats should come last. For instance the compressed instruction should come first, and non-compressed last, since the software stops at the first match.

1.6.2 The "I" format

This format changes the "R" format by merging `src2` with `extra-7` to give a 12 bit field where an immediate integer value caan be stored (up to $2^{12}-1 \rightarrow 4095$ values can be stored).

Figure 1.5: **I** Instruction layout

Software handling

The first instruction that the software tries has its `args` string: "Cs,Cw,Ct", we expect a source register in compressed format, i.e. register 8-15, followed by the *same* register. The second condition succeeds, and the software passes to the third argument: we expect a register, and we find the constant 1025. Nope, this instruction is not the one.

The next `addw` instruction to be tested has the string "Cs,Ct,Cw", a permutation of the above that fails also, for the same reasons.

More instructions are tried, with strings `d,Cu,Co` that fails, "d,s,t" that fails also since we have an immediate constant and not a register in the third position ('t' field). At last we arrive at an instruction with `args` field of `d,s,j`", i.e. a sign extended immediate ('j') in the third position. This time the software succeeds and we are done. Accessing the different fields is done with macros. Here is one example of a series of macros that extracts the immediate field of the immediate value in the instruction above

```
#define RV_X(x, s, n)  (((x) >> (s)) & ((1 << (n)) - 1))
```

This macro extracts `<n>` bits from `<x>`, beginning in bit position `<s>`. It has two parts:

1. The left side of the "and" operation that shifts the given number $<s>$ bits to the right to bring it to position zero, and

2. An expression that builds a mask of `<n>` 1 bits by shifting a 1 `<n>` positions to the right and subtracting one, what gives a power of two minus 1. A power of two minus 1 is a field full of ON bits in two's complement notation. For instance 1 << 3 → 8 (1000). You subtract 1 from that and you obtain 0111 (7), i.e. 3 bits "on", a mask to extract the lower 3 bits from a number.

```
#define RV_IMM_SIGN(x) (-(((x) >> 31) & 1))
```

This macro returns either -1 or 0, depending if the sign of the 32 bit number is negative or positive. Since -1 is 32 bits of "1" bits, it can be used to sign extend a number.

The two macros above are used in these new ones:

```
#define EXTRACT_ITYPE_IMM(x) (RV_X(x,20,12)|(RV_IMM_SIGN(x) << 12))
#define ENCODE_ITYPE_IMM(x) (RV_X(x, 0, 12) << 20)
```

The first macro extracts 12 bits from the given number ($<x>$) and sign-extends its sign. The second extracts the lower 12 bits of the value, and puts them at position 20-31 [8]

1.6.3   The "U" format

A variant of the **I** format featuring more space for immediate constants is the **U** format, that can hold immediate constants with 20 bits.

Figure 1.6:   **U** Instruction layout

The `lui`[9]. instruction loads an unsigned 20 bits immediate stored in the bits 12 to 31 of the instruction into the upper 20 bits of the destination and sets the lower 12 bits to zero. In C language notation we have: `dst = (imm20 << 12);`. The authors justify these choices with:

> In practice, most immediates are either small or require all XLEN bits. We chose an asymmetric immediate split (12 bits in regular instructions plus a special load upper immediate instruction with 20 bits) to increase the opcode space available for regular instructions.Immediates are sign-extended because we did not observe a benefit to using zero-extension for some immediates as in the MIPS ISA and wanted to keep the ISA as simple as possible.[10]

---

[8]It is a pity that machines implementing the boolean extension aren't widely available yet. I miss the ARM boolean instructions that will reduce many of those macros to a couple of instructions.

[9]`lui` stands for **l**oad **u**pper **i**mmediate

[10]Riscv ISA Architecture §2.2

Software handling

Looking up the `args` description for this instruction, we find the character string `"d,Cu"`. This means we should expect a register name, followed by a comma, and an immediate value to be able to use a `C` (compressed) instruction. But that doesn't work, our constant is beyond bounds of the compressed immediate.

The software continues its search for the correct instruction and we come to the next instruction in the list that has the `args` string `"d,u"`, without any compression requirements. This time a match is found, and necessary bits are inserted as shown in figure 1.6 page 28.

Obviously, loading an immediate constant that will be shifted by12 bits is seldom used. This is thought for loading the upper 20 bits of an *address*, then adding the lower 12 bits with another instruction. This constant was choosen in this example so that it has a 1 bit at the end of 10 bits, and 1 at the start to be visible in the drawing.

To extract the J type immediate we use the following macro:

```
1    #define EXTRACT_UTYPE_IMM(x) ((RV_X(x, 12, 20) << 12) | (RV_IMM_SIGN(x) << 32))
```

1.6.4 The "S" format

In this format, the `dst` field disappears and its bits are used to hold the lower 4 bits of an immediate value. An instruction that uses this format is the `sd` (store double word) instruction.

Figure 1.7: **S** Instruction layout

We use the instruction `sd ra,64(sp)` as example. This instruction means: Store the contents of the return address register (`ra`) at the memory address obtained by adding 64 to the contents of the sp register. We have here an address that is obtained by adding the contents of a register and a *displacement* that must fit into 12 bit. As you can see here, this is a much easier format than the ARM jungle of different types of offsets where you never really know which one to use. The Risc-V manual specifies that all offsets are signed.[11]

We have then for this instruction:

- `src1` is 0 0 0 1 0, or register 2.

- `src2` is 0 0 0 0 1, or register 1.

- The immediate is the concatenation of `imm[4:0]` and `imm[11:5]` i.e; 0 0 0 0 0 0 1 0 0 0 0 0 or 64.

---

[11]They say:

> Except for the 5-bit immediates used in CSR instructions, immediates are always sign-extended, and are generally packed towards the leftmost available bits in the instruction and have been allocated to reduce hardware complexity. In particular, the sign bit for all immediates is always in bit 31 of the instruction to speed sign-extension circuitry.

For extracting the immediate from the instruction we use the macro

```
#define EXTRACT_STYPE_IMM(x) \
(RV_X(x, 7, 5) | (RV_X(x, 25, 7) << 5) | (RV_IMM_SIGN(x) << 12))
```

This macro extracts five bits beginning at position seven, then 7 bits from position 25 upwards, shifted by 5 left, so that they come right after the first five. The whole is sign extended in the same way as explained in section 1.6.2 page 27.

1.6.5   The "B" format

Figure 1.8:   **B** Instruction layout

In this format, we have a 13 bit immediate for branches. The immediate represents the amount that will be added to the program counter to reach the specified location, in multiples of 2. Since the lowest bit of the immediate will be always zero, it has been replaced by bit 11 (the twelfth bit) adding one bit to the quantity being written. The range of the branch is $\pm$ 4K.

The different conditional branches are specified in the `extra-3` group, with

Table 1.3: Encoding of conditional branches

extra-3	Instruction	Description
0 0 0	beq	branch if equal
0 0 1	bne	branch if different
1 0 0	blt	branch if less than
1 0 1	bge	branch if greater/equal
1 1 0	bltu	branch if less than unsigned
1 1 1	bgeu	branch if greater equal unsigned

All these instructions share the same opcode: 99. The `extra-3` field is used to extend the opcode for different instructions.

The macro to access the immediate value is way more complicated due to the bit scrambling...

```
#define EXTRACT_BTYPE_IMM(x) ((RV_X(x, 8, 4) << 1) | \
(RV_X(x, 25, 6) << 5) | (RV_X(x, 7, 1) << 11) | (RV_IMM_SIGN(x) << 12))
```

1.6.6   The "J" format

The only difference between the U and J formats is that the 20-bit immediate is shifted left by 12 bits to form U immediates and by 1 bit to form J immediates. In the "J" format, the immediate represents an offset in pairs of 16 bit instructions from the current PC.

Why this scrambled layout? Citing the Risc-v manual:

Figure 1.9: **J** Instruction layout

Although more complex implementations might have separate adders for branch and jump calculations and so would not benefit from keeping the location of immediate bits constant across types of instruction, we wanted to reduce the hardware cost of the simplest implementations. By rotating bits in the instruction encoding of B and J immediates instead of using dynamic hardware muxes to multiply the immediate by 2, we reduce instruction signal fanout and immediate mux costs by around a factor of 2.

The scrambled immediate encoding will add negligible time to static or ahead-of-time compilation. For dynamic generation of instructions, there is some small additional overhead, but the most common short forward branches have straight-forward immediate encodings.

The macro to extract this monster from its hiding place looks like this

```
#define EXTRACT_JTYPE_IMM(x) ((RV_X(x, 21, 10) << 1)|(RV_X(x, 20, 1) << 11) | \
(RV_X(x, 12, 8) << 12) | (RV_IMM_SIGN(x) << 20))
#define ENCODE_JTYPE_IMM(x) ((RV_X(x, 1, 10) << 21)|(RV_X(x, 11, 1) << 20) | \
(RV_X(x, 12, 8) << 12) | (RV_X(x, 20, 1) << 31))
```

1.7 The compressed instructions

The Risc-v instructions are normally 32 bits in length. The "C" extension (C for **C**ompressed) encodes certain instructions in 16 bits, what leads to big savings in code size. These instructions aren't enabled by default in the assembler. You can enable them (if your machine actually supports them) with the instruction: `.option arch, +c`. Enabling them or not is not that important, since the linker will replace longer with shorter instruction whenever possible. For instance the jumps can't be really calculated until all the instructions are compressed, what only the linker can know.

The compressed instructions are enabled when one of these conditions is true:

- The compressed 16 bit instructions have the lowest 2 bits of the opcode set to either 00, 01, or 10.

- 32 bits instructions have their lowest two bits set to 11. The following 3 bits should have any value different from 111.

- The 48 bit instructions have their lowest 6 bits set to 011111. (5 bits set)

- 64 bit instructions have the 7 lower bits set to 0111111. (6 bits set)

The criteria for making a compressed instruction are as follows:

- The immediate or the address offset is small.

- One of the registers used is the `zero` register (x0), the return address register or link register `ra` (x1), or the stack pointer `sp` (x2).

- The destination and first source register are the same.

- The registers used belong to the 8 most popular ones, described with 3 bits in the table below[12].

Table 1.4: Compressed register numbers

number	000	001	010	011	100	101	110	111
int reg. number	x8	x9	x10	x11	x12	x13	x14	x15
ABI name	s0	s1	a0	a1	a2	a3	a4	a5
FP reg number	f8	f9	f10	f11	f12	f13	f14	f15
FP ABI name	fs0	fs1	fa0	fa1	fa2	fa3	fa4	fa5

There are nine different compressed instruction layouts.

In the table below the registers that use the 3 bit number are marked with a '.

Table 1.5: Compressed formats

Meaning	Code	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Register	CR	Extra-4				dst/src1					src2					op	
Immediate	CI	Extra-3			S	rd/rs1					immediate					op	
Store local	CSS	Extra-3			imm						rs2					op	
Wide imm	CIW	Extra-3			imm								rd'			op	
Load	CL	Extra-3			imm		rs1'			imm		rd'				op	
Store	CS	Extra-3			imm		rs1'			imm		rs2'				op	
Arithmetic	CA	Extra-6						rd'/rs1'			Extra-2		rs2'			op	
Branch	CB	Extra-3			offset		rs1'			offset						op	
Jump	CJ	Extra-3			jump target											op	

1.7.1 The compressed register (CR) format

Figure 1.10: Compressed **CR** Instruction layout

This format acccepts instructions where the destination and the first source register are the same. It has four fields, here from right to left, i.e. from bit 0 to 15:

1. OP: Bits 0-1. Value: 2.

2. Src2: Bits 2-6. The second source register. Note that it is specified in 5 bits, like dst/src1, so any register of the set of 32 is possible, except the zero register. In this case it is 13, i.e. register a3 (x13).

---

[12] Actually those numbers are just the normal register number modulo 8.

3. dst / src1: Bits 7-11. The source 1 and the destination register are the same. Also specified in 5 bits, in this case it is 12: the a2 (x12) register.

4. Extra-4: Bits 12-15. Value: 9. Complements the opcode. This field can have two values that correspond to `mv` (move) or, in the example, `add`.

The software side

The argument description for `addw,a2,a2,a3` is the character string `Cs,Cw,Ct`. The first argument is a compressed format source register (Cs), followed by a compressed format register that should be equal to the preceding one (Cw), followed by a compressed format second source register, (Ct).

The code for the 's' case in `riscv_ip` is as follows:

```
1   case 's':  /* RS1 x8-x15. */
2       if (!reg_lookup(&asarg,RCLASS_GPR,&regno)
3           || !(regno ≥ 8 && regno ≤ 15))
4       break;
5       INSERT_OPERAND(CRS1S,*ip,regno % 8);
6       continue;
```

It is a typical sample of the code in the encoder (`riscv_ip`). We search for a register name with `reg_lookup` and we ensure that is between 8 and 15. If that is not the case, the matching process for this instruction candidate fails, and we look for the next one (`break`).

If it is, we insert the operand in the right position and continue with this candidate.

Note that the identifier `CRS1S` doesn't appear in ANY macro, variable or enumeration in the whole program.

It is a literal name argument! When we look at the definition of `INSERT_OPERAND` we find:

```
1   #define INSERT_OPERAND(FIELD,INSN,VALUE) \
2       INSERT_BITS ((INSN).insn_opcode,VALUE,OP_MASK_##FIELD,OP_SH_##FIELD)
```

The `##` operand before the `FIELD` macro argument makes the preprocessor convert it to `OP_MASK_CRS1S` what is defined with `#define OP_MASK_CRS1S 0x7` in asm.h.

The first level expansion converts this to:

```
1   #define INSERT_OPERAND(FIELD,INSN,VALUE) \
2       INSERT_BITS ((INSN).insn_opcode,VALUE,OP_MASK_CRS1S,OP_SH_CRS1S)
```

The `INSERT_BITS` macro is defined as follows:

```
1   #define INSERT_BITS(STRUCT, VALUE, MASK, SHIFT) \
2       (STRUCT) = (((STRUCT) & ~((insn_t)(MASK) << (SHIFT))) \
3           | ((insn_t)((VALUE) & (MASK)) << (SHIFT)))
```

This macro has two parts, separated by an | (or) sign:

```
1   ((STRUCT) & ~((insn_t)(MASK) << (SHIFT))) and
2   ((insn_t)((VALUE) & (MASK)) << (SHIFT))
```

In the first one we set to zero all bits in the field that will be written. The second one introduces the bits into the right position. The *or* operation joins those parts into a single value.

The encoder works like an interpreter for a "language" of single letters that represent pieces of instruction fields. They indicate what to expect at the given position. Its actions can be only be "break" (discard the current candidate) or insert the correct bits and "continue" with it.

1.7.2   The compressed immediate (CI) format

These instructions perform operations between a register and a small immediate encoded in
only 6 bits. The register can't be the zero register, and the immediate can't be zero. There

Figure 1.11:   Compressed immediate **CI** Instruction layout

are four instructions that use the compressed immediate format. They differ in the `extra-3`
field. From least significant bit to the most significant one we have:

1. OP: Bits 0-1, always with value 1 for the CI format.

2. The immediate field, in bits 2 to 6 that encodes immediate bits 0 to 4. In the example
   above this is 27, 1 1 0 1 1 in binary.

3. The destination and the source register number over 5 bits. In the example we have
   14 since the register a4 has the number 14.

4. The sign of the immediate value in a single bit (index 12th).

5. The `Extra-3` field, that allows for 3 instructions to be distinguished: `addi`, `addiw`,
   and `addi16sp`. The last one adds a number of 16 bits quantities to the stack and is
   used to adjust the stack at the prologue or at the epilogue of a function. Since the
   stack must be aligned to a multiple of 16, there is no need to keep the lower 4 bits.
   This makes for adjustments of -512 to 496 bytes.

To access the immediate value we use

```
1    #define EXTRACT_CITYPE_IMM(x) (RV_X(x, 2, 5) | (-RV_X(x, 12, 1) << 5))
2    #define ENCODE_CITYPE_IMM(x) ((RV_X(x, 0, 5) << 2) | (RV_X(x, 5, 1) << 12))
```

The first macro uses the same technique for sign extending that our `RV_IMM_SIGN` uses (see
1.6.2 page 28). We just need another expression since the other was fixed for 32 bits.

1.7.3   The stack relative store (CSS) format

Five instructions use the CSS format:

1. `c.swsp` or store word to an offset from sp, the stack pointer.

2. `c.sdsp` or store double word (64 bits) to an offset from sp.

3. `c.fswsp` or store single precision (32 bits) to an offset from sp.

4. `c.fsdsp` or store double precision (64 bits) to an sp offset.

In our example instruction we have an `op` field of 2, an `src` field of 16 (10000) and the
cryptic "011000" sequence that is translated into 00110 (6 decimal) since the bits are scram-
bled: they are stored as bits 5 4 3 2 7 6 The macros to access the immediate displacement
here are:

Figure 1.12: Store to stack offset (**CSS**) instructions layout


```
1    #define EXTRACT_CSSTYPE_IMM(x) (RV_X(x, 7, 6) << 0)
2    #define ENCODE_CSSTYPE_IMM(x) (RV_X(x, 0, 6) << 7)
```

The encoding of instruction `c.swsp` needs only one source register: the source of the 32 bit data to store in memory. Any register will do since we have a register number in 5 bits. The value of the immediate displacement will be added to the stack pointer scaled by 4 to form the effective address. In the example above the 6 binary is scaled to 24. [13]

The argument description string is `"CV,CM(Cc)"`: We need a register name (`CV`), followed by a small constant (`CM`) that is a displacement (the parentheses) of the stack pointer (`Cc`). The constant value will be zero extended, since obviously negative offsets for the stack aren't very useful!

The reach of this instruction is $2^7 - 1$ values since we have 7 bits. Scaled by 4, i.e. 127 * 4 → 508.

And... "one more thing" as Steve Jobs liked to say, there is a problem with zero offsets from the stack pointer. Normally a zero offset is omitted, i.e. you do NOT write `sw a6,0(sp)`, you just write `sw a6,(sp)`. The handling of the `CM` directive tests for this with the function `riscv_handle_implicit_zero_offset`.

1.7.4 The wide immediate (CIW) format

This format is used to encode a constant in bits 5 to 12. It is used in the `addi4spn` instruction. The constant encoded in those 8 bits is scaled by 4, i.e. the two lower bits are implicit zeroes. The scaled value will be added to the stack pointer and written to the register whose index is stored in the 3 bits `rd'`. This instruction builds then pointers to values stored in the local stack frame.

Figure 1.13: Store to stack offset (**CIW**) instructions layout

1. The OP field is zero.

---

[13]By an unfortunate coincidence the scrambled bits of the constant are 011000, what is 24 in binary. Beware, nothing in this business is simple, and a 24 can be scrambled to 6, then scaled to 24 back again.

2. The destination (`rd'`) is 7, the register number in 3 bits of the a5 register

3. Now, this is more complicated to explain. The poor immediate bits are *scrambled*, i.e. they are **not** in the natural order but in the order: 5, 4, 9, 8, 7, 6, 2, 3. The bits 1 and 0 are implicitly zero. The quantity (128) has a single bit on at the position 7, what in our scrambled layout corresponds to bit 8. [14]. The Risc-V ISA manual justifies this saying:

> The immediate fields are scrambled in the instruction formats instead of in sequential order so that as many bits as possible are in the same position in every instruction, thereby simplifying implementations.[15]

The "simpliying" above refers to hardware simplification.

4. The Extra-3 field is zero.

The macros used to access the immediate are:

```
1    #define EXTRACT_CIWTYPE_ADDI4SPN_IMM(x) ((RV_X(x, 6, 1) << 2) |\
2    (RV_X(x, 5, 1) << 3) | (RV_X(x, 11, 2) << 4) | (RV_X(x, 7, 4) << 6))
3    #define ENCODE_CIWTYPE_ADDI4SPN_IMM(x) ((RV_X(x, 2, 1) << 6) |\
4    (RV_X(x, 3, 1) << 5) | (RV_X(x, 4, 2) << 11) | (RV_X(x, 6, 4) << 7))
```

The argument description string for this instruction is "Ct,Cc,CK"

1.7.5   The compressed load (CL) format

Figure 1.14:   Compressed load **CL** Instruction layout

1. The OP field is zero.

2. The destination register is 5 (a3).[16]

3. This field corresponds to an offset from a register. The constant should be aligned by a multiple of 4, since we are loading 4 bytes. The two lower bits then should be zero and they are implicit, i.e. they are absent from the encoding. The value is split between two bits at positions 5 and 6, and the rest in positions 10, 11, and 12. The two bits in positions 5 and 6 are scrambled, and bit 6 corresponds to bit 2 of the immediate and bit 5 is bit 6 of the immediate value, they are not consecutive.

4. The `rs1'` field contains 1 1 1, what corresponds to x15 (a5).

5. We have in bits 10, 11, and 12 the bits 3, 4, and 5 of the immediate value.

6. The `extra-3` field contains constant 2.

---

[14] The number 128 is 1000 0000 in binary. Bit 7 is one. In the scrambled order we have bit 7 in the fourth position of the immediate field, counting from left to right, as shown in the figure 1.13
[15] Risc-V Unprivileged ISA V20191213 §16.2
[16] These values are in table 1.4

1.8 The opcode table

The full table of opcodes (called `riscv_opcodes`) consists of entries with the following structure:

```
struct riscv_opcode {
const char *name;
```

The name of the instruction in lower case. This is also the used as the key to the hash table. Several instructions can share the same name, and they are recognized by their different arguments.

```
unsigned xlen_requirement;
```

The word bit length (32, 64, or 128) that is required to use this instruction. A zero here means no requirement.

```
enum riscv_insn_class insn_class;
```

The instruction class to which it belongs. For instance the instructions belonging to the basic integer operations are `INSN_CLASS_I` one of the member of the `enum riscv_insn_class`. This was used to decide whether or not this instruction is legal in the current machine architecture context, but this test has been dropped since we assume that the compiler will not generate instructions that are illegal for the target machine.

```
const char *args;
```

A string describing the arguments for this instruction. This string will be interpreted by the `riscv_ip` function in a rather big set of nested `switch` statements.

```
insn_t match;
insn_t mask;
```

The basic opcode for the instruction. When assembling, this opcode is modified by the arguments to produce the actual instruction that is used. If `pinfo` is `INSN_MACRO`, then this is 0. Otherwise the `mask` field is a bit mask used to isolate the relevant portions of the opcode when disassembling. If `pinfo` is `INSN_MACRO` then this field contains the macro identifier, encoded as a member of an anonymous enumeration and casted to an integer.

```
int (*match_func) (const struct riscv_opcode *op, insn_t word);
```

A function to determine if a word corresponds to this instruction. Usually, this computes `((word & mask) == match)`.

```
unsigned long pinfo;
```

Additional information about the instruction. They are:

Table 1.6: Opcode flags

Symbol	Description
INSN_ALIAS	Just an alias, for example "mv" for "addi dest,src,zero
INSN_BRANCH	Unconditional branch
INSN_CONDBRANCH	Conditional branch
INSN_JSR	Jump to a subroutine
INSN_DREF	Data reference

Table 1.6: Opcode flags

`INSN_V_EEW64`	Instruction allowed only when the machine is a 64 bit machine or more
`INSN_XX_BYTE`	5 different data size specifiers, for XX=1, 2, 4, 8, or 16 bytes

```
};
```

The field `args` above needs more explanation. It is a one (or more) letters that represent the type of argument that can be expected in an instruction. This can be a register, a constant within a certain range, or other things. During assembly, the assembler reads and interprets this character string to weed out wrong choices or emit warnings, and to verify that all constrains are met.

The table below should document all the letters used by the `riscv_ip` function. They are listed in the order they appear there; only for the first level. If a letter has a continuation (for instance for the compressed instructions), the secondary switch statement is explained in another table[17].

Table 1.7: Opcode arguments letters

Char	Description
\0	End of the argument string. Here are done the final checks, for instance that this instruction corresponds to the bit length of the machine (64 bit instructions can't be done in a 32 bit machine). It checks also if the end of the argument string coincides with the end of the actual arguments present. If everything goes well it sets the errors to zero and branchs to the end of the `riscv_ip` function.
C	Compressed format instructions. This leads to a nested switch statement, since all the compressed argument descriptions begin with a C letter. This switch is described in table 1.9 page 40.
V	Vector instructions. This leads to a nested switch statement too.
,	Synchronization. Arguments are separated by commas. The software tests this and ignores the separators.
()[]	Displacement or index. Same behavior as for commas.
<	Shift amount for shifts less than 32.
>	Shift amount for 0 to word length - 1.
Z	CSRRxI Immediate. **C**ontrol and **S**tatus **R**egisters are specified in a different instruction format. For this to work, you have to have access to a CPU with the 'z' extension.
E	Control register number. This is used only in privileged instructions.
m	Rounding mode. This argument expects a character string representing the rounding mode. It can be one of `"rne"`,`"rtz"`, `"rdn"`,`"rup"`, `"rmm"`, 0,0,`"dyn"`. See table 1.8 page 39.
PQ	Fence predecessor or successor
d	Destination register.
s	First source register. Also called src1 in the documentation.
t	Second source register. The '`t`' is for `target`. It is also called src2 in the documentation.
r	RS3
D	Floating point destination register
S	Floating point source 1.

---

[17]Nested tables are as difficult to read as nested `switch` statements.

Table 1.7: Opcode arguments letters

T	Floating point source 2
U	Floating point source 1 and 2
R	Floating point RS3.
F	Expects a bit field, that is defined by the following character. Used in the `.inst` directive only.
I	`M_LI` macro. Immediate value.
A	Requests a symbol
B	Requests a symbol or a constant.
j	Sign extended immediate.
q	Expects a register store displacement.
o	Expects a load displacement.
l	Used for thread local storage.
p	PC relative offset
0	Expects a zero displacement. For instance: `lr.w a5,0(sp)`.
u	Expects a 20 bit immediate
a	20 bit relative offset.
c	Call using the global object table
O	Opcode field
y	bs immediate for branch offsets.
Y	rnum immediate
z	Expects a zero
W	Various operands
X	Integer immediate
Xu	eXtract unsigned n bits starting at position m. These arguments look like this: `Xu2@25`, meaning e**X**tract 2 bits starting at bit 25.

Below is the set of rounding modes for the `m` parameter. It has been taken from the Sifive site[18]. Edited in May 27th 2020.

Table 1.8: Accepted rounding modes for the 'm' parameter

Binary Value	Mnemonic	Meaning
000	rne	Round to Nearest, ties to Even
001	rtz	Round towards Zero
010	rdn	Round Down (towards $-\infty$)
011	rup	Round Up (towards $+\infty$)
100	rmm	Round to Nearest, ties to Max Magnitude
101		Invalid. Reserved for future use.
110		Invalid. Reserved for future use.
111	dyn	In instruction's rm field, selects dynamic rounding mode; In Rounding Mode register, Invalid.

These rounding modes are recognized in the assembler using the `riscv_rm` table, a simple table of 8 character strings.

The C (compressed) instructions are differentiated by the following letters:

---

[18]https://observablehq.com/@nschwass/riscv-f-extension-single-precision-floating-point-instruction
The URL seems truncated but it is not...

Table 1.9: Compressed instruction types

Char	Description	Char	Description
s	Source register 1 (x8-x15)	w	Source 1 and destination when they are the same
t	Source 2 with x8-x15	x	Source 2 and destination are the same. x8-x15 only.
U	Source 1 and destination the same.	v	Source 2
c	Source 1 constrained to be sp	z	Source 2 should be the zero register
>	Shift amount between 0 and word length - 1	5	Five bit field
6	Six bit numeric field	8	Eight bit field
j	Non-zero immediate	k	Immediate (possibly zero)
l	Load immediate (64 bits)	m	Load immediate
n	Immediate offset from SP	o	`C.addiw`, `c.li`, and `c.andi` allow zero immediate. C.addi allows zero immediate as hint. Otherwise this is same as 'j'.
K	scaled by 4 stack addend	L	Stack offset scaled by 16
M	Scaled by 4 stack displacement(32 bits store)	N	Data reference with offset from stack scaled by 4(64 bits store)
u	Immediate for jumps	v	Immediate for jumps
S	Floating point source 1 x8-x15	D	Floating point source 2 x8-x15
T	Floating point source 2	F	Field of 6, 4, 3, or 2 bits
V	Second source integer register rs2		

This is an example for an instruction entry in the opcodes table:

```
{"addi",0,INSN_CLASS_C,"Ct,Cc,CK",MATCH_C_ADDI4SPN,MASK_C_ADDI4SPN,\
match_c_addi4spn,INSN_ALIAS},
```

After parsing the name of the instruction, the `riscv_ip` function examines entries in the opcode table starting with the first one that has this name. It copies this entry into temporary storage because it will modify it later (using the `create_insn` function).

Then, it uses the letter in the `args` character string to check if there is a match. If there is, it stores immediately the bits into the instruction copy. But, as mentioned above, if there isn't any match, all the work is discarded and `riscv_ip` starts over using a saved pointer to the start of the arguments.

This way it ensures that eventually, the good instruction will be discovered, if at all. It is a slow process, since in many cases 4 other 5 instructions will be parsed and discarded until the correct one is found. Since the order of the opcodes is crucial the most used instructions can be the last ones to be found, what compounds the problem.

Several solutions can be imagined to speed up things, but the question arises if the speed of the assembler encoding is really the limiting factor for the compilation process. In a very cheap riscv machine assembling a 3.6Mb file takes 1.7 seconds, including the time for i/o from disk.

1.9 Writing the object file

After we have encoded all instructions and setup all the static data, processed all the assembler directives, we arrive at the end of the file, and we start preparing for writing the result of our efforts: the object file.

This file is written according to the ELF (**E**xecutable and **L**ink **F**ormat.)[19] standard. This file format is extensively described in a lot of documentation floating in the internet, so it is not necessary to repeat all that here.

Before we start writing out things we must finish the assembling process.

- We have a long list of "fragments", each holding a piece of the final section... we have to stitch all that together.

- We have some symbols that still haven't got a specific location. We should resolve them.

- We have to prepare to write the file header and the section headers.

- We have symbols in an internal format. We have to prepare to write them out in the ELF symbol format.

- References to symbols (fixups) must be resolved as far as it is possible. Of course some symbols are just externals, and can't be resolved anyway.

1.9.1 Write the object file

The `write_object_file` function is a very long one (more than 250 lines). Here is a detailed account of it:

- `subsegs_finish` This function does mainly two things:

  1. Correctly align the section.
  2. Finish the last fragment, so that there isn't any half done fragment.

- `riscv_pre_output_hook` This function finishes optimizations of the `eh_frame` output. Basically, if a subtraction from two symbols is performed, it is feasible to substitute the subtraction by a constant when the two symbols are in the same fragment. Sometimes, however, it is impossible to know if that is the case. In that case the optimization is postponed to the end of the assembly. This is done here.

- The assembler creates some sections to store its own data. They need to be discarded now, since they aren't needed any more. Once we do that, the sections need to be renumbered since we have thrown away some.

- `chain_frchains_together` This function manipulates the next and previous pointer of the fragment chains to make a single list. Now, since we have chained everything in a single list, any new relocations must be done not relative to a fragment, but relative to the start of the big list. We record that we have done the fragment reorganization in the variable `frags_chained`. This global variable is used in the function `fix_new_internal` when making a new fixup:

      fixS **seg_fix_rootP = (frags_chained ? &seg_info(now_seg)→fix_root
                                           : &frchain_now→fix_root);

- `merge_data_into_text`. If the user specified (with the `-R` flag) that data sections should go into the text segment to make the data read-only, we should merge the data and the text sections. This is done now.

- we keep calling `relax_segment` until we record that there isn't any more changes.

---

[19]Unix is fond of mythological names: We have magic numbers, Elfs, dwarfs, daemons...

```
1          rsi.pass = 0;
2          while (1) {
3              rsi.changed = 0;
4              map_over_sections(relax_seg,&rsi);
5              rsi.pass++;
6              if (!rsi.changed)
7                  break;
8          }
```

`rsi` is a variable of type `struct relax_seg_info`[20]. The function `map_over_sections` just calls the function given in argument for each section in the output file.

- `size_seg`. Now that the address and size of all fragments is known, we can calculate the total size of each segment.

- `dwarf2dbg_final_check`. This is interesting stuff. There is a proposal from Alexandre Oliva[21] that introduces the concept of "view numbers" where the same program counter can belong to several views. The underlying need for this are inlined functions, where the inlined code can belong to the current function, or it can be understood as part of the inlined function, allowing the debugger to trace through the inlined function as if it were a normal function call.[22]

- `create_obj_attrs_section` creates a section to hold all program attributes. The attributes should refer to the CPU type where the program can run.

- All relocations refer to symbols. So we have to resolve symbols before doing the relocations. this is done

```
1          if (symbol_rootP) {
2              symbolS      *symp;
3
4              for (symp = symbol_rootP; symp; symp = symbol_next(symp))
5              resolve_symbol_value(symp);
6          }
7          resolve_local_symbol_values();
8          resolve_reloc_expr_symbols();
```

The `resolve_symbol_value` function tries to determine the value of a possibly very complex expression and assigning it to the symbol.

The `resolve_local_symbol_value` organizes a traversal of the hash symbol table to resolve all local symbols.

---

[20] A very simple structure:
`struct relax_seg_info {int pass; int changed;}`
The `pass` member is incremented but never used. It is there to allow debugging infinite loops that could arise.
[21] https://www.fsfla.org/~lxoliva/
[22] The whole proposal text is here:
This proposal introduces a new implicit column to the line number table, namely "view numbers", so that multiple program states can be identified at the same program counter, and extends loclists with means to add view numbers to address ranges, enabling locations to start or end at specific views.
This may improve debug information, enabling generators to indicate inlined entry points and preferred breakpoints for statements even if instructions associated with the corresponding source locations were not emitted at the given PC, and to emit variable locations that indicate the initial values of inlined arguments, and side effects of operations as they would be expected to take effect from the source code, even when multiple statements have their side effects all encoded at the same PC: with view numbers, debug information consumers may be able to logically advance the perceived program state, so as to reflect user-expected changes specified in the source code, even if the operations were reordered or optimized out in the executable code.

- `elf_frob_file_before_adjust` will go through all symbols and will eliminate unneeded versions of versioned symbols.

- `adjust_reloc_syms` will go through all symbols and try to replace the references to symbols by references to the section symbol + offset.

- `fix_segment`. This function will go through all fixups of a segment and resolve those that can be resolved at this stage. For instance if a fragment's address has been resolved any fixup mentioning this address can be resolved too. Or when a symbol has been resolved, the fixup can be eliminated.

- Now it's time to write the symbol table. The code goes through all symbols checking that:

  1. Local labels are defined.
  2. Splice out symbols that should be ignored, like symbols that were equated to bss or to undefined symbols.
  3. `elf_frob_symbol` Will take care of symbol versioning and associated complexities...
  4. Take care of "warning" symbols, i.e. symbols that are there just to generate a warning. They are just skipped.
  5. Take care of the infinite possibilities of bugs... For instance there could be symbols that were emitted before an alignment that ended as a zero byte alignment. They are unnecessary. Get rid of them.

- `set_symtab`. This function counts the symbols, and allocates a table that will be used to store the symbols to be written out.

- `elf_frob_file`. This function does two things:

  1. In the case we are emitting `stabs` debug information, fill the header with the number of stabs, and other information.
  2. Do the checks necessary for putting in the elf file flags, the necessary description of the target machine.

- `write_relocs`. Write out all relocations.

- `elf_frob_file_after_relocs`. If we have a group of sections, and we have established the number of relocations, it could be that a section has no longer any relocations or that the number of relocations has changed. In that case the size of the group must be adjusted.

- Once the relocations have been prepared for writing, we can compress the debug section, if necessary. This must be done before anything is written out since it makes the size of the file change.

- `write_contents`. this function organizes the actual writing out of the data. It writes the fixups, the section contents and the fill data to align sections. This is done using the `set_section_contents` function. This function makes some checks and then calls `elf_set_section_contents`.

  This one makes some further checks, copies the contents into the image of the section in RAM and calls `generic_set_section_contents` that makes some checks and positions the file pointer at the correct position, then finally calls `bfd_bwrite` that will send the data to the disk with `fwrite`.

Described like that, this whole bunch of stacked procedures seems bloated but it is not. Each one takes a piece of the work. The GAS code is written by defensive programmers and defensive programming is not a bad idea. It pays when you have clear error messages and not bad results. Bugs provoked by missing sanity tests are very difficult to find, bugs with clear error messages spare you the time consuming search for "where is the bug?". They pop up with an error message and you instantly know where the problem is.

1.10   Assembler directives

Directives are defined in a table of structures of type `pseudo_typeS`:

Listing 1.7: struct pseudo_typeS

```
1 typedef struct _pseudo_type {
2     /* Assembler mnemonic in lower case, without the implicit dot '.' */
3     const char *poc_name;
4     /* Function that will be called to handle this directive */
5     void (*poc_handler) (int);
6     /* Value to pass to handler. */
7     int poc_val;
8 } pseudo_typeS;
```

The assembler defines several tables of this structures. We have the main one, `potable` and several others: `cfi_pseudo_table` for the debug information, `elf_pseudo_table` for the directives concerning the object code format, and a `riscv_pseudo_table` for several riscv specific directives.

All of them will be called from `read_a_source_file` function. Here is the relevant code snippet:

```
1  if (*s == '.') {
2      /* PSEUDO - OP. WARNING: Next_char may be end-of-line. We lookup the pseudo-op
3       * table with s+1 because we already know that the pseudo-op begins with a '.' */
4      pop = str_hash_find(po_hash,s + 1);
5      if (pop && !pop→poc_handler)
6          pop = NULL;
7  // ... code elided
8      /* Input_line is restored. Input_line_pointer→1st non-blank char after
9       * pseudo-operation. */
10     (*pop→poc_handler) (pop→poc_val);
11 }
```

The `po_hash` table is built when the assembler starts, containing the different tables mentioned above. The function that does this is very simple:

Listing 1.8: pop_insert

```
1  static void pop_insert(const pseudo_typeS * table)
2  {
3      const pseudo_typeS *pop;
4      for (pop = table; pop→poc_name; pop++) {
5          if (str_hash_insert(po_hash,pop→poc_name,pop,0) ≠ NULL) {
6              if (!pop_override_ok)
7                  as_fatal("error constructing %s pseudo-op table",
8                      pop_table_name);
9          }
10         //else printf("%s\n",pop→poc_name);
11     }
12 }
```

Just a loop inserting each member of the given table. The variable `pop_override_ok` is a global that will be zero if we don't accept any insertions with the same name.

That function will be called from `pobegin`, that looks like this:

Listing 1.9: pobegin

```
1  static void pobegin(void)
2  {
3      po_hash = str_htab_create();
4      pop_table_name = "md"; /* Do the target-specific pseudo ops. */
5      pop_override_ok = 0;  /* Do not accept any shadowing */
6      pop_insert(riscv_pseudo_table);
7      pop_table_name = "obj"; /* Object specific. Skip any already present */
8      pop_override_ok = 1;
9      pop_insert(elf_pseudo_table);
10     pop_table_name = "standard";/* Now portable ones. Skip any already present */
11     pop_insert(potable);
12     pop_table_name = "cfi";   /* Now CFI ones. */
13     pop_insert(cfi_pseudo_table);
14 }
```

This code ensures that machine specific directives shadow any object or standard directives since they are inserted first. The global variable `pop_table_name` is used for error messages only, as we have seen in the code of `pop_insert`[23].

1.10.1 .align, .p2align, p2alignw, p2alignl

Entries in the table:

```
1      {"align",s_align_ptwo,0},
2      {"p2align",s_align_ptwo,0},
3      {"p2alignw",s_align_ptwo,-2},
4      {"p2alignl",s_align_ptwo,-4},
```

These four entries lead to calls to the same function, albeit with different arguments.

```
1      void s_align_ptwo(int arg) { s_align(arg,0); }
```

`s_align` receives two arguments. The first one, if positive, defines a default alignment. If negative, it defines a length of a fill pattern. The second argument, if positive, should be interpreted as a byte boundary, not as a power of two. Now, if the first argument was negative, the second argument should contain the fill pattern.

All arguments are optional. If none is given, the alignment defaults to the argument that will be given to `s_align_ptwo`.

The `s_align` function calls eventually `do_align`. The comment at the start of this function says it all:

```
1      /* Guts of .align directive: N is the power of two to which to align. A value
2       * of zero is accepted but ignored: the default alignment of the section will
3       * be at least this. FILL may be NULL,or it may point to the bytes of the fill
4       * pattern. LEN is the length of whatever FILL points to,if anything. If LEN
5       * is zero but FILL is not NULL then LEN is treated as if it were one. MAX is
6       * the maximum number of characters to skip when doing the alignment,or 0 if
7       * there is no maximum. */
```

But we aren't done yet. `do_align` calls `md_do_align` that is actually a macro:

---

[23]Looking at this code I do not quite understand why there isn't an additional parameter to `pop_insert` instead of a global variable. Probably it is difficult to modify the syntax for all back-ends of GAS.

```
1    #define md_do_align(N, FILL, LEN, MAX, LABEL)   \
2    if ((N) ≠ 0 && !(FILL) && subseg_text_p (now_seg)) \
3    {                                                  \
4        if (riscv_frag_align_code (N))           \
5        goto LABEL;                              \
6    }
```

The actual call sequence looks like this:

```
1    md_do_align(n,fill,len,max,just_record_alignment);
```

Yes, there is *still* another level. And in this level we discover that we just can't align anything. The riscv linker changes the size of some instructions, allowing compressed instructions where possible, what will change the adresses of all subsequent instructions. So, the only thing that `riscv_frag_align_code` can do is just emit an alignment relocation that will tell the linker that this fragment needs to be aligned.

Obviously, all this lengthy process could be simplified a lot, but I have tried to keep the original structure, it may be useful to understand GAS in the context of other machines.

1.10.2   .ascii, .asciiz, .string, .string8, .string16, .string32, .string64

All these directives lead to the `stringer` function. The entries are as follows:

```
1    {"ascii",stringer,8 + 0},
2    {"asciz",stringer,8 + 1},
3    {"string8",stringer,8 + 1},
4    {"string16",stringer,16 + 1},
5    {"string32",stringer,32 + 1},
6    {"string64",stringer,64 + 1},
```

The `stringer` receives an odd argument when it should append a zero to its output. The numbers represent how many bytes should it use for each character. The input is done by following `input_line_pointer` that is a global pointer to the assembler text. `stringer`'s code is easy to follow, so it is not further described here.

1.10.3   .bss

Changes (if necessary) the current section to he bss. This section contains uninitialized data and will set to zero at the program's start by the loader. This directive will calls `obj_elf_bss`, a small function that realizes this change.

```
1  /* Change to the .bss section. */
2  static void obj_elf_bss(int i ATTRIBUTE_UNUSED)
3  {
4      int     temp;
5      obj_elf_section_change_hook();
6      temp = get_absolute_expression(); // Optional subsection. Normally blank
7      subseg_set(bss_section,(subsegT) temp);
8      demand_empty_rest_of_line();
9  }
```

Function `obj_elf_section_change_hook` remembers the section before the change so that a `.section previous` directive can find it. See §1.10.21, page 60 for `subseg_set`.

1.10.4   .byte, .dc, .dc.a, .dc.b, .dc.d, .dc.l, .dc.s, .dc.w, etc

```
1    {"byte",cons,1},
2    {"dc",cons,2},
```

```
3      {"dc.a",cons,0},
4      {"dc.b",cons,1},
5      {"dc.d",float_cons,'d'},
6      {"dc.l",cons,4},
7      {"dc.s",float_cons,'f'},
8      {"dc.w",cons,2},
9      {"hword",cons,2},
10     {"int",cons,4},
11     {"octa",cons,16},
12     {"quad",cons,8},
13     {"short",cons,2},
14     {"long",cons,4},
15     {"quad",cons,8},
16     {"word",cons,2},
17     {"2byte",cons,2},
18     {"4byte",cons,4},
19     {"8byte",cons,8},
20     {"half",cons,2},
```

GAS likes to be compatible. The consequence of that is the above list. All those directives lead to the same function. You can write a two byte constant with .short, .dc, .dc.w, .hword, .2byte and .half. [24]

So, what does this cons function do?

It is a fairly simple function, consisting in a loop reading expressions separated by commas. In the original code, the crucial lines look like this:

```
1      do {
2          TC_PARSE_CONS_RETURN_TYPE ret = TC_PARSE_CONS_RETURN_NONE;
3          ret = TC_PARSE_CONS_EXPRESSION(&exp,(unsigned int)nbytes);
4
5          if (rva) {
6              if (exp.X_op == O_symbol)
7                  exp.X_op = O_symbol_rva;
8              else
9                  as_fatal(("rva without symbol"));
10         }
11         emit_expr_with_reloc(&exp,(unsigned int)nbytes,ret);
12         ++c;
13     } while (*input_line_pointer++ == ',');
```

The problem with macros such as those here (lines 2 and 3), is that they make impossible to know what is going on actually in the program. Translated into C, these two lines expand into:

```
1      do {
2          bfd_reloc_code_real_type ret = BFD_RELOC_NONE;
3          ret = (expr(0,&exp,expr_normal),BFD_RELOC_NONE);
4          ... // The rest is the same
5      }
```

Line 2 shows that ret is a member of the enumeration bfd_reloc_code_real_type that is assigned zero.

Line 3 is a comma expression, that in its first statement evaluates a call to expr, that reads an expression from input_line_pointer and in the second (and last) one evaluates to a constant that is assigned to the ret variable.

Besides this small problem, cons doesn't present any big difficulties.

---

[24]The directives .2byte, .4byte, etc are used by gcc mainly within the debug information.

1.10.5   .data

Tells the assembler to change (if necessary) to the data section. This directive is handled by
the `s_data` function:

Listing 1.10: s_data

```
 1 static void s_data(int ignore ATTRIBUTE_UNUSED)
 2 {
 3     segT        section;
 4     int     temp;
 5
 6     temp = get_absolute_expression();
 7     if (flag_readonly_data_in_text) {
 8         section = text_section;
 9         temp += 1000;
10     } else section = data_section;
11     subseg_set(section,(subsegT) temp);
12     demand_empty_rest_of_line();
13 }
```

If the data section is readonly, a special subsegment in the text section is used.[25]
See §1.10.21, page 60 for `subseg_set`.

1.10.6   debug, extern, format, lflags, name, noformat, spc, xref

All those directives have only *one* thing in common: they are completely **ignored** by the
GNU assembler. It just advances the line pointer to the end of the line.
Why this?
As you guessed, it is just a compatibility feature.

```
 1     {"debug",s_ignore,0},
 2     {"extern",s_ignore,0},/* We treat all undef as ext. */
 3     {"format",s_ignore,0},
 4     {"lflags",s_ignore,0},/* Listing flags. */
 5     {"name",s_ignore,0},
 6     {"noformat",s_ignore,0},
 7     {"spc",s_ignore,0},
 8     {"xref",s_ignore,0},
```

As the comment shows, declaring a symbol *extern* doesn't do anything. The assembler
declares all undefined symbols `extern`. This implies that if a misspelled name appears in
your assembler program you will see it at link time, not at assembly time. No big deal
anyway.
More problematic is ignoring directives like `xref` or `debug`. These directives are expected
to *do* something, and silently accepting and ignoring them will provoke in people that expect
some result from their directives to search in vain **why** the assembler is not doing what they
have written.
This is worst than a clear error message: "unknown directive". Much worst. That is why
those directives aren't accepted any more in `tiny-asm`, except the **extern** one, because that
one *does* what the user is expecting.

1.10.7   equ, equiv, eqv, set

`.equ symbol, expression`

---

[25]It could be possible to set the flags of the data section to read-only, but GAS prefers this methods for
portability reasons... Not all systems probably support that.

This directive sets the value of symbol to expression. It is synonymous with '.set';

This is something similar to

`#define␣name␣another_name`

in C. There are some subtleties though. The `equiv` directive will complain if the first symbol is already defined. The `eqv` directive announces to the assembler that the right hand side is a forward reference.

```
1    {"equ",s_set,0},
2    {"equiv",s_set,1},
3    {"eqv",s_set,-1},
4    {"set",s_set,0},
```

The `s_set` function is simple to follow.

1.10.8   globl

```
     .global symbol[, symbol, symbol, ...]
```

`.global` makes the symbol visible to ld. If you define symbol in your partial program, its value is made available to other partial programs that are linked with it. Otherwise, symbol takes its attributes from a symbol of the same name from another file linked into the same program.

In the `potable` we have:

```
1  Table: (potable)
2    {"global",s_globl,0},
3    {"globl",s_globl,0},
```

Unix has a big problem with vowels. They are shunned everywhere. Why write `globl`? Is the absence of a poor vowel *really* that shorter? Or is the necessary effort of *remembering its absence* when writing the program (taking precious memory space in the brain) even costlier?

Well, at least the assembler lets you decide, you can use both.

Coming back to our source code, the `s_globl` function is a very simple and short one. It just scans names and adds the `EXTERNAL` bit to each of the symbols scanned in a loop (not shown).

```
1    if ((name = read_symbol_name()) == NULL)
2        return;
3    symbolP = symbol_find_or_make(name);
4    S_SET_EXTERNAL(symbolP);
```

1.10.9   attach_to_group

```
Syntax:
     .attach_to_group <name>
Table: (elf_pseudo_table)
     {"attach_to_group",obj_elf_attach_to_group,0},
```

This will attach the current section to the named group. If the group doesn't exist it will be created. The `obj_attach_to_group` function just changes a pointer and the flags of the current section. The relevant lines (without error checking etc) of this function are:

```
1    elf_group_name(now_seg) = gname;
2    elf_section_flags(now_seg) |= SHF_GROUP;
```

1.10.10   .comm, .common, .lcomm

Only the directive `.comm` and `.lcomm` are documented in the official documentation.

```
Syntax:
.comm symbol , length
Table: (elf_pseudo_table)
{"comm",obj_elf_common,0},
{"common",obj_elf_common,1},
{"lcomm",obj_elf_lcomm,0},
```

`.comm` declares a common symbol named symbol. When linking, a common symbol in one object file may be merged with a defined or common symbol of the same name in another object file. If ld does not see a definition for the symbol–just one or more common symbols–then it will allocate length bytes of uninitialized memory. length must be an absolute expression. If ld sees multiple common symbols with the same name, and they do not all have the same size, it will allocate space using the largest size.

`.lcomm` (local common) has the same syntax as `comm` but the symbol is just declared in the bss section and not make visible.

`.common` is a synonym for `comm` even if it receives a different argument because actually... the argument is ignored!

```
1  static void obj_elf_common(int is_common ATTRIBUTE_UNUSED)
2  {
3      s_comm_internal(0,elf_common_parse);
4  }
```

The function `s_comm_internal` is mostly parsing and error checking. The essential lines are at the end:

```
1      S_SET_VALUE(symbolP,(valueT) size);
2      S_SET_EXTERNAL(symbolP); // This is absent in lcomm
3      S_SET_SEGMENT(symbolP,bfd_com_section_ptr);
```

1.10.11   hidden

```
Syntax:
    .hidden symbol-name [, symbol-name, ...]
```

Sets the visibility of a symbol, i.e. if it is visible for modules outside the one being assembled. This directive implies *protected* as well.

It is handled by the `obj_elf_visibility` function.

1.10.12   .ident

```
Syntax:
.ident "A string"
Table: elf_pseudo_table
{"ident",obj_elf_ident,0},
```

This directive writes any string into the comments section of the file. For instance:

```
.ident "I love you Barbie"
```

Assembling your file, you can display it to your girlfriend with:

```
star64:~/tiny-asm$ asm sample.s
star64:~/tiny-asm$ objdump -s -j .comment a.out

a.out:     file format elf64-littleriscv

Contents of section .comment:
0000 0049206c 6f766520 796f7520 42617262  .I love you Barb
0010 696500                               ie.
```

She will be surely greatly impressed... The `obj_elf_ident` function creates the `.comments` section if it is not already present. Then, it calls the stringer for parsing. You can write any number of these comments.

1.10.13 `insn`

```
Syntax:
.insn type, operand [,...,operand_n]
.insn insn_length, value
.insn value
Table: riscv_pseudo_table
{"insn",s_riscv_insn,0},
```

This directive assembles an unknown instruction into the instruction stream. For instance, using the first type of syntax, let's say you want to want to issue the instruction `add a0,a1,a2`. First, you have to look up what type of instruction it is. It is an "R" type of instruction. You write as first argument "r".

After the type, you should give the fields of the R format that are fixed: the opcode, the extra-3 and the extra-7 fields. In this case both are zero. And then, you should give the arguments of the instruction, i.e. the register names.

You should write then:

```
1      .insn r 0x33, 0, 0, a0,a1,a2
```

Note that there isn't any comma between the "r" and the 0x33! The "r" is understood as a part of the opcode.

Now where does this 0x33 come from?

If you go to the opcode table, and search for the "`add`" entries, you will see several of them. You should choose this one:

```
1      {"add",0,INSN_CLASS_I,"d,s,t",MATCH_ADD,MASK_ADD,match_opcode,0},
```

since the other ones further up are compressed (`INSN_CLASS_C`) and we do not want compression. The opcode is in the `MATCH_ADD` field, that is defined in asm.h to be... 0x33. After the two zeroes of the bit fields associated with class "R" we write the 3 required register names.

How can we know that this is OK?

Easy: just write following assembler program:

```
1      add a0,a1,a2
2      .insn r 0x33, 0, 0, a0,a1,a2
```

Then assemble it, and then display the contents with

```
1      star64:~/tiny-asm$ objdump -d sample.o
2
3      sample.o:    file format elf64-littleriscv
4
```

```
5     Disassembly of section .text:
6
7     0000000000000004 <main>:
8     4:  00c58533          add a0,a1,a2
9     8:  00c58533          add a0,a1,a2
```

We find the 0x33 in the lower 7 bits of the opcode field.

The other syntax variants of the directive are trivial.

Another example: the instruction `addw a0,a1,a2`. The entry in the opcode table is:

␣`{"addw",64,INSN_CLASS_I,"d,s,t",MATCH_ADDW,MASK_ADDW,match_opcode,0},`

We look the constant `MATCH_ADDW` in asm.h, what gives 0x3b. So, as shown in 1.4 page 26, the two fields "extra-3" and "extra-7" are zero. We write then:

```
1     addw a0,a1,a2
2     .insn r 0x3b, 0, 0, a0,a1,a2
```

and when disassembling we get:

```
1     4:  00c58533          add  a0,a1,a2
2     8:  00c58533          add  a0,a1,a2
3     c:  00c5853b          addw a0,a1,a2
4     10: 00c5853b          addw a0,a1,a2
```

The `s_riscv_insn` function essentially just calls `riscv_ip`. The lookup of the "r" letter yields an entry into the `riscv_insn_types` table, that looks like this:

␣`{"r",0,INSN_CLASS_I,"O4,F3,F7,d,s,t",0,0,match_opcode,0},`

where we see the length of the instruction (4 bytes) and the names of the 3 and 7 bits extra fields. Then, we find the usual denominations ("d,s,t") that we discussed when analyzing the string arguments to each opcode, see table 1.7 page 39.

Conclusion This is quite difficult stuff, because precisely the point of an assembler is to avoid you to encode manually the instructions. It is a *very* error prone process. And in the end if you write:

```
        .word 0xc58533
```

it will work in the same way. The justification advanced by the GNU folks is that in future versions of the assembler you will *not* see this as just data, but as a real instruction.

Maybe. But I think a more real justification is that the riscv architecture itself allows for instruction extensions, and has a whole part of the instruction space available for standard or non-standard extensions to the accepted opcodes. The existence of an `insn` extension here, would allow the assembler to assemble code that uses those extensions.

1.10.14   internal

Syntax:
`.internal symbol-name [, symbol-name, ...]`

Sets the visibility of a symbol, i.e. if it is visible for modules outside the one being assembled. This directive implies *protected* as well.

It is handled by the `obj_elf_visibility` function. function.

1.10.15   loc

Syntax:
`    .loc fileno lineno [column] [options]`
`Table:  elf_pseudo_table`
`      {"loc",dwarf2_directive_loc,0},`

Ahhh the old days, when everything was simple and clear! Remember when the debug information for the line number was just a triplet of address, file, line?

Say goodbye to that now, and welcome to DWARF[26]. The line number is a series of instructions to an interpreted language executed by a state machine.

Yes, you read correctly.

Conceptually we have a table of addresses, each one with as many properties as desired:

address	source file	source line	source column	state-ment?	basic block	... other columns
0x40260	1	23	12	0	0	
0x40264	1	23	12	1	1	

.
.
.

"... we design a byte-coded language for a state machine and store a stream of bytes in the object file instead of the matrix. This language can be much more compact than the matrix. When a consumer of the line number information executes, it must "run" the state machine to generate the matrix for each compilation unit it is interested in."[27]

The arguments for the `.loc` directive then, are as follows:

- `fileno`. The file index in the assembler's file table.
- `lineno`. Line number.
- `column`. This field is optional.
- `options`. They are the following:

    - `basic_block` This instruction represents the start of a basic block.[28]
    - `prologue_end`. End of the setup of the stack frame. This changes the state of the interpreter. In C it corresponds to the opening brace of a function.
    - `is_stmt␣value` Start of a statement sequence.
    - `isa␣value` Sets the instruction set architecture register to *value*
    - An unsigned integer identifying the block to which the current instruction belongs. Discriminator values are assigned arbitrarily by the DWARF producer and serve to distinguish among multiple blocks that may all be associated with the same source file, line, and column. Where only one block exists for a given source position, the discriminator value should be zero. This is necessary because the compiler can move instructions around to keep the pipeline busy. Then, instructions belonging a one or several blocks could be mixed.
    - `view`. This is not in the 4th edition of the DWARF standard nor in the 5th. It has been added later probably. The documentation says:

        This option causes a row to be added to .debug_line in reference to the current address (which might not be the same as that of the following assembly instruction), and to associate value with the view register in the .debug_line state machine. If value is a label, both the view register and the label are set to the number of prior .loc directives at the same program location. If value is the literal 0, the view register is set to zero, and the assembler asserts that there aren't any prior .loc directives at the

---

[26] Critiques to DWARF abound. See for instance: https://tobast.fr/doc/publications/oopsla19-dwarf.pdf

[27] DWARF Debugging Information Format Version 4, page 108

[28] A basic block is a sequence of instructions where only the first instruction may be a branch target and only the last instruction may transfer control. A procedure invocation is defined to be an exit from a basic block.

same program location. If value is the literal -0, the assembler arrange
for the view register to be reset in this row, even if there are prior .loc
directives at the same program location.

Crystal clear isn't? [29]

The function `dwarf2_directive_loc` is interesting as an example of the functions used
to parse data within the assembler. To make things a bit clearer I have added comments to
everything.

Listing 1.11: Parsing .loc directive

```
1 static void   dwarf2_directive_loc(int dummy ATTRIBUTE_UNUSED)
2 {
3     /* If we see two .loc directives in a row,force the first one to be output now.*/
4     if (dwarf2_loc_directive_seen) dwarf2_emit_insn(0);
5     offsetT filenum = get_absolute_expression();                 ①    get_absolute_expression
6     SKIP_WHITESPACE();
7     offsetT line = get_absolute_expression();
8     /* error checking: */
9     if (filenum < 1) {
10        /* DWARF5 specifies that a file number of zero indicates that
11           the file is unknown */
12        if (filenum == 0 && dwarf_level < 5) dwarf_level = 5;
13        /* All other values are just nonsense */
14        if (filenum < 0 || DWARF2_LINE_VERSION < 5) {
15            as_bad("file number less than one");
16            return;
17        }
18    }
19    if ((valueT) filenum ≥ files_in_use || files[filenum].filename == NULL) {
20        as_bad("unassigned file number %ld",(long)filenum);
21        return;
22    }
23    gas_assert(debug_type == DEBUG_NONE);                         ②    debug_type
24    current.filenum = filenum;
25    current.line = line;                                          ③    current
26    current.discriminator = 0;
27    SKIP_WHITESPACE();
28    /* test for an optional column number */
29    if (ISDIGIT(*input_line_pointer)) {
30        /* We have the optional column number */
31        current.column = get_absolute_expression(); SKIP_WHITESPACE();
32    }
33    /* Now we start parsing the "options" field */
34    while (ISALPHA(*input_line_pointer)) {



35        char        *p,c = get_symbol_name(&p);                  ④    get_symbol_name
36        offsetT       value;
37        if (strcmp(p,"basic_block") == 0) {
38            current.flags |= DWARF2_FLAG_BASIC_BLOCK;
39            *input_line_pointer = c; // Restore character
40        } else if (strcmp(p,"prologue_end") == 0) {
```

---

[29]There is no other documentation anywhere that would state what this thing does in a more understand-
able way... Sorry.

```
41             if (dwarf_level < 3) dwarf_level = 3;
42             current.flags |= DWARF2_FLAG_PROLOGUE_END;
43             *input_line_pointer = c;
44         } else if (strcmp(p,"epilogue_begin") == 0) {
45             if (dwarf_level < 3) dwarf_level = 3;
46             current.flags |= DWARF2_FLAG_EPILOGUE_BEGIN;
47             *input_line_pointer = c;
48         } else if (strcmp(p,"is_stmt") == 0) { // is_stmt <boolean value>
49             (void)restore_line_pointer(c);
50             value = get_absolute_expression();
51             if (value == 0) current.flags &= ~DWARF2_FLAG_IS_STMT;
52             else if (value == 1) current.flags |= DWARF2_FLAG_IS_STMT;
53             else { as_bad("is_stmt value not 0 or 1"); return; }
54         } else if (strcmp(p,"isa") == 0) { // "isa" numbers are defined by the ABI
55             if (dwarf_level < 3) dwarf_level = 3;
56             (void)restore_line_pointer(c);
57             value = get_absolute_expression();
58             if (value ≥ 0) current.isa = value;
59             else {
60                 as_bad("isa number less than zero");
61                 return;
62             }
63         } else if (strcmp(p,"discriminator") == 0) {
64             (void)restore_line_pointer(c);
65             value = get_absolute_expression();
66             if (value ≥ 0) current.discriminator = value;
67             else {
68                 as_bad(("discriminator less than zero"));
69                 return;
70             }
71         } else if (strcmp(p,"view") == 0) {
72         /* Now we parse the mysterious "view" statement. */
73             symbolS      *sym;
74             (void)restore_line_pointer(c);
75             SKIP_WHITESPACE();
76             if (ISDIGIT(*input_line_pointer) || *input_line_pointer == '-') {
77                 /*
78                  * Now, we expect either "0" or "-0"
79                  */
80                 bool      force_reset = *input_line_pointer == '-';
81                 value = get_absolute_expression();
82                 if (value ≠ 0) {
83                     as_bad("numeric view can only be asserted to zero"); return;
84                 }
85                 if (force_reset && force_reset_view) sym = force_reset_view;
86                 else {
87                     sym = symbol_temp_new(absolute_section,&zero_address_frag,value);
88                     if (force_reset) force_reset_view = sym;
89                 }
90             } else { // We have a symbol that will be put into the "view" register.
91                 char          *name = read_symbol_name();
92                 // We silently accept .loc view followed by nothing, without
93                 // any warning or error.
94                 if (!name) return;
95                 sym = symbol_find_or_make(name);
96                 free(name); // read_symbol_name allocates memory for its result
97                 if (S_IS_DEFINED(sym) || symbol_equated_p(sym)) {
98                     if (S_IS_VOLATILE(sym)) sym = symbol_clone(sym,1);
```

⑤   restore_line_pointer

```
 99                    else if (!S_CAN_BE_REDEFINED(sym)) {
100                        as_bad("symbol '%s' is already defined",S_GET_NAME(sym));
                           return; }
101                    }
102                    S_SET_SEGMENT(sym,undefined_section); S_SET_VALUE(sym,0);
103                    symbol_set_frag(sym,&zero_address_frag);
104                }
105                current.u.view = sym;
106            } else {
107                as_bad("unknown .loc sub-directive '%s'",p);
108                (void)restore_line_pointer(c); return;
109            }
110            /* This macro differs from SKIP_WHITESPACE in that it ignores a double quotes
111             * after the name */
112            SKIP_WHITESPACE_AFTER_NAME();
113        }
114        demand_empty_rest_of_line();                          6      demand_empty_rest_of_line
115        dwarf2_any_loc_directive_seen = dwarf2_loc_directive_seen = true;
116        /* If we were given a view id,emit row now */
117        if (current.u.view) dwarf2_emit_insn(0);
118 }
```

1. The function `get_absolute_expression` reads a constant from the global line pointer
   and sets it to just after the last character of the constant. If any error occurs, it emits
   an error message and returns zero. If the expression is absent, it returns zero without
   any error message. This makes many things default to a convenient zero.

2. The value in the global variable `debug_type` will be turned off by the function
   `dwarf2_directive_filename`, and if we don't have a dwarf style `.file` directive in
   between, then `files_in_use` will be zero and the error in line 15 will trigger.
   Note: The global `debug_type` will be left to zero, effectively disabling the emission of
   any debug information by the assembler.

3. `current` is a structure of type `dwarf2_line_info` that holds the current context. We
   update it AFTER all error checking is done, to preserve a correct context in case of
   an error

4. The `get_symbol_name` function parses a symbol using `input_line_pointer`. It writes
   a zero immediately after the expected symbol and returns the value of the character
   at the position where zero was written. Its result is left in its pointer argument, that
   will point to the start of the symbol.

5. `restore_line_pointer` writes the previous character into the line pointer, advances
   to the next character and if it is a double quote, it ignores it by advancing again.

6. `demand_empty_rest_of_line` advances the line pointer to the next newline character.
   If there is anything in that part of the line it will complain with an error.

1.10.16   local

```
Syntax:
    .local symbol,symbol,...
   Table: elf_pseudo_table
   {"local",obj_elf_local,0},
```

This directive makes the given symbol a local symbol, not visible to other modules. Since
all symbols are local unless declared extern or undefined, the utility of this is not clear.
    The important lines of `obj_elf_local` are:

```
1    symbolP = get_sym_from_input_line_and_check();
2    S_CLEAR_EXTERNAL(symbolP);
3    symbol_get_obj(symbolP)→local = 1; // See below
```

The function `symbol_get_obj` returns a pointer to a small structure that keeps several disjoint pieces of information about a symbol, among them, whether it is a local symbol. We can't access directly the field because of local symbols precisely.

```
1  /* Get a pointer to the object format information for a symbol. */
2  static struct elf_obj_sy *symbol_get_obj(symbolS * s)
3  {
4      if (s→flags.local_symbol)
5          s = local_symbol_convert(s);
6      return &s→x→obj;
7  }
```

If the symbol is local, we have to convert it first. See §1.4.2, page 17 to see where this piece fits in the general schema of things.

1.10.17 .option

Syntax:
.option <option-name>
Table: riscv_pseudo_table
{"option",s_riscv_option,0},

This handles the update of several riscv related options. The example given in the GAS documentation runs as follows:

```
1    .option push
2    .option norelax
3    la gp, __global_pointer$
4    .option pop
```

In the "relaxation" process, the assembler tries to find shorter, compressed, sequences for instructions. It tries to substitute loading a global directly, for a shorter sequence that loads the address from an offset from the `__global_pointers` table. The problem arises when you want to load the address of the `__global_pointers` table itself. In that case you do NOT want the assembler to pick an offset since the `__global_pointers` table is not loaded. Then, you disable for a single instruction, this feature and all goes well.

Of course this happens only to people that are writing the startup code, or other assembler wizards. This kind of fiddling is *for them only*. Please do not mess around with any of this things yourself.

The code for `s_riscv_options` is trivial: a long series of:

```
1    if (strcmp(name,"push") == 0) { /* code for push option */}
2    else if (strcmp(name,"pop") == 0 {/* code for pop option */})
3    etc...
```

Other interesting values for `.option` are:

- `pic` or `nopic`. Enable or disable the position independent code generation. This corresponds to the `-fPic` flag in gcc.
- `rvs` or `norvc`. Enable or disable the compressed instructions generation.
- `relax` or `norelax`. Enable or disable relaxation.
- `csr-check` or `nocsr-check`. Enables or disables checking when using the CSR registers.
- Etc. There are many other obscure things to peruse here: https://sourceware.org/binutils/docs/as/RISC_002dV_002dDirectives.html

1.10.18   org

Syntax:
```
.org new-location-counter , fill byte
```

Advance the location counter of the current section to *new-location-counter*. It should be either an absolute expression or an expression with the same section as the current subsection. That is, you can't use .org to cross sections: if it has the wrong section, the .org directive is ignored. To be compatible with former assemblers, if the section of new-lc is absolute, as issues a warning, then pretends the section of new-lc is the same as the current subsection.

1.10.19   protected

Syntax:
```
.protected symbol-name [, symbol-name, ...]
```

Sets the visibility of a symbol, i.e. if it is defined for modules outside the one being assembled, the definition in this module will be used.

It is handled by the `obj_elf_visibility` function. function.

1.10.20   reloc

The documentation of GAS says about this directive:

Syntax:
```
.reloc offset, reloc_name[, expression]
```

> Generate a relocation at `offset` of type `reloc_name` with value `expression`. If offset is a number, the relocation is generated in the current section. If offset is an expression that resolves to a symbol plus offset, the relocation is generated in the given symbol's section. `expression`, if present, must resolve to a symbol plus addend or to an absolute value, but note that not all targets support an addend. e.g. ELF REL targets such as i386 store an addend in the section contents rather than in the relocation. This low level interface does not support addends stored in the section.

The last part of the description needs maybe a clarification. In the x86 systems, the addend to the relocation is stored in the data itself, so the program loader should only add the load address. This makes constructing relocations with an addend impossible.

Why is this directive necessary? Mystery, the official documentation gives no examples, and (with my limited imagination) I just can't figure out its use.[30]

Well, the only way of figuring out this, is to use it and see what it does. I write this in C:

```
1    long double mm = 3.1415926534564321;
2    int main(void) {}
```

I compile it with: `gcc -c -S tld.c` and obtain a `tld.s` assembler file:

```
1    .file   "tld.c"
2    .size   mm, 16
3    mm:
4    .word   0
```

---

[30]In the documentation of the ARM assembler I found a similar RELOC directive that (seems) to force the assembler to put either a symbol or the preceding instruction at a specific address, like the `.org` directive, but I am not sure

```
 5    .word  -1610612736
 6    .word  -1253836416
 7    .word  1073779231
 8    .text
 9    .globl main
10    .type  main, @function
11    main:
12    jr ra
```

We have then, a long double in the data section. I start gdb:

```
(gdb) print &mm
$1 = (<data variable, no debug info> *) 0x2aaaaac010 <mm>
```

OK, now I add the line: `.reloc 8,BFD_RELOC_32,mm` after the last `.word` in the definition of mm. I start gdb with the new program and...

```
(gdb) print &mm
$1 = (<data variable, no debug info> *) 0x2aaaaac010 <mm>
```

The address is the same, the contents of the long double constant are the same, nothing changed. Weird.

Next thing: Change the text segment? I add the same reloc directive just before the `jr ra` at the end of `main`. Now I obtain:

```
(gdb) b main
Breakpoint 1 at 0x66c
(gdb) run
Starting program: /home/jacob/tiny-asm/tld-reloc
/home/jacob/tiny-asm/tld-reloc: error while loading shared libraries:\
unexpected reloc type 0x01
[Inferior 1 (process 1474) exited with code 0177]
```

Great! Now something seems to have changed. I can't run the program. The relocation is probably disturbing something in the program loader.

Conclusion

1) Do not mess around with this unless you know exactly what you are doing...

2) If you know what you are doing... please let me know.

😜 !

1.10.21 text

Tells the assembler to change (if necessary) to the text section. Data and instructions will go at the end of that section.[31]

The change is handled by the `s_text` function:

Listing 1.12: s_text function

```
1 static void s_text(int ignore ATTRIBUTE_UNUSED)
2 {
3     int temp = get_absolute_expression();
4     subseg_set(text_section,(subsegT) temp);
5     demand_empty_rest_of_line();
6 }
```

---

[31] Remember that "text" in this context has *nothing* to do with a text format, in the usual sense of the word. There is no text sequences here, unless you put a text sequence yourself.

The function `subseg_set` is used in several other functions to change the current section/segment.

Listing 1.13: Code of subseg_set

```
1  static void subseg_set(segT secptr,subsegT subseg)
2  {
3      if (!(secptr == now_seg && subseg == now_subseg))
4          subseg_set_rest(secptr,subseg);
5  }
6  static void subseg_set_rest(segT seg,subsegT subseg)
7  {
8      frchainS      *frcP;   /* crawl frchain chain */
9      frchainS      **lastPP; /* address of last pointer */
10     frchainS      *newP;   /* address of new frchain */
11     segment_info_type *seginfo;
12
13     if (frag_now && frchain_now)
14         frchain_now→frch_frag_now = frag_now;                          (1)
15     subseg_change(seg,(int)subseg);                                    (2)
16     seginfo = seg_info(seg);                                           (3)
17     /* Should the section symbol be kept? Yes. */
18     seg→symbol→flags |= BSF_SECTION_SYM_USED;                          (4)
19     /* Attempt to find or make a frchain for that subsection. We keep the
20      * list sorted by subsection number. */
21     for (frcP = *(lastPP = &seginfo→frchainP); frcP ≠ NULL;
22             frcP = *(lastPP = &frcP→frch_next))
23             if (frcP→frch_subseg ≥ subseg)
24                 break;
25     if (frcP == NULL || frcP→frch_subseg ≠ subseg) {
26     /* Not found. Make a new. This should be the only code that creates a frchainS.*/
27         newP = (frchainS *) obstack_alloc(&frchains,sizeof(frchainS));
28         newP→frch_subseg = subseg;
29         newP→fix_root = NULL;
30         newP→fix_tail = NULL;
31         obstack_begin(&newP→frch_obstack,CHUNKSIZE);
32         obstack_alignment_mask(&newP→frch_obstack) = __alignof__(fragS) - 1; (5)
33         newP→frch_frag_now = frag_alloc(&newP→frch_obstack);
34         newP→frch_frag_now→fr_type = rs_fill;
35         newP→frch_cfi_data = NULL;
36         newP→frch_root = newP→frch_last = newP→frch_frag_now;
37         *lastPP = newP; // Insert in chain
38         newP→frch_next = frcP;
39         frcP = newP;
40     }
41     frchain_now = frcP;
42     frag_now = frcP→frch_frag_now;
43  }
```

1. Make sure that `frchain_now` has a correct pointer in `frch_frag_now`.

2. `subseg_change` is a small function that sets the global variables `now_seg` and `now_subseg` to the values given, and, if necessary, allocates the `seg_info` structure.

3. `seg_info` is just a macro that accesses the structure in the `userdata` of the bfd.

4. The original code used a function call and was just too complicated for setting a flag. It was: `if (bfd_keep_unused_section_symbols(stdoutput))` that returned always true...

5. The gnu C construct `__alignof__` has an equivalent in the C standard of 2011: `_Alignof`. In the code above it should be `_Alignof(fragS)`.

1.10.22 uleb128, sleb128

```
Syntax:
.uleb128 value
.sleb128 value
Table: riscv_pseudo_table
{"uleb128",s_riscv_leb128,0},
{"sleb128",s_riscv_leb128,1},
```

These instructions encode a number using a special format. There is also a general directive for all machines that has the same syntax.

To encode an unsigned number:

1. Split the number in 7 bit chunks

2. Read the 7 bits of the lowest significant bits into a byte.

3. Set the most significant bit of the byte to 1 if more bytes follow, to zero otherwise.

4. Output 1 byte and shift the value right by 7 bits.

Listing 1.14: output_uleb128

```
1  static unsigned int output_uleb128(char *p,valueT value)
2  {
3      char          *orig = p;
4      unsigned byte;
5  
6      do {
7          byte = (value & 0x7f);
8          value >≥  7;
9          if (value ≠ 0)
10         /* More bytes to follow. */
11             byte |= 0x80;
12         *p++ = byte; // If value was zero, byte is zero
13     } while (value ≠ 0);
14     return p - orig;
15 }
```

A signed number has a different encoding. Example: Encode -98765432

1. Ignore the minus sign. Binary representation is
   0101 1110 0011 0000 1010 0111 1000, a 27 bit number padded to 28 with zero.

2. Negate all bits, what gives:
   1010 0001 1100 1111 0101 1000 0111

3. Add 1, what gives:
   1010 0001 1100 1111 0101 1000 1000

4. Split into 7 bit groups:
   1010000 1110011 1101011 0001000

5. Add high 1 bit in all but the most significant one
   01010000 11110011 11101011 10001000 → 0x50F3EB88

The code for this is written in a quite complicated way, maybe because the code doesn't do step 1 above or because some machine under some OS is behaving badly...

Listing 1.15: output_sleb128

```
1  static inline unsigned int output_sleb128(char *p,offsetT value)
2  {
3      char            *orig = p;
4      int     more;
5
6      do { unsigned byte = (value & 0x7f);
7      /* Sadly,we cannot rely on typical arithmetic right shift behaviour. Fortunately,
8       * we can structure things so that the extra work reduces to a noop on systems
9       * that do things "properly". */
10         value = (value >> 7)|~(-(offsetT) 1 >> 7);
11         more = !(((((value == 0) && ((byte & 0x40) == 0))
12         || ((value == -1) && ((byte & 0x40) ≠ 0))));
13         if (more) byte |= 0x80;
14         *p++ = byte;
15     } while (more);
16     return p - orig;
17 }
```

1.10.23 Other directives

In general, the code for handling directives is simple and easy to follow. There is no need to detail all of that here.

1.11   The cfi directives

CFI stands for **C**all **F**rame **I**nformation. The objective of these directives is to furnish to a debugger enough information so that at any address within the program, the layout of the stack is clear.

The C++ language uses also this kind of information for another purpose: to rewind the stack, looking for a procedure that will *catch* an exception that has been thrown somewhere in the program. To be able to reconstruct the stack at any moment, big tables are generated, that give the stack unwinding machinery all the information needed to rewind the stack.

Before we get into the details, we need to explain some concepts. We begin with the concept of the *stack frame*, i.e. the portion of the stack used by the currently running function. When a function call is executed, both the riscv CPU and the ARM cpu copy the address of the next instruction into a special register. At the end of the called function, the last instruction that is executed is a jump to the address stored into that register.

Other machines like the x86 family, do not have a link register and the machine pushes the return address into the stack, decreasing the stack by the address size and writing into the new space the return address. Under the riscv/ARM RISC machines we have a link register that allows to avoid (sometimes) to store the return address in memory.

The stack address at the moment of the call is called **C**anonical **F**rame **A**ddress or CFA.

The first thing the called procedure does is to save the permanent registers that it will use. All machines have in their ABI a list of registers that are preserved across calls (the permanent registers) and other scratch registers that are used freely, without any obligation to preserve their contents. A procedure then, needs to store the current values of those registers in the stack to be able to restore them at the end to their previous values.

To be able to reconstruct the data that is active at procedures higher in the stack, the debugger or the stack unwinding machinery must restore the values of the saved registers, so the addresses and register numbers must be stored in the tables for each procedure. Starting

with the current instruction pointer, the debugger restores the values of the previous CFA, virtually returning from a procedure, what allows it to show the values of all the variables of that procedure, and so on.

The debug information is independent of the type of machine being used, what complicates further things.

- Compilers can duplicate the epilogue to avoid executing a jump instruction to a common one.
- Sometimes a procedure uses a frame pointer register, sometimes they use directly the stack pointer.
- Within the prologue or epilogue, the stack can change. Some compilers will use a push instruction for each register saved, some others will subtract from the stack a fixed amount, and save the registers at fixed offsets from the stack or frame pointer.
- Sometimes a preserved register will be saved in a scratch register, and restored later without using a stack frame...
- Some machines use a bit-mask for saving the registers in a single instruction.
- Etc. There are many other special conditions, weird designs that needed not to be mentioned here.

1.11.1 Concepts

The `.eh_frame` section contains two things: a CIE (Common Information Entry) and several FDEs or Frame Description Entry) records.

The CIE

Table 1.11: Common Information Entry fields

Field	Description
Length	A 4 byte unsigned value indicating the length in bytes of the CIE structure, not including the Length field itself. If Length contains the value 0xffffffff, then the length is contained in the Extended Length field. If Length contains the value 0, then this CIE shall be considered a terminator and processing shall end.
Extended Length	Optional, see above. In practice, this is never used.
CIE-ID	A 4 byte value that is used to distinguish between CIEs and FDEs. In CIEs it will be always zero.
Version	This is a single byte and should be 1.
Augmentation	This is a series of byte codes that are interpreted (sounds familiar?) See below.
Code alignment	An unsigned leb128 encoded value that represents the units used in the "advance location" instructions in this CIE and its associated FDEs.

Table 1.11: Common Information Entry fields

return address register	This field is only mentioned in the MaskRay blog. All other official documents do not mention it[32]
Data alignment factor	Similar to the code alignment factor above
Augmentation length	Unsigned leb128 encoded value. This field is only present if the augmentation string contains the 'z' character.
Augmentation data	A block of data, that is interpreted according to the augmentation string.
The augmentation string characters	
'z'	Indicates there is some data there. Must be the first character.
'L'	The FDEs contain pointers to language specific data. This is a single byte that indicates how those pointers are encoded.
'P'	This indicates the presence of two items: 1) A single byte that specifies how the second item, a pointer, is encoded. 2) The second item is encoded according to the type of encoding described by the first, and it represents a pointer to a **personality** routine, i.e. some routine that will be used to unwind the stack according to the language preferences.
'S'	An associated FDE describes a signal frame, i.e. an interrupt procedure [33].

The FDE

FDE stands for Frame Description Entry.

Table 1.12: FDE fields

Field	Description
Length	In 4 bytes
Extended Length	Same specs as in the CIEs above
CIE pointer	A 4 byte unsigned value that when subtracted from the offset of the the CIE Pointer in the current FDE yields the offset of the start of the associated CIE.
Program Counter begin	This is a pointer encoded according to the method specified by the 'R' character in the CIE[34]
PC range	An absolute value that tells how long the code section is.
Augmentation length	Unsigned leb128 encoded value that contains the length of the following data

---

[32]The riscv specification mentions explicitly that other registers could contain the return address.

> There is no dedicated stack pointer or subroutine return address link register in the Base Inte-gerISA; the instruction encoding allows any x register to be used for these purposes. However, the standard software calling convention uses register x1 to hold the return address for a call, with register x5 available as an alternate link register. The standard calling convention uses register x2 as the stack pointer.

RISC-V Unprivileged ISA V20191214-draft, page 14 For RISC-V machines then, this field *could* be useful. In any case, the software representation has a field "return column".

[33]This letter is not mentioned in the Linux Standard Base specifications release 5, but it is mentioned in the MaskRay blog.

[34]... as far as I have understood this mess.

Table 1.12: FDE fields

Augmentation data	Contains pointers encoded according to the prescriptions of the CIE
Call frame instructions	A set of call frame instructions.

Software representation

A CIE will be described by the following structure in asm.h:

```
1 struct cie_entry {
2     struct cie_entry *next;
3     symbolS        *start_address;
4     unsigned   return_column;
5     unsigned   signal_frame;
6     unsigned char fde_encoding;
7     unsigned char per_encoding;
8     unsigned char lsda_encoding;
9     expressionS personality;
10    struct cfi_insn_data *first,*last;
11 };
```

The `cie_entry` structure will be built in the function `cfi_finish` .
An FDE is described by the following structure:

```
1 struct fde_entry {
2     struct fde_entry *next;        Linked list
3     symbolS        *start_address;   start
4     symbolS        *end_address;     end
5     struct cfi_insn_data *data;
6     struct cfi_insn_data **last;
7     unsigned char per_encoding;    Always DW_EH_PE_omit
8     unsigned char lsda_encoding;   Always DW_EH_PE_omit
9     int     personality_id;        Not supported in riscv
10    expressionS personality;
11    expressionS lsda;
12    unsigned   return_column;
13    unsigned   signal_frame;
14    int     eh_header_type;
15    /* Compact unwinding opcodes,not including the PR byte or LSDA. */
16    int     eh_data_size;
17    uint8_t        *eh_data;
18    symbolS        *eh_loc;          Not used in riscv
19    int     sections;
20 };
```

The constructor is `cfi_new_fde`. It receives a label symbol as argument, and the fde will start at that label. Calls `alloc_fde_entry` to allocate and fill the new structure with default values. The default "return column" is 1, as the ABI specifies[35].

1.11.2 An example

Let's see how the debug information is organized with a simple example. Given the following C program:

---

[35]This is a misnomer. It is not a "column" but a register number actually. Columns in the virtual table correspond to register numbers.

```
1  #include <stdio.h>
2  int main(void)
3  {
4      printf("hello\n");
5  }
```

... sorry for this lack of any imagination. Now, if we compile this with:

```
1  star64:~/tiny-asm $ gcc -c -S -g hello.c
```

We obtain then:

Listing 1.16: hello.s

```
1  star64:~/tiny-asm cat hello.s
2      .file  "hello.c"        set the file name
3      .option pic              see §1.10.17 page 57
4      .text                    assemble in the text section
5  .Ltext0:
6      .cfi_sections .debug_frame See §1.11.3 page 67.
7      .file 0 "/home/jacob/tiny-asm" "hello.c"
8      .section   .rodata       assemble in the read only section
9      .align 3                 align to multiple of 8 (2³)
10 .LC0:
11     .string "hello"          see §1.10.2 page 46
12     .text
13     .align 1
14     .globl  main
15     .type   main, @function
16 main:
17 .LFB0:                       "main" will be known as LFB0 in some debug statements
18     .file 1 "hello.c"
19     .loc 1 3 1               See 1.10.15, page 52.
20     .cfi_startproc           first executable instruction of "main"
21     addi    sp,sp,-16        reserve space for stack frame
22     .cfi_def_cfa_offset 16 record that with CFI
23     sd  ra,8(sp)             store return address at sp+8
24     sd  s0,0(sp)             store previous frame pointer at (sp).
25     .cfi_offset 1, -8        return address is at s0-8. See §1.11.6 page 70
26     .cfi_offset 8, -16       previous frame pointer is at s0-16
27     addi    s0,sp,16         set s0 (frame pointer)
28     .cfi_def_cfa 8, 0        See §1.11.8 page 70
29     .loc 1 4 2               Start line 4 col 2 in the C text above
30     lla a0,.LC0              load the address of .LC0 into a0
31     call    puts@plt         call puts (and not printf)
32     li  a5,0                 put zero into scratch register a5
33     .loc 1 5 1               we start line 5 col 1 of the program text
34     mv  a0,a5                put the zero into the result register
35     ld  ra,8(sp)             restore the return address
36     .cfi_restore 1           tell that to CFI
37     ld  s0,0(sp)             restore the frame pointer
38     .cfi_restore 8           tell that to CFI
39     .cfi_def_cfa 2, 16       See §1.11.8 page 70
40     addi    sp,sp,16         restore the stack
41     .cfi_def_cfa_offset 0  tell that to CFI
42     jr  ra                   jump to the return address
43     .cfi_endproc             tell CFI that we returned
44 .LFE0:                       alias for the end of "main"
45     .size   main, .-main     subtract from the current position the address of "main"
46                              label. That will be the size of this procedure.
```

```
47  # Further lines snipped
```

We see here that there are only 7 `.cfi_*` directives used. In bigger files, for instance in `asm.c` we find that the only directives used are exactly the same ones. And that file makes around 35 000 lines. We will document here those ones that are used by gcc. The other are documented in the GAS documentation.

Let's go to each of those `cfi` directives in detail.

1.11.3 cfi_sections

```
Syntax:
.cfi_sections <section_list>
Table: cfi_pseudo_table
 {"cfi_sections",dot_cfi_sections,0},
```

The directive `.cfi_sections` is used to specify the type of format that should be used: whether CFI directives should emit `.eh_frame` section, `.debug_frame` section and/or `.sframe` section. To emit multiple sections, specify them together in a list. For example, to emit both `.eh_frame` and `.debug_frame`, use `.eh_frame, .debug_frame`. The default if this directive is not used is `.cfi_sections .eh_frame`.

The `.eh_frame` is required for exceptions to work. It must contain sufficient info to unwind from all the places where exception may be raised, but doesn't have to include anything beyond that. For example, it does not need to contain info needed to unwind through function prologue or epilogue, since no exception can be raised there.

The `.debug_frame` (and other `.debug_*` sections) is only needed for debugging (and also for "self-aware" programs which unwind their own stack on e.g. crashes). It should contain sufficient info for debugger to unwind the stack from arbitrary place in the program, though in practice it may not.

The differences between the two formats are:[36]

- `.eh_frame` is based on `.debug_frame` introduced in DWARF v2.
- `.eh_frame` has the flag of `SHF_ALLOC` (indicating that a section should be part of the process image) but `.debug_frame` does not, so the latter has very few usage scenarios.
- `.debug_frame` supports DWARF64 format (supports 64-bit offsets but the volume will be slightly larger) but `.eh_frame` does not support (in fact, it can be expanded, but lacks demand)
- In the CIE (Common Information Entry) of `.debug_frame`, augmentation instead of `augmentation_data_length` and `augmentation_data` is used.
- The version field in CIEs is different.
- The meaning of `CIE_pointer` in FDEs is different. `.debug_frame` indicates a section offset (absolute) and `.eh_frame` indicates a relative offset. This change made by `.eh_frame` is great. If the length of `.eh_frame` exceeds 32-bit, `.debug_frame` has to be converted to DWARF64 to represent `CIE_pointer`. Relative offsets do not need to worry about this issue (if the distance between FDE and CIE exceeds 32-bit, add a CIE OK)
- In `.eh_frame`, augmentation typically includes R and the FDE encoding is `DW_EH_PE_pcrel | DW_EH_PE_sdata4` for small code models of AArch64, PowerPC64, x86-64.
- `initial_location` has 4 bytes in GCC (even if `-mcmodel=large`). In `.debug_frame`, 64-bit architectures need 8-byte `initial_location`. Therefore, `.eh_frame` is usually smaller than an equivalent `.debug_frame`

_____

[36]see https://maskray.me/blog/2020-11-08-stack-unwinding

1.11.4   cfi_startproc

.cfi_startproc is used at the beginning of each function that should have an entry in
.eh_frame.

Syntax:
.cfi_startproc [simple]
Table: cfi_pseudo_table
{"cfi_startproc",dot_cfi_startproc,0}

The .cfi_startproc directive is handled by dot_cfi_startproc, that performs following
actions:

- Verifies that an cfi_endproc has been issued or that we are at the start of the program.
- Allocates and initializes a new FDE.
- If present parses the simple argument, and sets an internal flag accordingly.
- If simple wasn't present, it generates the initial instructions for the virtual machine,
  in this case it sets the stack pointer to x [37].

1.11.5   cfi_def_cfa_offset

Syntax:
.cfi_def_cfa_offset offset
Table: cfi_pseudo_table
{"cfi_def_cfa_offset",dot_cfi,DW_CFA_def_cfa_offset}

.cfi_def_cfa_offset modifies a rule for computing CFA. Register remains the same,
but offset is new. Note that it is the absolute offset that will be added to a defined register
to compute CFA address. In the example of hello.s line 22 we see that the new offset is
emitted right after we subtract 16 from the stack. Right after that instruction, the CFA is
16 bytes from the value of sp, obviously.
    This instruction (and several others) are handled by the dot_cfi function that receives
as its argument the instruction for the virtual machine.
    This function does the following:

- Check that a previous cfi_startproc has been issued.
- If the last address wasn't the current address, emit an instruction to advance to the
  current address.
- And now... a big switch statement that will perform the actions needed for each
  instruction.

In this case (DW_CFA_def_cfa_offset) the code is:

```
1    case DW_CFA_def_cfa_offset:
2        offset = cfi_parse_const();
3        cfi_add_CFA_def_cfa_offset(offset);
4        break;
```

The function cfi_add_CFA_def_cfa_offset is as follows:

```
1  /* Add a DW_CFA_def_cfa_offset record to the CFI data. */
2  static void cfi_add_CFA_def_cfa_offset(offsetT offset)
3  {
4      cfi_add_CFA_insn_offset(DW_CFA_def_cfa_offset,offset);
```

---

[37]This instruction is repeated for each procedure in the program. It would be much easier to set this
information in the CIE, since there isn't any program that will switch the stack register on a procedure
basis...

```
5       frchain_now→frch_cfi_data→cur_cfa_offset = offset;
6 }
7
8 static void cfi_add_CFA_insn_offset(int insn,offsetT offset)
9 {
10      struct cfi_insn_data *insn_ptr = alloc_cfi_insn_data();
11
12      insn_ptr→insn = insn;
13      insn_ptr→u.i = offset;
14 }
```

Each action is split in several functions, a side-effect of object oriented design. The function `alloc_cfi_insn_data` allocates space for a new data packet.

These data packets are defined like this:

Listing 1.17: cfi_insn_data

```
1 struct cfi_insn_data {
2     struct cfi_insn_data *next;         Linked list
3     int     insn;                        The instruction in question
4     union {                              Depending on the instruction, only one
5         struct {                          of these fields is active.
6             unsigned   reg;
7             offsetT    offset;
8         }       ri;
9         struct {
10            unsigned   reg1;
11            unsigned   reg2;
12        }       rr;
13        unsigned   r;
14        offsetT    i;
15        struct {
16            symbolS       *lab1;
17            symbolS       *lab2;
18        }       ll;
19        struct cfi_escape_data *esc;
20        struct {
21            unsigned   reg   ,encoding;
22            expressionS exp;
23        }        ea;
24        const char    *sym_name;
25    } u;
26 };
```

The function `alloc_cfi_insn_data` let us see immediately how everything is organized:

```
1 static struct cfi_insn_data *alloc_cfi_insn_data(void)
2 {
3     struct cfi_insn_data *insn = XCNEW(struct cfi_insn_data);
4     struct fde_entry *cur_fde_data = frchain_now→frch_cfi_data→cur_fde_data;
5
6     *cur_fde_data→last = insn;      Link the new item in the linked list
7     cur_fde_data→last = &insn→next;
8     SET_CUR_SEG(insn,is_now_linkonce_segment());
9     return insn;
10 }
```

The macro `XCNEW` is just a call to `xcalloc` with a corresponding `sizeof` its argument, that should be a type. It is saved as the current FDE data pointer, added to the linked list. And that is all.

No? You want me to explain to you the impressing code
SET_CUR_SEG(insn,is_now_linkonce_segment());
Well, I don't know what it should do, since in asm.h we have the definition:
#define␣SET_CUR_SEG(structp,seg)␣(void)␣(0␣&&␣seg) [38]

So, all that complex statement is actually nothing!

1.11.6 cfi_offset

Syntax:
.cfi_offset register, offset
Table: cfi_pseudo_table
{"cfi_offset",dot_cfi,DW_CFA_offset}

The previous value of register is saved at offset *offset* from the CFA. Processing goes to
dot_cfi (see above in cfi_def_cfa_offset). The relevant lines in dot_cfi are:

```
1    case DW_CFA_offset:
2        reg1 = cfi_parse_reg();
3        cfi_parse_separator();
4        offset = cfi_parse_const();
5        cfi_add_CFA_offset(reg1,offset);
6    break;
```

1.11.7 cfi_restore

Syntax:
.cfi_restore register [, register]
Table: cfi_pseudo_table
{"cfi_restore",dot_cfi,DW_CFA_restore}

The argument is a list of one or more registers. Again, we use the workhorse dot_cfi.
The relevant lines are below:

```
1     case DW_CFA_restore:
2     for (;;) {
3         reg1 = cfi_parse_reg();
4         cfi_add_CFA_restore(reg1);
5         SKIP_WHITESPACE();
6         if (*input_line_pointer ≠ ',')
7             break;
8         ++input_line_pointer;
9     }
10    break;
```

1.11.8 cfi_def_cfa

Syntax:
.cfi_def_cfa register, offset
Table: cfi_pseudo_table
{"cfi_def_cfa",dot_cfi,DW_CFA_def_cfa},

.cfi_def_cfa defines a rule for computing CFA as: take address from register and add
offset to it. The relevant lines in dot_cfi are:

---

[38]Yes, I should eliminate all those fake statements from the code of tiny-asm... but I haven't since it is
quite a lot of work, to find them, and to get rid of them. In other CPUs that statement does something,
surely.

```
1    case DW_CFA_def_cfa:
2        reg1 = cfi_parse_reg();
3        cfi_parse_separator();
4        offset = cfi_parse_const();
5        cfi_add_CFA_def_cfa(reg1,offset);
6    break;
```

1.11.9 `.cfi_endproc`

Syntax:
.cfi_endproc
Table: cfi_pseudo_table
{"cfi_endproc",dot_cfi_endproc,0},

`.cfi_endproc` is used at the end of a function where it closes its unwind entry previously opened by `.cfi_startproc` and emits it to `.eh_frame`.

The `dot_cfi_endproc` procedure is as follows:

```
1  static void dot_cfi_endproc(int ignored ATTRIBUTE_UNUSED)
2  {
3      if (!cfi_test_startproc()) return;
4      last_fde = frchain_now→frch_cfi_data→cur_fde_data;
5
6      cfi_end_fde(symbol_temp_new_now());
7      demand_empty_rest_of_line();
8
9      cfi_sections_set = true;
10     if ((cfi_sections & CFI_EMIT_target) ≠ 0)
11         tc_cfi_endproc(last_fde);
12 }
```

- Requires a previous open `startproc`
- sets globals like `last_fde`, a variable that is set, kept current, but *never used*. It is there just for fun.
  Or not?
  Actually, it is used when `SUPPORT_COMPACT_EH` is defined. Since this is not supported under the riscv version of GAS, what you see are just leftovers of its former self... [39]
- `cfi_end_fde` sets several globals to mark the end of a function.
- `tc_cfi_endproc` is `#defined` as nothing, so the last two lines are empty.

1.12 The actual instructions

OK, we know now how to build tiny-asm, how to write directives, how the operations are encoded, let's start now to *do* something with that knowledge. Let's see how the common operations are done.

We will start by showing programs generated by the C compiler. It is the best way to get a feeling for this machine, its instructions and its possibilities.

---

[39]It can be asked why these variables are still there if they do not fill any purpose. There are several reasons why. The first is that they do not cost a lot of space or execution time. And the second is that, of course, maybe tiny-asm will one day support compact `eh_frames`, and if the skeleton of places where the variable is set and updated is erased, that would be impossible. And the third one is that I haven't found the time to enclose all usages of that variable in a conditional compilation like `SUPPORT_COMPACT_EH` what would be actually the correct solution.

1.12.1 Loads, stores and addition

Here we cover the basics: loading data from memory, performing an operation, and storing
the result in memory again. The riscv is a RISC machine, i.e. like the ARM, it can't work
directly on data in memory like the x86 family. Data must be first loaded into memory,
before it can be used for calculations.

Consider the following C program:

```
1 int main(void)
2 {
3     short sa=5,sb=6,sc=sa+sb; // 16 bit addition
4     int ia=5,ib=6,ic=ia+ib; // 32 bit
5     long long lla=5,llb=6,llc=lla+llb; // 64 bit
6     float fa=5,fb=6,fc=fa+fb; // single precision
7     double da=5,db=6,dc=da+db;// double precision
8     return sc+ic+llc+fc+dc; // Should be 55 isn't it?
9 }
```

We translate this with:
`gcc -c -S add.c` obtaining the following assembler file:

Listing 1.18: add.s, no optimizations

```
1     .file  "add.c"        // Standard instructions at the beginning of any file.
2     .option pic           // PC relative code
3     .text                 // Ensure code section
4     .align 4              // Align to multiple of 16 bits (2^4 bytes)
5     .globl main           // Visible outside this module
6     .type  main,@function // Debug statement
7 main:
8     addi   sp,sp,-112     // add immediate -112 to the value in the stack
9     sd s0,104(sp)         // Store doubleword: old frame pointer (s0)
10    addi   s0,sp,112      // Setup the new frame pointer
```

At this point the prologue of this function is finished. The old value of the frame pointer
has been saved and a new one established. We start compiling the first C statement.

```
11                         // short sa=5,sb=6,sc=sa+sb; // 16 bit addition
12    li a5,5              // Put constant 5 in a5
13    sh a5,-18(s0)        // Store halfword (16 bits)
14    li a5,6              // Put 6 into a5
15    sh a5,-20(s0)        // Store it
16    lhu a4,-18(s0)       // Load half word unsigned
```

Note that the compiler uses the "`lhu`" instruction for loading an *unsigned* instead of the
correct one `lh` that does a sign extension and is used for loading signed data, as it should
be since we have declared the data as `short` and not `unsigned short`!

```
17    lhu a5,-20(s0)       // Same as above
18    addw   a5,a4,a5      // At last! 32 bit aaddition
19    slli   a5,a5,48      // shift left a5 48 bits
20    srli   a5,a5,48      // shift right a5 48 bits
21    sh  a5,-22(s0)       // Store 16 bits: store halfword, sh
```

The compiler emits code to load the data as unsigned, do the addition, and select the lower
16 bits. We can see better what is going on if we follow this sequence in the debugger but
using -5 instead of a positive constant.

```
=> 0x2aaaaaa63e <main+22>: lhu a5,-20(s0)
(gdb) print/x $a4
$2 = 0xfffb
```

```
=> 0x2aaaaaa642 <main+26>: addw a5,a5,a4
(gdb) print/x $a5
$3 = 0x6
=> 0x2aaaaaa644 <main+28>: slli a5,a5,0x30
(gdb) print/x $a5
$4 = 0x10001
=> 0x2aaaaaa646 <main+30>: srli a5,a5,0x30
(gdb) print/x $a5
$5 = 0x1000000000000
=> 0x2aaaaaa648 <main+32>: sh a5,-22(s0)
(gdb) print/x $a5
$6 = 0x1
```

We see now that the addition was done in an unsigned form, producing 0x10001, that after the shifts was converted to 1. So, -5 + 6 → 1. We are saved for this time... [40]

```
22                          // int ia=5,ib=6,ic=ia+ib; // 32 bit
23    li  a5,5              // Same as before: 5 into a5
24    sw  a5,-28(s0)        // Initialize "ia" to 5
25    li  a5,6              // Put 6 into a5
26    sw  a5,-32(s0)        // Store it into "ib"
27    lw  a5,-28(s0)        // load ia
28    mv  a4,a5             // copy it to a4
29    lw  a5,-32(s0)        // load "ib"
30    addw   a5,a4,a5       // Do the addition
31    sw  a5,-36(s0)        // store the result
32                          // long long lla=5,llb=6,llc=lla+llb; // 64 bit
33    li  a5,5              // load 5
34    sd  a5,-48(s0)        // Store doubleword this time
35    li  a5,6              //
36    sd  a5,-56(s0)        // Same
37    ld  a4,-64(s0)        // Load "lla" into a4 (directly this time)
38    ld  a5,-48(s0)        // Load "llb" into a5
39    add a5,a4,a5          // 64 bit addition
40    sd  a5,-64(s0)        // Store 64 bits
41                          // float fa=5,fb=6,fc=fa+fb; // single precision
42    lla a5,.LC0           // Load the address of.LC0 into a5
43    flw fa5,0(a5)         // Load single precision from the address in a5
44    fsw fa5,-68(s0)       // Store it at "fa"
45    lla a5,.LC1           // Same process for "fb".
46    flw fa5,0(a5)
47    fsw fa5,-72(s0)       // "fb" at -72
48    flw fa4,-68(s0)       // Load fa4 with "fa"
49    flw fa5,-72(s0)       // Load fa5 with "fb"
50    fadd.s fa5,fa4,fa5    // Add single precision
51    fsw fa5,-76(s0)       // Store result at -76
52                          // double da=5,db=6,dc=da+db; double precision
53    lla a5,.LC2           // Load the address of .LC2 into a5
54    fld fa5,0(a5)         // Load double precision from that address
55    fsd fa5,-88(s0)       // Initialize "da"
56    lla a5,.LC3           // Same for "db"
57    fld fa5,0(a5)
```

---

[40]Why does the compiler do this instead of loading everything as signed and doing a signed addition? Nobody knows, at least not me. Note that we are using the compiler without any optimizations, we will see later what happens when some of those are turned on.

In any case, the sequence of loading sign extended 16 bit data and making a 32 bit addition gives exactly the same results.

```
58      fsd fa5,-96(s0)
59      fld fa4,-88(s0)       // Load "da" into fa4
60      fld fa5,-96(s0)       // Load "db" at fa5
61      fadd.d fa5,fa4,fa5    // Add double precision
62      fsd fa5,-104(s0)      // Store result
63                            // return sc+ic+llc+fc+dc; Should be 55
64      lh  a5,-22(s0)        // "sc" into a5
65      sext.w a5,a5          // Sign extend it
66      lw a4,-36(s0)         // "ic" goes into a4
67      addw   a5,a4,a5       // Add both a,nd accumulate into a5
68      sext.w a5,a5          // Sign extend result to 64 bits
69      mv a4,a5              // Copy it to a4
70      ld a5,-64(s0)         // Load "llc" to a5
71      add a5,a4,a5          // Add accumulating into a5
72      fcvt.s.l   fa4,a5     // Convert integer in a5 into float in fa4
73      flw fa5,-76(s0)       // Load "fc" into fa5
74      fadd.s fa5,fa4,fa5    // Add single precision fa4 and fa5
75      fcvt.d.s   fa4,fa5    // Convert that result into double precision
76      fld fa5,-104(s0)      // Load "dc" into fa5
77      fadd.d fa5,fa4,fa5    // Add double precision fa4+fa5 -→ fa5
78      fcvt.w.d a5,fa5,rtz   // Truncate result into a45
79      sext.w a5,a5          // Sign extend
80      mv a0,a5              // Put result into the result register
81                            // Start of epilogue --------------------
82      ld s0,104(sp)         // Restore previous frame pointer
83      addi   sp,sp,112      // Restore stacl
84      jr ra                 // Jump to return address
85                            // End of code of "main"
86      .size  main, .-main   // Compute size of main at assembly time
87      .section   .rodata    // New section: read only data
88      .align 2              // Align to 4 byte boundary
89  .LC0:
90      .word  1084227584     // 5.0 in single precision
91      .align 2
92  .LC1:
93      .word  1086324736     // 6.0 in single precision
94      .align 3              // Align to 8 byte boundary
95  .LC2:
96      .word  0
97      .word  1075052544     // 5.0 in double precision
98      .align 3
99  .LC3:                     // 6.0 in double precision
100     .word  0
101     .word  1075314688
102                           // End of module add.o GNU specific stuff follows
103     .ident "GCC: (GNU) 11.3.0"
104     .section   .note.GNU-stack,"",@progbits
```

This simple program allows us to see the instructions in action. How data is loaded from, and written to memory, how to convert from integer to floating point and vice versa, and how to add in several formats.

- Load from memory all integer data into the a5 register
- Once in memory, copy the data to its eventual destination.
- Target the result of the operations into a5, to save it into memory.

What happens with higher optimization levels? Trying with
gcc -c -S -O1 add.c we obtain:

```
1 main:
2     li  a0,55
3     ret
```

WOW... there is nothing left even at the lowest optimization level. To avoid this we change the program like this:

```
1 cat add1.c
2 int main(int argc,char *argv[])
3 {
4     short sa=argc,sb=6,sc=sa+sb;
5     int ia=argc,ib=6,ic=ia+ib;
6     long long lla=argc,llb=6,llc=lla+llb;
7     float fa=argc,fb=6,fc=fa+fb;
8     double da=argc,db=6,dc=da+db;
9     return sc+ic+llc+fc+dc;
10 }
```

The compiler can't possibly know what "argc" will contain and will be forced to do the hard work.

This yields the following program:

Listing 1.19: add1.s

```
1 main:                     // argc is in a0 (first argument)
2     addiw   a5,a0,6       // add argc + 6. Result in a5
3     slliw   a5,a5,16      // 16 bit left shift of result
4     sraiw   a5,a5,16      // 16 bit right shift of result. "sa" is in a5
5     addiw   a4,a0,6       // 32 bit add of argc and 6
6     addw    a5,a5,a4      // Accumulate addition into a5
7     addi    a4,a0,6       // Add 64 bits argc + 6 into a4
8     add a5,a5,a4          // Accumulate into a5
9     fcvt.s.l   fa5,a5     // Convert sum sc+ic+llc to double
10    fcvt.s.w   fa4,a0     // Convert argc into float in fa4
11    flw fa3,.LC0,a5       // auipc instruction
```

This instruction, that the assembly code of gcc represents as "lw" is actually the "auipc" instruction that was introduced to the specifications in 2014, version 2.0. "auipc" adds a 20 bit upper immediate to the program counter to form an address where the data will be loaded. This constant will be filled by the linker, that can establish the definitive distance between the program counter and the variable in question[41].

If you look at the entry of "auipc" in the opcodes table you will find:

```
{"auipc",0,INSN_CLASS_I,"d,u",MATCH_AUIPC,MASK_AUIPC,match_opcode,0},
```

Now, looking at table §1.7 page 39 you will see that the 'u' letter means a 20 bit immediate will be supplied. Our label ".LC1" is precisely that.

But, I hear your question, how come that I see "lw" in the assembler source text and an "auipc" instruction gets written out ???

Well, that the magic of tiny-asm. It will be explained below, after we finish with this small program.

```
12    fadd.s fa4,fa4,fa3    // Add single precision: fa4 = fa4 + fa3
13                          // fa4 contains argc in single precision
14                          // fa3 contains 6
15    fadd.s fa5,fa5,fa4    // Accumulate in fa5 that contains the sum of sc+ic
16    fcvt.d.s   fa5,fa5    // Convert from single precision to double precision.
17    fcvt.d.w   fa4,a0     // Convert argc to double precision
```

---

[41]**A**dd **U**pper **I**mmediate to **Program C**ounter → **AUIPC**.

```
18      fld fa3,.LC1,a5        // The same auipc instruction to acces 6.0 in double prec.
19      fadd.d fa4,fa4,fa3     // Double precision add: fa4 = argc+6.0
20      fadd.d fa5,fa5,fa4     // Add toaccumulator fa5
21      fcvt.w.d a0,fa5,rtz    // Convert to integer
22      sext.w a0,a0           // sign extend
23      ret                   // Done.
24      .size  main, .-main
25      .section   .rodata.cst4,"aM",@progbits,4
26      .align 2
27  .LC0:
28      .word  1086324736
29      .section   .rodata.cst8,"aM",@progbits,8
30      .align 3
31  .LC1:
32      .word  0
33      .word  1075314688
```

We see here what it means to optimize:

- The compiler keeps all data in registers, there isn't even a stack frame.
- More operations do actual calculations than loading or storing data from/to memory. In the unoptimized version of add.c we have only 15 out of 76 instructions that do arithmetic. In the optimized version we have 15 out of 33, mainly because there are so few loads and no stores
- Use of more advanced instructions

Load and store instructions in short

Table 1.13: Standard load and store operations

Loads	Description		Stores	Description
lb	Load 8 bits. Sign extension.		sb	Store 8 bits
lbu	Load 8 bits. Zero extension			
lh	Load 16 bits Sign extension		sh	Store 16 bits
lhu	Load 16 bits. Zero extension			
lw	Load 32 bits Sign extension		sw	Store 32 bits
lwu	Load 32 bits Zero extension			
ld	64 bit load		sd	Store 64 bits
lui rd,imm20	load upper immediate: Load a 20 bit address constant into rd.			
auipc rd,imm20	Adds the 20 bit immediate to the program counter and stores the result in rd.			

Syntax
```
load  rd , imm12(rs1)
store rs1, imm12(rs2)
```

The words load and store stands for one of the first seven instructions above. The imm12 is always sign extended.

Addressing modes

- Absolute addressing.

```
        lui  a0, %hi(message)
        addi a0, %lo(message)
```

The `%hi` and the `%lo` constructs mean the higher 20 and the lower 12 bits of the address.

- Relative addressing

```
        auipc  a0, %pcrel_hi(msg + 1)
        addi   a0, a0, %pcrel_lo(message)
```

- GOT (Global Object Table) relative addressing

```
.L1:
        auipc  a0, %got_pcrel_hi(message)
        ld     a0, %pcrel_lo(.L1)(a0)
```

Note that the last two are the same: either PC relative or GOT relative, the instructions

Recognizing addressing modes

The assembler recognizes these keywords using tables of the following structure:

```
1 struct percent_op_match {
2     const char *str; // Name without the percentage sign
3     bfd_reloc_code_real_type reloc; // Relocation type invoked
4 };
5 const struct percent_op_match percent_op_utype[];
6 const struct percent_op_match percent_op_itype[];
7 const struct percent_op_match percent_op_stype[];
8 const struct percent_op_match percent_op_rtype[];
9 const struct percent_op_match percent_op_null[];
```

These tables will be used in the function `parse_relocation` to recognize (or not) a relocation directive.

```
1 /* Return true if *STR points to a relocation operator. When returning true, move
2  * *STR over the operator and store its relocation code in *RELOC. Leave both *STR
3  * and *RELOC alone when returning false. */
4 bool parse_relocation(char **str,bfd_reloc_code_real_type * reloc,
5                       const struct percent_op_match *percent_op)
```

This function will set up a pointer to the first table, and will scan each name in all the tables, assuming they are in consecutive order. The last "table" is a terminator with only zeroes. It is crucial then, that, unaware of this, you insert something in between those tables. That would totally screw up things...

    `parse_relocation` will be called when parsing an expression that should yield a small immediate constant ot offset. Its single use will be in `my_getSmallExpression`.

```
1 /* Parse string STR as a 16-bit relocatable operand. Store the expression in
2  * *EP and the relocation,if any,in RELOC. Return the number of relocation
3  * operators used (0 or 1).
4  *
5  * On exit,EXPR_PARSE_END points to the first character after the expression. */
6 size_t my_getSmallExpression(expressionS * ep, bfd_reloc_code_real_type * reloc,
7                              char *str,const struct percent_op_match *percent_op)
```

Now, this function, `my_getSmallExpression` will be called from two places:

1. `my_getOpcodeExpression`, a function used in `riscv_ip`.

2. `riscv_ip` directly, and in an extensive fashion.
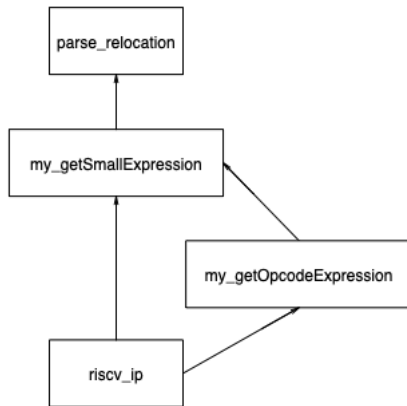
Figure 1.15: Who calls the `parse_relocation`hn,vc=:m:= function

1.12.2   Digression: assembler macros

We have seen above that the expression `flw fa3,.LC0,a5` gets translated into two instructions:

```
1     auipc a5,0x0
2     flw fa3,52(a5)
```

Looking at the opcode table, we find that there several entries for the "flw" instruction.

```
1 {"lw",0,INSN_CLASS_C,"d,Cm(Cc)",MATCH_C_LWSP,MASK_C_LWSP, \
2             match_rd_nonzero,INSN_ALIAS|INSN_DREF|INSN_4_BYTE},
3 {"lw",0,INSN_CLASS_C,"Ct,Ck(Cs)",MATCH_C_LW,MASK_C_LW,\
4             match_opcode,INSN_ALIAS|INSN_DREF|INSN_4_BYTE},
5 {"lw",0,INSN_CLASS_I,"d,o(s)",MATCH_LW,MASK_LW, \
6             match_opcode,INSN_DREF|INSN_4_BYTE},
7 {"lw",0,INSN_CLASS_I,"d,A",0,(int)M_LW,match_never,INSN_MACRO},
```

At line 7, we find an instruction whose flag field has the `INSN_MACRO` set. In the `md_assemble` function, we find the sequence:

```
    if (insn.insn_mo→pinfo == INSN_MACRO)
        macro(&insn,&imm_expr,&imm_reloc);
    else
        append_insn(&insn,&imm_expr,imm_reloc);
```

If this instruction is actually a macro, expand it, if not, append the new instruction.
    What does the `macro` procedure do?

- It decomposes its arguments into 4 parts: the destination register (rd), the two source registers (rs1 and rs2), and a mask. According to the mask, different actions are performed. In our case we have `M_LW`, as wecan see in line 7 of the opcodes listing above. In the same line we find that the function for matching the opcode is `match_never` a function that will always fail, excluding that the macro will be understood as another opcode. [42]
- Using the mask value, it dispatches in a long `switch` statement for each mask. In our case:

```
    case M_LW:
        pcrel_load(rd,rd,imm_expr,"lw",
```

---
[42]`M_LW` is a member of an anonymous enumeration defined in asm.h

```
                         BFD_RELOC_RISCV_PCREL_HI20, BFD_RELOC_RISCV_PCREL_LO12_I);
             break;
```

`pcrel_load` and its companion `pcrel_store` call `pcrel_access` with slightly different arguments:

Listing 1.20: pcrel load and store

```
1 void pcrel_load(int destreg,int tempreg,expressionS * ep,const char *lo_insn,
2         bfd_reloc_code_real_type hi_reloc, bfd_reloc_code_real_type lo_reloc)
3 {
4     pcrel_access(destreg,tempreg,ep,lo_insn,"d,s,j",hi_reloc,lo_reloc);
5 }
6
7 void pcrel_store(int srcreg,int tempreg,expressionS * ep,const char *lo_insn,
8          bfd_reloc_code_real_type hi_reloc,bfd_reloc_code_real_type lo_reloc)
9 {
10     pcrel_access(srcreg,tempreg,ep,lo_insn,"t,s,q",hi_reloc,lo_reloc);
11 }
```

And, to make a disgression within a disgression, long and explicit type names can be nice, but sometimes they can lead to *really* verbose code... What if we substitute in the code above the long names with something like `Reloc` ?

Listing 1.21: pcrel load and store improved

```
1     void pcrel_load(int destreg,int tempreg,expressionS * ep,const char *lo_insn,
2               Reloc hi_reloc, Reloc lo_reloc)
```

Is this code less lisible?

Anyway, both functions call `pcrel_acces` [43]

```
1 static void pcrel_access(int destreg,int tempreg,expressionS * ep,
2                  const char *lo_insn,const char *lo_pattern,
3            bfd_reloc_code_real_type hi_reloc,bfd_reloc_code_real_type lo_reloc)
4 {
5     expressionS ep2;
6     ep2.X_op = O_symbol; // expression is a symbolic expression
7     ep2.X_add_symbol = make_internal_label(); // Symbol to attach the relocation
8     ep2.X_add_number = 0;
9     macro_build(ep,"auipc","d,u",tempreg,hi_reloc); // First insn
10     macro_build(&ep2,lo_insn,lo_pattern,destreg,tempreg,lo_reloc); // Second
11 }
```

`pcrel_access` builds a symbolic expression and calls `macro_build` twice. The first one to build the `auipc` instruction, and the second for the actual load using the temporary register. The function `macro_build` receives as arguments:

1. An expression.

2. A name for the instruction to generate.

3. A format string that will be used, in a similar manner to printf, as a template for the extraction of the corresponding arguments from the rest.

---

[43] The problem with `bfd_reloc_code_real_type` (besides the fact that is a pain to type!) is that many of the words used do not convey any new information... Real type? Are other types "unreal"? What did they want to say?

In our case we give it first the `.LC0` label, the name of the first instruction that we want to generate ("auipc"), and a format string of 'd' and 'u'.

The meaning of those letters is as follows:

Table 1.14: Macro letter arguments

Letter	action
'V'	Vector macro. It needs a further letter for fully specifying which action is needed.
'd'	`INSERT_OPERAND(RD,insn,va_arg(args,int)); continue;`
's'	`INSERT_OPERAND(RS1,insn,va_arg(args,int)); continue`
't'	`INSERT_OPERAND(RS2,insn,va_arg(args,int)); continue`
'q','u' and 'j'	`r=va_args(args,int); continue;` "r" is the relocation type".

Then, just before exiting, `macro_build` will call: `append_insn(&insn,ep,r);`

Let's see the output of `objdump` when we ask for disassembly and relocations:

```
$ # First we generate the object file
$ ./asm -o tauipc.o tauipc.s
$ # Let's look at it with objdump
$ objdump -d -r tauipc.o
0000000000000000 <main>:
0:  00000697            auipc   a3,0x0
0: R_RISCV_PCREL_HI20 .LC0
0: R_RISCV_RELAX  *ABS*
4:  0006a787            flw fa5,0(a3) # 0 <main>
4: R_RISCV_PCREL_LO12_I .L0
4: R_RISCV_RELAX  *ABS*
```

As expected, we have a relocation of 20 bits and another one for the next instruction for the lower 12 bits. There are also 'relax"' relocations, that we will meet later, when we study relocations.

1.12.3 Subtraction

Replacing all additions with subtractions in our C source doesn't change much to the overall shape of the program. The subtraction instructions are:

- `sub`. 64 bit subtraction.
  Syntax: `sub rd,rs1,rs2`
  Operation: `rd ← rs1 - rs2`.

- The `subw` instruction does a 32 bit subtraction. Same syntax and operation as above.

- The `fsub.s` does a single precision subtraction
  Syntax: `fsub.s fd,fs1,fs2`
  Operation: `fd ← fs1 - fs2`

- The `fsub.d` instruction does a double precision subtraction. Same as above.

1.12.4 Comparisons

- `slti`. Set less than immediate. (Signed)
  Syntax: `slti rd,rs1,immediate`
  Operation: `rd ← (rs1 < immediate) ? 1 : 0`

- `sltiu` Set less than immediate unsigned.
  Syntax:
  `sltiu rd, rs1, imm`
          `sltu rd,rs1,rs2`
  Operation: rd ← (rs1 < rs2/imm) ? 1 : 0

- The pseudo instruction `SEQZ rd,rs` sets rd to 1 if rs is equal to zero. This is actually an alias for `SLTIU rd,rs,1`.

- `flt.s` and `flt` perform floating point comparisons for single and double precision floating point respectively.
  Syntax: `flt rd,fsrc1,fsrc2`
   Operation: rd ← fsrc1 < fsrc2) ? 1 : 0
  `rd` is an integer register, `fsrc1` and `fsrc2` are floating point.

- `feq` and `feq.s` do an equality comparison.
  Syntax: `feq rd,fsrc1,fsrc2`
   Operation: rd ← (fsrc1 == fsrc2) ? 1 : 0
  `rd` must be an integer register, `fsrc1` and `fsrc2` are floating point.

An instruction alias that uses subtraction is `neg` that is actually just `sub rd,x0,rs1` i.e. subtract rs1 from zero.

1.12.5 Multiplication and Division

Multiplication

These instructions are present if the processor implements the 'M' extension.

- `mul` performs a 64 by 64 bits multiplication, returning the lower 64 bits.

- `mulh` performs a signed 64 bit by a signed 64 bit multiplication and returns the higher 64 bits of the result.

- `mulhu` multiplies unsigned by unsigned 64 bit quantities and returns the upper 64 bits.

- `mulhsu` multiplies a signed rs1 by an unsigned rs2 and returns the higher 64 bits. [44]

- `mulw` is a 32 bit multiplication. The lower 32 bits are returned, with sign extension.

XuanTie-OpenC910

This processor features several new instructions for multiplication.

Table 1.15: Thead Multiplication extensions

Instruction	Operation	Description
`th.mula` `rd,rs1,rs2`	$rd \leftarrow rd + (rs1 \times rs2)$	Accumulate in `rd`
`th.mulah` `rd,rs1,rs2`	$t[0:31] \leftarrow rd + (rs1[0:15] \times rs2[0:15])$ $rd \leftarrow sign\_extend(t)$	Accumulate with result of 16 bit multiplication.

---

[44]In a multiple precision context, this instruction can be used to multiply the higher 64 bits that contain the sign, with the lower 64 bits of the other multiplicand, that has no sign.

Table 1.15: Thead Multiplication extensions

th.mulaw rd,rs1,rs2	$t[0 : 31] \leftarrow rd + (rs1[0 : 31] \times rs2[0:31])$ $rd \leftarrow sign\_extend(t)$	Accumulate with result of 32 bit multiplication.
th.muls rd,rs1,rs2	$rd \leftarrow rd - (rs1 \times rs2)$	Subtract from rd the result of the multiplication
th.mulsh rd,rs1,rs2	$t[0 : 31] \leftarrow rd - (rs1[0 : 15] \times rs2[0:15])$ $rd \leftarrow sign\_extend(t)$	Subtract from rd result of 16 bit multiplication.
th.mulsw rd,rs1,rs2	$t[0 : 31] \leftarrow rd - (rs1[0 : 31] \times rs2[0:31])$ $rd \leftarrow sign\_extend(t)$	Subtract from rd result of 32 bit multiplication.

These operations are encoded using a modified form of the "R" format. Here is the encoding for the `mula` instruction for instance: `mula a0,a1,a2`.

Figure 1.16: Modified C910 **R** Instruction layout

As you can see, the `extra-7` field of the "R" format has been split into a 5+2 bit field. The meaning of those 2 bits is described in §1.12.14 page 88

Division

A change of the standard allows now to implement processors that have multiplication but not division. For those that do feature division, we have:

- `div` features a signed 64 bit division with rounding towards zero.
  Syntax:
   `div rd,rs1,rs2`
  Operation:
  `rd ← rs1 / rs2`

- `divu` Unsigned division. Syntax and mode of operation the same as DIV.

- `rem` Signed remainder
  Syntax:
   `rem rd,rs1,rs2`
  Operation:
  `rd ← rs1 % rs2`

- `remu` Unsigned remainder. Same as `REM` but for unsigned data.

- `remw` and `remuw` 32 bit versions.

Division by zero returns a result with all bits set, without any trap. [45]

The riscv ISA doesn't provide an instruction for calculating the remainder and the division with only one division operation. The sequence: `DIV[U] rdq, rs1, rs2; REM[U] rdr, rs1, rs2` ()where rdq can't be the same as rs1 or rs2) is proposed for optimization.

1.12.6 Shifts

Table 1.16: Standard shift operations

Syntax	Operation
`slli rd,rsrc1,imm`	Shift left logical immediate.
`srli rd,rsrc1,imm`	Shift right logical immediate (Shifts in zeros)
`srai rd,rsrc1,imm`	Shift right arithmetic (propagating the sign bit).
`sll rd,rsrc1,rsrc2`	Shift left logical (shifts in zeroes).
`sllw rd,rsrc1,rsrc2`	As `sll` but works on lower 32 bits.
`srl rd,rsrc1,rsrc2`	Shift right logical (shifts in zeroes).
`srlw rd,rsrc1,rsrc2`	As `srl` but works on lower 32 bits.
`sra rd,rsrc1,rsrc2`	Shift right arithmetic (propagating the sign bit).
`sraw rd, rsrc1, rsrc2`	As `sra` but works on lower 32 bits.

In all this instructions `rsrc1` is the quantity to be shifted, and `rsrc2` or `imm` contain the number of bits to shift.

1.12.7 Control flow

Incondional Jumps

Table 1.17: Standard incondional jumps

Pseudo instruction	Base instruction	Operation
`j label`	`jal x0 label`	Jump incondicional pc← pc+sign_extend(imm20 * 2)
`jal fn`	`jal x1,fn`	Call subroutine
`jr register`	`jalr x0,register`	Call function pointer in register

The `jal` instructions uses the 'j' instruction format (See §1.6.6 page 30). The offset immediate (in multiples of 2 bytes) is added to the current program counter value to form the target address. It has a reach of 1MB forward or backwards.

---

[45]The riscv standard justifies this with:

> We considered raising exceptions on integer divide by zero, with these exceptions causing a trap in most execution environments. However, this would be the only arithmetic trap in the standard ISA (floating-point exceptions set flags and write default values, but do not cause traps) and would require language implementors to interact with the execution environment's trap handlers for this case. Further, where language standards mandate that a divide-by-zero exception must cause an immediate control flow change, only a single branch instruction needs to be added to each divide operation, and this branch instruction can be inserted after the divide and should normally be very predictably not taken, adding little runtime overhead.The value of all bits set is returned for both unsigned and signed divide by zero to simplify the divider circuitry.

> The value of all 1s is both the natural value to return for unsigned divide, representing the largest unsigned number, and also the natural result for simple unsigned divider implementations. Signed division is often implemented using an unsigned division circuit and specifying the same overflow result simplifies the hardware.

The indirect jumps through a register are **not** in multiples of two bytes, beware. The address must be the real address of the target.

1.12.8 Conditional expressions

Table 1.18: Standard conditional expressions

Inst	Operation
`beq rs1, rs2, label`	if (rs1 = rs2) pc←pc+sign_extend(imm12 << 1)
`bge rs1, rs2, label`	if (rs1 ≥ rs2) pc←pc+sign_extend(imm12 << 1)
`bgeu rs1, rs2, label`	if (rs1 ≥ rs2) pc←pc+sign_extend(imm12 << 1)
`blt rs1, rs2, label`	if (rs1 ≤ rs2) pc←pc+sign_extend(imm12 << 1)
`bltu rs1, rs2, label`	if (rs1 ≤ rs2) pc←pc+sign_extend(imm12 << 1)
`bne rs1, rs2, label`	if (rs1 ≠ rs2) pc←pc+sign_extend(imm12 << 1)

All these instructions have a range of $\pm 4K$.

Exercise 1: *The instruction* `bgt` *is an alias. How would you build it from the other instructions?*

Exercise 2: *Write a small program that uses a conditional branch.*

Exercise 3: *Disassemble the program. What you see instead of bgt?*

Exercise 4: *How is the change achieved? Look at the source asm.c.*

1.12.9 And, Or, Xor

Table 1.19: Standard boolean instructions

Inst	Operation
`and rd,rsrc1,rsrc2`	$rd \leftarrow rsrc1 \wedge rsrc2$
`andi rd,rsrc1,imm12`	$rd \leftarrow rsrc1 \wedge imm12$
`or rd,rsrc1,rsrc2`	$rd \leftarrow rsrc1 \vee rsrc2$
`ori rd,rsrc1,imm12`	$rd \leftarrow rsrc1 \vee imm12$
`xor rd,rsrc1,rsrc2`	$rd \leftarrow rsrc1 \oplus rsrc2$
`xori rd,rsrc1,imm12`	$rd \leftarrow rsrc1 \oplus imm12$

Exercise 5: *Use the XOR instruction to invert all bits in an integer register*

1.12.10 Reading timers

The "Zinctr" extension prescribes at least **3** counters/timers that should be present in all implementations.

- Cycles.The `rdcycle` pseudo instruction reads the low XLEN bits of the cycle special register which holds the number of clock cycles executed by the processor core on which the hardware thread is running from an arbitrary start time somewhere in the past, probably, when the machine was powered on.
- Time. The `rdtime` instructions returns the wall clock time since start, sometime in the past.
- Instructions retired. The `rdinstret` instruction returns the number of instructions retired, i.e. executed (roughly) since some time in the past.

Reading standard counters

Table 1.20: Counter reading

Instruction	Description
`rdtime rd`	Reads a 64 bit timer counter
`rdcycle rd`	Reads a 64 bit cycle counter
`rdinstret rd`	Reads a 64 bit counter for the number of instructions retired, i.e. executed

The `rd` placeholder represents a 64 bit register.

Exercise 6: *Write a program in assembler to print these 3 counters.*

Exercise 7: *Try to verify that time corresponds to a time measure*

1.12.11 CSR instructions

"CSR" stands for **C**ontrol and **S**tatus **R**egister. These registers are used primarily in the privileged part of the instruction set, but there are some uses in the unprivileged instructions (the subject of this book).

1.12.12 Boolean instructions

These instructions correspond to the "Zbb". In the opcode table they have `INISN_CLASS_ZBB` in the class field. extension.

The instructions that work only in a 32 bit environment have been excluded.

The Sifive U74-MC supports the standard Zbb extension. The XuanTie-OpenC910 has two somehow similar instructions, "`ff1`" and "`ff0`".

Note: GCC doesn't recognize these instructions in machines using the U74 CPU. You have to force it by adding the (undocumented) option `-march=rv64gc_zbb` to the compilation command line. These problems do not affect tiny-asm: it will generate the correct instructions without problems.

Table 1.21: Zbb boolean extension instructions

Inst.	Description
`bclr rd,rs1,rs2`	Returns rs1 with a single bit cleared at the index specified in rs2. The index is read from the lower log2(XLEN) bits of rs2.
`bclri rd,rs1,nr`	Returns rs1 with a single bit cleared at the index specified in nr. The index is read from the lower log2(XLEN) bits of nr.
`bext rd, rs1, rs2`	Returns a single bit extracted from rs1 at the index specified in rs2. The index is read from the lower log2(XLEN) bits of rs2.
`bexti rd, rs1, nr`	Returns a single bit extracted from rs1 at the index specified in nr. The index is read from the lower log2(XLEN) bits of nr.
`binv rd, rs1, rs2`	Returns rs1 with a single bit inverted at the index specified in rs2
`binvi rd, rs1,nr`	Returns rs1 with a single bit inverted at the index specified in nr.
`bset rd, rs1,rs2`	Returns rs1 with a single bit set at the index specified in rs2.
`bseti rd, rs1,nr`	Returns rs1 with a single bit set at the index specified in nr.
`clmul rd,rs1,rs2`	clmul produces the lower half of the 2×XLEN carry-less product.
`clmulh rd, rs1, rs2`	Produces the upper half of the 2× XLEN carry-less product.

Table 1.21: Zbb boolean extension instructions

`clmulr`	Produces bits $2\times$ XLEN-2 to XLEN-1 of the $2\times$XLEN carry-less product. [46]
`clz rd,rs1`	Counts the number of 0 bits before the first 1 bit, starting at the most significant bit and progressing to bit 0. If the input is 0, the output is 64. If the most-significant bit of the input is 1, the output is 0.
`clzw rd,rs1`	Counts the number of 0 bits before the first 1 bit, starting at bit 31 and progressing to bit 0. If the least-significant word is 0, the output is 32. If the most-significant bit of the word is 1, the output is 0.
`ctz rd,rs1`	Counts the number of 0 bits before the first 1 bit, starting at the least-significant bit and progressing to the most-significant bit. If the input is 0, the output is 64. If the least-significant bit of the input is 1, the output is 0.)
`ctzw rd,rs1`	Counts the number of 0 bits before the first 1 bit, starting at the least-significant bit and progressing to the most-significant word. If the least significant word is 0, the output is 32. If the least significant bit of the input is 1, the output is 0.
`cpop rd,rs1`	Counts the number of 1 bits in the source register. This operations is also known as "population count" or "Hamming weight".
`cpopw rd,rs1`	Counts the number of 1 bits in the least-significant word of the source register
`max rd,rs1,rs2`	Returns the larger of two signed integers.
`maxu rd,rs1,rs2`	Returns the larger of two unsigned integers.
`min rd,rs1,rs2`	Returns the smaller of two signed integers.
`minu rd,rs1,rs2`	Returns the smaller of two unsigned integers.
`orc.b rd,rs`	Combines the bits within each byte using bitwise logical OR. This sets the bits of each byte in the result rd to all zeros if no bit within the respective byte of rs is set, or to all ones if any bit within the respective byte of rs is set.
`orn rd,rs1,rs2`	performs the bitwise logical OR operation between rs1 and the bitwise inversion of rs2. $rd \leftarrow rs1 \mid \sim rs2$
`rev8 rd,rs1`	Reverses the order of the bytes in rs1.
`rol rd,rs1,rs2`	Rotate left. Performs a rotate left of rs1 by the amount in least-significant 6 bits of rs2.
`rolw rd,rs1,rs2`	Rotate left. Performs a rotate left of rs1 by the amount in least-significant 5 bits of rs2. The resulting 32 bit vvalue is sign extended to 64.
`ror rd,rs1,rs2`	Rotate right. Uses the least significant 6 bits of rs2 for the amount to rotate.
`rori rd,rs1,imm`	Rotate right immediate. Uses the least significant 6 bits of *imm* for the amount to rotate.
`sext.b rd,rs`	Sign-extends the least-significant byte in the source to 64 by copying the most-significant bit in the byte (i.e., bit 7) to all of the more-significant bits.

---

[46]The clmulr instruction is used to accelerate CRC calculations. The r in the instruction's mnemonic stands for reversed, as the instruction is equivalent to bit-reversing the inputs, performing a clmul, then bit-reversing the output.

Table 1.21: Zbb boolean extension instructions

`sext.h rd,rs`	Sign-extends the least-significant 16 bits in the source to 64 by copying the most-significant bit in the byte (i.e., bit 7) to all of the more-significant bits.
`sh1add rd,rs1,rs2`	Shifts rs1 left by 1 and adds it to rs2.
`sh2add rd,rs1,rs2`	Shifts rs1 left by 2 and adds it to rs2.
`sh1add_uw rd,rs1,rs2`	This instruction performs an 64 bit wide addition of two addends. The first addend is rs2. The second addend is the unsigned value formed by extracting the least-significant word of rs1 and shifting it left by 1 place.
`sh2add_uw rd,rs1,rs2`	Same as above but the shift is by 2 places.
`sh3add_uw rd,rs1,rs2`	Same as above but the shift is 3 places.
`slli.uw rd, rs1, imm6`	Takes the least-significant word of rs1, zero-extends it, and shifts it left by the immediate.
`xnor rd, rs1, rs2`	Performs the bit-wise exclusive-NOR operation on rs1 and rs2. `rd = ~(rs1 ^ rs2);` [47]

Exercise 8: *Use the "max" instruction to calculate the absolute value of a signed integer.*

The Zbkb extension

Table 1.22: Zbkb instructions

Inst.	Description
`pack rd,rs1,rs2`	Packs the XLEN/2-bit lower halves of rs1 and rs2 into rd, with rs1 in the lower half and rs2 in the upper half.
`packh rd,rs1,rs2`	Packs the least-significant bytes of rs1 and rs2 into the 16 least-significant bits of rd, zero extending the rest of rd.
`packw rd,rs1,rs2`	packs the low 16 bits of rs1 and rs2 into the 32 least-significant bits of rd, sign extending the 32-bit result to the rest of rd.
`rvb rd,rs1`	Reverses the order of the bits in every byte of a register.
`xperm.b rd, rs1, rs2,`	The xperm.b instruction operates on bytes. The rs1 register contains a vector of 8 8-bit elements. The rs2 register contains a vector of 8 8-bit indexes. The result is each element in rs2 replaced by the indexed element in rs1, or zero if the index into rs2 is out of bounds. This instruction is in the extension `Zxbkx`.
`xperm.n rd, rs1, rs2`	The xperm.n instruction operates on nibbles. The rs1 register contains a vector of XLEN/4 4-bit elements. The rs2 register contains a vector of XLEN/4 4-bit indexes. The result is each element in rs2 replaced by the indexed element in rs1, or zero if the index into rs2 is out of bounds. This instruction is in the extension `Zxbkx`.
`zext.h rd, rs`	This instruction zero-extends the least-significant halfword of the source to XLEN by inserting 0's into all of the bits more significant than 15.

---

[47]The XNOR operation of two inputs returns 1 if the two inputs are equal, zero otherwise.

1.12.13   Pause instruction

Syntax:
pause

The **pause** instruction is a HINT that indicates the current hart's rate of instruction retire-
ment should be temporarily reduced or paused. The duration of its effect must be bounded
and may be zero. No state is changed. The standard says about this:

> Software can use the PAUSE instruction to reduce energy consumption while
> executing spin-wait code sequences. Multithreaded cores might temporarily re-
> linquish execution resources to other harts when PAUSE is executed. It is rec-
> ommended that a PAUSE instruction generally be included in the code sequence
> for a spin-wait loop.

Exercise 9: *Calculate how long takes a pause instruction in your machine*

1.12.14   Floating point

Floating point operations are controlled with the status register, **fcsr**. It is a 32-bit read-
/write register that selects the dynamic rounding mode for floating-point arithmetic opera-
tions and holds the accrued exception flags.

- Bit 0: NX Inexact
- Bit 1: UF Underflow
- Bit 2: OF Overflow
- Bit 3: DZ Divide by zero
- Bit 4: NV Invalid operation
- Bits 5-7: Rounding mode. This will be used when the instruction uses the dynamic
  rounding mode. See §1.24 page 89.
- Bits 8-31 Reserved.

The **fcsr** register can be read and written with the FRCSR and FSCSR instructions,
which are assembler pseudo instructions, built on the underlying CSR access instructions.

The fields of the **csr** can also be accessed individually. The instruction **frrm** reads the
rounding mode field. The instruction **fsrm** writes to it. In a similar fashion **frflags** and
fsflags read and write to the flags field.

Syntax:
    frrm rd   rd← rounding mode
    fsrm rd,rs1   rounding mode ← rs1, rd ← old rounding mode
    frflags rd   rd ← flags
    fsflags rd,rs1   flags ← rs1, rd ← old flags

Exercise 10: *Write an assembler program to show the CSR flags in the console*
Exercise 11: *Write a subroutine that returns the flags of the CSR as a 32 bit integer*
Floating point can have 4 possible precision settings: half (16 bits), single (32), double
(64) and quad(128). Most machines implement single and double precision, some implement
half precision, and (till now) none has implemented 128 bit precision.

Encodings

Floating point instructions use a slightly modified "R" format. Bits 25 and 26 define a
"format" field that is used to differenciate between half, single and double precision.
    In the figure above we have:

- Bits 0-6, the opcode, 83 (0x53).

Figure 1.17: Modified **R** Instruction layout

- Bits 7-11, the destination, 10 (0xA)
- Bits 12-14, the rounding mode, 7 (Dynamic rounding mode)
- Bits 15-19, the first source register, 11 (a1)
- Bits 20-24, the second source register 12 (a2)
- Bits 25-26, The format, in this case 0, single precision
- Bits 27-31, The code for the operation, in this case 1, the code for FADD or FSUB.

Table 1.23: Format bits (Bits 25-26)

Bits value	Instruction mnemonic	Description
0 0	S	Single precision
0 1	D	Double precision
1 0	H	Half precision
1 1	Q	128 bit Quad precision

Table 1.24: Rounding mode bits (Bits 12-14)

Bits value	Mode mnemonic	Description
000	RNE	Round to nearest
001	RTZ	Round to zero.
010	RDN	Round down towards $-\infty$
011	RUP	Round up towards $+\infty$
100	RMM	Round to nearest. Ties towards max magnitude
101	Reserved	
110	Reserved	
111	DYN	If in the instruction, selects dynamic rounding mode. If in the rounding mode register, it is *reserved*

The recognition of the rounding modes for an instruction is done in the case for the letter 'm', in the `riscv_ip` function. It uses the `riscv_rm` table of rounding modes.

```
1 /* Table of legal rounding modes. */
2 const char *const riscv_rm[8] = {"rne","rtz","rdn","rup","rmm",0,0,"dyn"};
3 /* Code snippet for case 'm': rounding modes in riscv_ip */
4     case 'm':  /* Rounding mode. */
```

```
 5          if (arg_lookup(&asarg,riscv_rm,
 6              ARRAY_SIZE(riscv_rm),&regno)) {
 7              INSERT_OPERAND(RM,*ip,regno);
 8              continue;
 9          }
10      break;
```

Floating point instructions

Single precision floating point is called extension "F", double "D" and half "H". There is provision for a "Q" extension for 128 bit numbers but none of the machines I have used implements that yet.

In the instructions below, an addddress is formed by adding the contents of the source register with a sign extended imm12.

Table 1.25: Floating point load/store instructions

Instruction	Description
`flw frd,imm12(fs1)`	Load single precision data from address at into frd.
`fsw fs2,imm12(rs1)`	Store single precision data from fs2 at address
`fld frd,imm12(fs1)`	Load double precision data from address at into frd
`fsd fs2,imm12(rs1)`	Store double precision data from fs2 at address
`flh frd,imm12(fs1)`	Load half precision data from address at into frd
`fsh fs2,imm12(rs1)`	Store half precision data from fs2 at address

In the instructions above the data will be moved without any changes.

Table 1.26: Floating point arithmetic instructions

Instruction	Description
`fadd.{h|s|d} frd,frs1,frs2`	Add. $frd \leftarrow frs1 + frs2$
`fsub.{h|s|d} frd,frs1,frs2`	Subtraction. $frd \leftarrow frs1 - frs2$
`fmul.{h|s|d} frd,frs1,frs2`	Multiplication. $frd \leftarrow frs1 \times frs2$
`fdiv.{h|s|d} frd,frs1,frs2`	Division. $frd \leftarrow frs1 \div frs2$
`fmadd.{h|s|d} fd,fs1,fs2,fs3`	Fused multiply add. $fd \leftarrow (fs1 \times fs2) + fs3$
`fmsub.{h|s|d} fd,fs1,fs2,fs3`	Fused multiply subtract. $fd \leftarrow (fs1 \times fs2) - fs3$
`fnmadd.{h|s|d} fd,fs1,fs2,fs3`	Fused negative multiply add simple precision. $fd \leftarrow (-fs1 \times fs2) - fs3$ [48]
`fnmsub.{h|s|d} fs1,fs2,fs3`	Fused negative subtraction simple precision. $fs1 \leftarrow (-fs1 \times fs2) + fs3$

---

[48] The official riscv manual acknowledges that fnmadd is a **misnomer**. They try to justify this error with:

> The FNMSUB and FNMADD instructions are counter intuitively named, owing to the naming of the corresponding instructions in MIPS-IV. The MIPS instructions were defined to negate the sum, rather than negating the product as the RISC-V instructions do, so the naming scheme was more rational at the time. The two definitions differ with respect to signed-zero results. The RISC-V definition matches the behavior of the x86 and ARM fused multiply-add instructions, but unfortunately the RISC-V FNMSUB and FNMADD instruction names are swapped compared to x86 and ARM.
>
> RISC-V Unprivileged ISA V20191214-draft page 77

In my opinion, this "explanation" doesn't explain why this misnomer is maintained...

Table 1.27: Floating point square root, min, max instructions

Instruction	Description
`fsqrt.{h|s|d} rd,rs1`	Square root. $rd \leftarrow \sqrt{rs1}$
Minimum/Maximum	
`fmin.{h|s|d} rd,rs1,rs2`	Minimum of two inputs.rd← rs1<rs2 ? rs1 : rs2
`fmax.{h|s|d} rd,rs1,rs2`	Maximum of two inputs.rd← rs1<rs2 ? rs2 : rs1
`fsgnj.{s|d|h} fd,fs1,fs2`	Sign injection of fs2 into fs1. $fd[xlen-1] \leftarrow rs2[xlen-1]$ $fd[0..xlen-2] \leftarrow rs1[0..xlen-2]$
`fsgnjn.{s|d|h}} fd,fs1,fs2`	Sign injection of neg(fs2) into fs1. $fd[xlen-1] \leftarrow\ !\ rs2[xlen-1]$ $fd[0..xlen-2] \leftarrow rs1[0..xlen-2]$
`fclass.{h|s|d} rd,fs1`	Classify fs1, returning a classification in rd, that must be an integer register. The bits in the result are explained in table §1.28, page 91. Only one bit will be set.

Table 1.28: `fclass` results

Bit number	Meaning
0	$-\infty$
1	$rs1 < 0$
2	subnormal $rs1 < 0$
3	$rs1 \equiv -0$
4	$rs1 \equiv 0$
5	subnormal $rs1 > 0$
6	$rs1 > 0$
7	$+\infty$
8	signaling NAN
9	quiet NAN

Table 1.29: Floating point conversion instructions

Instruction	Description
`fcvt.w.s rd,rs1`	Converts a single-precision floating-point number to a signed 32-bit integer. Sign-extends the32-bit result to the destination register width.
`fcvt.s.w rd,rs1`	Converts a signed 32-bit integer to a single-precision floating-point number
`fcvt.wu.s rd,rs1`	Converts a single-precision floating-point number to an unsigned 32-bit integer. Sign-extends the 32-bit result to the destination register width.
`fcvt.s.wu rd,rs1`	Converts a unsigned 32-bit integer to a single-precision floating-point number
`fcvt.l.s rd,rs1`	Converts a single-precision floating-point number to a signed 64-bit integer.
`fcvt.s.l rd,rs1`	Converts a signed 64-bit integer to a single-precision floating-point number

Table 1.29: Floating point conversion instructions

`fcvt.lu.s rd,rs1`	Converts a single-precision floating-point number to an unsigned 64-bit integer.
`fcvt.s.lu rd,rs1`	Converts a unsigned 64-bit integer to a single-precision floating-point number
Other precisions	
`fcvt.{l|w|lu|wu}.{s|d|h}`	Convert floating point to integer in the different sizes and precisions
`fcvt.{s|d|h}.{l|lu|w|wu}`	Convert integer to floating point in the different sizes and precisions
`fcvt.{h|s|d}.{h|s|d}`	Convert between different floating point formats.
`fmv.x.w rd,rs1`	Moves the single-precision value in floating-point register rs1represented in IEEE 754-2008 encoding to the lower 32 bits of integer register rd. The higher 32 bits of the destination register are filled with copies of the floating-point number's sign bit.
`fmv.w.x`	Moves the single-precision value encoded in IEEE 754-2008 standard encoding from the lower 32 bits of integer register rs1 to the floating-point register rd.

Absent from the table of instructions above are the ones introduced in 2023: the "Zfa" extension, that will make possible to load some immediates into fp registers, minimum/maximum operations with NANs and others. [49]

Floating-point compare instructions (`feq, flt, fle`) perform the specified comparison between floating-point registers (rs1 = rs2, rs1 < rs2, rs1 ≤ rs2) writing 1 to the integer register rd if the condition holds, and 0 otherwise.

Table 1.30: Floating point comparison instructions

Instruction	Description
`feq.{h|s|d} rd,fs1,fs2`	Equality. $rd \leftarrow (fs1 = fs2)$
`flt.{h|s|d} rd,fs1,fs2`	Less than comparison. $rd \leftarrow (fs1 < fs2)$
`fle.{h|s|d} rd,fs1,fs2`	Less equal comparison. $rd \leftarrow (fs1 \leq fs2)$

`flt.{h|s|d}` and `fle.{h|s|d}` perform signaling comparisons: they set the invalid operation exception flag if either input is NaN. `feq` performs a quiet comparison: it only sets the invalid operation exception flag if either input is a *signaling* NaN. For all three instructions, the result is 0 if either operand is NaN.

1.12.15   Atomic instructions

These instructions read, modify, and write memory atomically to provide synchronisation across several RISC-V harts in the same memory space. The theoretical model whereupon riscv is based is described in the paper **Memory consistency and event ordering in scalable shared-memory multiprocessors**.In *Proceedings of the 17th Annual International Symposium on Computer Architecture*, pages 15–26, 1990.

Complex atomic memory operations on a single memory word or doubleword are performed with the load-reserved (`lr`) and store-conditional (`sc`) instructions.

---

[49]They are not supported in tiny-asm, nor do they have any implementation in actual hardware yet.

Table 1.31: load reserved/store conditional instructions

Instruction	Description
`lr.{d|w}d rd,rs1`	Loads a double or a single word from the address in rs1, places the sign-extended value in rd, and registers a reservation set, a set of bytes that subsumes the bytes in the addressed double word. For `lr.w` the value is extended to 64 bits.
`sc.{d|w} rd,rs1,rs2`	Conditionally writes a word in rs2 to the address in rs1: the `sc.{w|d}` succeeds only if the reservation is still valid and the reservation set contains the bytes being written. If it succeeds, the instruction writes the word in rs2 to memory, and it writes zero to rd. If it fails, the instruction does not write to memory, and it writes a nonzero value to rd. Regardless of success or failure, executing an `sc` instruction invalidates any reservation held by this hart.

The aq and rl bits

The riscv specifications explain:

> To follow more closely the specifications in the paper cited above for release consistency, each atomic instruction has two bits, `aq` and `rl`, used to specify additional memory ordering constraints as viewed by other riscv harts.The bits order accesses to one of the two address domains, memory or I/O, depending on which address domain the atomic instruction is accessing.

> If both bits are clear, no additional ordering constraints are imposed on the atomic memory operation.If only the aq bit is set, the atomic memory operation is treated as an acquire access,i.e., no following memory operations on this RISC-V hart can be observed to take place before the acquire memory operation.

> If only the rl bit is set, the atomic memory operation is treated as a release access, i.e., the release memory operation cannot be observed to take place before any earlier memory operations on this RISC-V hart. If both the aq and rl bits are set, the atomic memory operation is sequentially consistent and cannot be observed to happen before any earlier memory operations or after any later memory operations in the same riscv hart and to the same address domain.[50]

These bits can be specified by writing their state after the instruction, for instance:
```
amoxor.d rd,rs1,rs2,aq
amoxor.d rd,rs1,rs2,aqrl
amoxor.d rd,rs1,rs2,rl
```

Encodings

New bits are: bit 26, for `aq` and bit 25, for `rl`. The width of the instruction is written in bits 12 to 14. A value of 3, for instance, means $2^3 \rightarrow 8$.

---

[50] §10.1, page 55

Figure 1.18: Atomic instructions layout

Table 1.32: Atomic Memory Operation instructions

Instruction	Description
`amoswap.{w|d}` `rd,rs1,rs2,[aq|aqrl|rl]`	Loads a data value from the address in rs1 , place the value into register rd , swaps the loaded value and the original value in rs2 , then stores the result back to the address in rs1
`amoadd.{w|d}` `rd,rs1,rs2,,[aq|aqrl|rl]`	Loads a data value from the address in rs1 , place the value into register rd , adds the loaded value and the original value in rs2 , then stores the result back to the address in rs1
`amoand.{w|d}` `rd,rs1,rs2,[aq|aqrl|rl]`	Loads a data value from the address in rs1 , place the value into register rd , does and AND operation between the loaded value and the original value in rs2 , then stores the result back to the address in rs1
`amoor.{w|d}` `rd,rs1,rs2,[aq|aqrl|rl]`	Loads a data value from the address in rs1 , place the value into register rd , does an OR operation between the loaded value and the original value in rs2 , then stores the result back to the address in rs1
`amoxor.{w|d}` `rd,rs1,rs2,[aq|aqrl|rl]`	Loads a data value from the address in rs1 , place the value into register rd , does an XOR operation between the loaded value and the original value in rs2 , then stores the result back to the address in rs1
`amomax[u].{w|d}` `rd,rs1,rs2,[aq|aqrl|rl]`	Loads a data value from the address in rs1 , place the value into register rd , finds the signed (or the unsigned) maximum value between the loaded value and the original value in rs2 , then stores the result back to the address in rs1
`amomin[u].{w|d}` `rd,rs1,rs2,[aq|aqrl|rl]`	Loads a data value from the address in rs1 , place the value into register rd , finds the signed (or unsigned) minimum value between the loaded value and the original value in rs2 , then stors the result back to the address in rs1

1.13   Instructions specific to the Thead processor

All these instructions are prefixed with the letters `"th."`.

Table 1.33: Thead instructions

Instruction	Description
`th.addsl rd,rs1,rs2,imm2`	$rd \leftarrow rs1 + (rs2 << imm2)$ Add with shifted register.
`th.ext rd,rs1,imm1,imm2`	$rd \leftarrow rs1[imm1 : imm2]$. Extract bits `imm1` to `imm2` with sign extension
`th.extu rd,rs1,imm1,imm2`	$rd \leftarrow rs1[imm1 : imm2]$. Extract bits `imm1` to `imm2` with zero extension
`th.ff0 rd,rs`	Finds the first bit with the value of 0 from the highest bit of rs1 and writes the result back into the rd register. If the highest bit of rs1 is 0, the result 0 is returned. If all the bits in rs1 are 1, the result 64 is returned.
`th.ff1 rd,rs`	Finds the first bit with the value of 1 from the highest bit of rs1 and writes the index of this bit back into rd. If the highest bit of rs1 is 1, the result 0 is returned. If all the bits in rs1 are 1, the result 64 is returned.

`th.rev rd,rs1`	Reverses the bytes in rs1.
	 `rd[7]← rs[0] rd[6]← rs[1] rd[5]← rs[2]`
	 `rd[4]← rs[3] rd[3]← rs[4] rd[2]← rs[5]`
	 `rd[1]← rs[6] rd[0]← rs[7]`

`th.revw rd,rs1`	Reverses the bytes in lower word of rs1.
	 `rd[3]← rs[0] rd[2]← rs[1] rd[1]← rs[2]`
	 `rd[0]← rs[3]`

`th.tst rd,rs1,imm6`	rd contains the bit at position imm6 of rs1.
`th.tstnbz rd, rs1`	Tests for a zero byte in rs1. Each byte of rd will be either 0xff (the corresponding byte is zero) or 0 (the corresponding byte is different than zero)
`th.lbia rd,(rs1),imm5,imm2`	Post-increment. $signExt(rd \leftarrow mem[rs1])$; $rs1 \leftarrow rs1 + imm5 << imm2$ rd and rs1 must be different registers.
`th.lbib rd,(rs1),imm5,imm2`	Pre-increment. $rs1 \leftarrow rs1 + imm5 << imm2$; $signExt(rd \leftarrow mem[rs1])$ rd and rs1 must be different registers.
`th.lbuia rd,(rs1),imm5,imm2`	Post-increment. $zeroExt(rd \leftarrow mem[rs1])$; $rs1 \leftarrow rs1 + imm5 << imm2$ rd and rs1 must be different registers.
`th.lbuib rd,(rs1),imm5,imm2`	Pre-increment. $rs1 \leftarrow rs1 + imm5 << imm2$; $zeroExt(rd \leftarrow mem[rs1])$ rd and rs1 must be different registers.
`th.ldd rd1,rd2, (rs1),imm2`	Load pair of registers. $address \leftarrow rs1 + zero\_extend(imm2 << 4)$ $rd1 \leftarrow mem[address + 7 : address]$ $rd2 \leftarrow mem[address + 15 : address + 8]$

Table 1.33: Thead instructions

`th.ldia rd,(rs1),imm5,imm2`	Load byte with post increment $rd \leftarrow signExt(mem[rs1 + 7 : rs1])$ $rs1 \leftarrow rs1 + signExt(imm5 << imm2)$
`th.ldib rd,(rs1),imm5,imm2`	Load byte with pre increment $rs1 \leftarrow rs1 + signExt(imm5 << imm2)$ $rd \leftarrow signExt(mem[rs1 + 7 : rs1])$
`th.lhia rd,(rs1),imm5,imm2`	Load half word with post increment $rd \leftarrow signExt(mem[rs1 + 1 : rs1])$ $rs1 \leftarrow rs1 + signExt(imm5 << imm2)$
`th.lhib rd,(rs1),imm5,imm2`	Load half word with pre increment $rd \leftarrow signExt(mem[rs1 + 1 : rs1])$ $rs1 \leftarrow rs1 + signExt(imm5 << imm2)$
`th.lhuia rd,(rs1),imm5,imm2`	Load half word with post increment and zero extend. $rd \leftarrow zeroExt(mem[rs1 + 1 : rs1])$ $rs1 \leftarrow rs1 + signExt(imm5 << imm2)$
`th.lhib rd,(rs1),imm5,imm2`	Load half word with pre increment and zero extend. $rd \leftarrow zeroExt(mem[rs1 + 1 : rs1])$ $rs1 \leftarrow rs1 + signExt(imm5 << imm2)$
`th.lrb rd,rs1,imm2`	Load sign extended byte with shifted register. $rd \leftarrow signExt(mem[rs1 + (rs2 << imm2)])$
`th.lrbu rd,rs1,imm2`	Load zero extended byte with shifted register. $rd \leftarrow zeroExt(mem[rs1 + (rs2 << imm2)])$
`th.lrd rd,rs1,imm2`	Load double word with shifted register. $rd \leftarrow mem[rs1 + (rs2 << imm2)]$
`th.lrh rd,rs1,imm2`	Load half word with sign extend and shifted register. $rd \leftarrow mem[rs1 + (rs2 << imm2)]$
`th.lrhu rd,rs1,imm2`	Load half word with zero extend and shifted register. $rd \leftarrow mem[rs1 + (rs2 << imm2)]$
`th.lrw rd,rs1,imm2`	Load word with sign extend and shifted register. $rd \leftarrow mem[rs1 + (rs2 << imm2)]$
`th.lrwu rd,rs1,imm2`	Load word with zero extend and shifted register. $rd \leftarrow mem[rs1 + (rs2 << imm2)]$
`th.lurb rd,rs1,rs2,imm2`	Load byte, shift it, then sign extend result. $rd \leftarrow signExt(mem[rs1 + zeroExt(rs2[0 : 31])] << imm2)$
`th.lurbu rd,rs1,rs2,imm2`	Load byte, shift it, then zero extend result. $rd \leftarrow zeroExt(mem[rs1 + zeroExt(rs2[0 : 31])] << imm2)$
`th.lurd rd,rs1,rs2,imm2`	Load double word, shift the result. $rd \leftarrow mem[rs1 + zeroExt(rs2[0 : 31])] << imm2$
`th.lurh rd,rs1,rs2,imm2`	Load half word, shift the result. $rd \leftarrow signExt(mem[rs1 + zeroExt(rs2[0 : 31])] << imm2)$
`th.lurhu rd,rs1,rs2,imm2`	Load half word, shift the result. $rd \leftarrow zeroExt(mem[rs1 + zeroExt(rs2[0 : 31])] << imm2)$
`th.lurw rd,rs1,rs2,imm2`	Load word, shift the result. $rd \leftarrow signExt(mem[rs1 + zeroExt(rs2[0 : 31])] << imm2)$

Table 1.33: Thead instructions

`th.lurwu rd1,rd2,(rs1),imm2`	Load word register pair. $address \leftarrow rs1 + zeroExt(imm2 << 3)$ $rd1 \leftarrow signExt(mem[address + 3 : address])$ $rd2 \leftarrow signExt(mem[address + 7 : address + 4])$
`th.lwd rd,rs1,(rs2),imm2`	

1.14 Pseudo instructions

The accepted policy under risc-v is the opposite to the ARM64 assembler. Under ARM64, the assembler will issue an error if an instruction alias expands to more than one instruction. Here, it is quite the opposite, the assembler (and above all, the linker) is responsible for expanding high level macros.

Table 1.34: Pseudo instructions

Pseudo	Base instruction	Meaning
`beqz rs, offset`	`beq rs, x0, offset`	Branch if = zero
`bgez rs, offset`	`bge rs, x0, offset`	Branch if $\geq$zero
`bgt rs, rt, offset`	`blt rt, rs, offset`	Branch if $>$
`bgtu rs, rt, offset`	`bltu rt, rs, offset`	Branch if $>$, unsigned
`bgtz rs, offset`	`blt x0, rs, offset`	Branch if $>$ zero
`ble rs, rt, offset`	`bge rt, rs, offset`	Branch if $\leq$
`bleu rs, rt, offset`	`bgeu rt, rs, offset`	Branch if $\leq$, unsigned
`blez rs, offset`	`bge x0, rs, offset`	Branch if $\leq$ zero
`bltz rs, offset`	`blt rs, x0, offset`	Branch if $<$ zero
`bnez rs, offset`	`bne rs, x0, offset`	Branch if $\neq$ zero
`call offset`	`auipc x1, offset[31:12];` `jalr x1, x1, offset[11:0]`	Call far-away subroutine
`fabs.d rd, rs`	`fsgnjx.d rd, rs, rs`	Double-precision absolute value. Just an alias.
`fabs.s rd, rs`	`fsgnjx.s rd, rs, rs`	Single-precision absolute value
`fence`	`fence iorw, iorw`	Fence on all memory and I/O
`fl{w|d} rd, symbol, rt`	`auipc rt, symbol[31:12];` `fl{w|d} rd, symbol[11:0](rt)`	Floating-point load global
`fmv.d rd, rs`	`fsgnj.d rd, rs, rs`	Copy double-precision register
`fmv.s rd, rs`	`fsgnj.s rd, rs, rs`	Copy single-precision register
`fneg.d rd, rs`	`fsgnjn.d rd, rs, rs`	Double-precision negate
`fneg.s rd, rs`	`fsgnjn.s rd, rs, rs`	Single-precision negate
`fs{w|d} rd,symbol, rt`	`auipc rt, symbol[31:12];` `fs{w|d} rd, symbol[11:0](rt)`	Floating-point store global
`j offset`	`jal x0, offset`	Jump
`jal offset`	`jal x1, offset`	Jump and link
`jalr rs jalr`	`x1, rs, 0`	Jump and link register
`jr rs`	`jalr x0, rs, 0`	Jump register

Table 1.34: Pseudo instructions

l{b|h|w|d} rd, symbol	auipc rd, symbol[31:12]; l{b|h|w|d} rd, symbol[11:0](rd)	Load global
la rd, symbol	auipc rd, symbol@GOT[31:12]; lw|d rd, symbol@GOT[11:0](rd)	Load address With .option pic
la rd, symbol	auipc rd, symbol[31:12]; addi rd, rd, symbol[11:0]	Load address With .option nopic (Default)
li rd, immediate	Myriad sequences	Load immediate
lla rd, symbol	auipc rd, symbol[31:12]; addi rd, rd, symbol[11:0]	Load local address
mv rd, rs	addi rd, rs, 0	Copy register
neg rd, rs	sub rd, x0, rs	Two's complement
negw rd, rs	subw rd, x0, rs	Two's complement word
nop	addi x0, x0, 0	No operation
not rd, rs	xori rd, rs, -1	Ones' complement
pause	fence w, 0	PAUSE hint
ret	jalr x0, x1, 0	Return from subroutine
s{b|h|w|d} rd, symbol, rt	auipc rt, symbol[31:12]; s{b|h|w|d} rd, symbol[11:0](rt)	Store global
seqz rd, rs	sltiu rd, rs, 1	Set if = zero
sext.b rd, rs	slli rd, rs, XLEN - 8; srai rd, rd, XLEN - 8	Sign extend byte It will expand to another instruction sequence when B extension is available
sext.h rd, rs	slli rd, rs, XLEN - 16; srai rd, rd, XLEN - 16	Sign extend half word It will expand to another instruction sequence when B extension is available
sext.w rd, rs	addiw rd, rs, 0	Sign extend word
sgtz rd, rs	slt rd, x0, rs	Set if > zero
sltz rd, rs	slt rd, rs, x0	Set if < zero
snez rd, rs	sltu rd, x0, rs	Set if ≠ zero
tail offset	auipc x6, offset[31:12]; jalr x0, x6, offset[11:0]	Tail call far-away subroutine.
zext.b rd, rs	andi rd, rs, 255	Zero extend byte
zext.h rd, rs	slli rd, rs, XLEN - 16; srli rd, rd, XLEN - 16	Zero extend half word It will expand to another instruction sequence when B extension is available
zext.w rd, rs	slli rd, rs, XLEN - 32; srli rd, rd, XLEN - 32	Zero extend word It will expand to another instruction sequence when B extension is available

Exercise 12: *Mismatch between source and disassembly. Explain*
Consider this program:

```
1     .globl main
2 main:
3     la a0,.L2
4     jr ra
5 .L2:
```

When we assemble this with `gcc -o addr.o addr.s` we obtain an object file. When we disassemble it, we obtain:

```
1 Disassembly of section .text:
2
3 0000000000000000 <main>:
4 0:   00000517            auipc   a0,0x0
5 4:   00050513            mv  a0,a0
6 8:   00008067            ret
```

Explain why this mismatch.

1.15   The full opcode table

Instruction	Extension	Parameters	Flags
add	C	Cc,Cc,CL	alias
add	C	Ct,Cc,CK	alias
add	C	d,CU,CV	alias
add	C	d,CU,Co	alias
add	C	d,CV,CU	alias
add	C	d,Cz,CV	alias
add	I	d,s,j	alias
add	I	d,s,t	–
add	I	d,s,t,1	–
add.uw	Zba	d,s,t	–
addi	C	Cc,Cc,CL	alias
addi	C	Ct,Cc,CK	alias
addi	C	d,CU,Cj	alias
addi	C	d,CU,z	alias
addi	C	d,CV,z	alias
addi	C	d,Cz,Co	alias
addi	I	d,s,j	–
addiw	C	d,CU,Co	alias
addiw	I	d,s,j	–
addw	C	Cs,Ct,Cw	alias
addw	C	Cs,Cw,Ct	alias
addw	C	d,CU,Co	alias
addw	I	d,s,j	alias
addw	I	d,s,t	–
aes32dsi	Zknd	d,s,t,y	crypto
aes32dsmi	Zknd	d,s,t,y	crypto
aes32esi	Zkne	d,s,t,y	crypto
aes32esmi	Zkne	d,s,t,y	crypto
aes64ds	Zknd	d,s,t	crypto
aes64dsm	Zknd	d,s,t	crypto
aes64es	Zkne	d,s,t	crypto
aes64esm	Zkne	d,s,t	crypto
aes64im	Zknd	d,s	crypto
aes64ks1i	Zknd |Zkne	d,s,Y	crypto
aes64ks2	Zknd | Zkne	d,s,t	crypto
amoadd.d	A	d,t,0(s)	dref | 8-byte
amoadd.d.aq	A	d,t,0(s)	dref | 8-byte
amoadd.d.aqrl	A	d,t,0(s)	dref | 8-byte
amoadd.d.rl	A	d,t,0(s)	dref | 8-byte
amoadd.w	A	d,t,0(s)	dref | 4-byte
amoadd.w.aq	A	d,t,0(s)	dref | 4-byte
amoadd.w.aqrl	A	d,t,0(s)	dref | 4-byte
amoadd.w.rl	A	d,t,0(s)	dref | 4-byte
amoand.d	A	d,t,0(s)	dref | 8-byte
amoand.d.aq	A	d,t,0(s)	dref | 8-byte
amoand.d.aqrl	A	d,t,0(s)	dref | 8-byte
amoand.d.rl	A	d,t,0(s)	dref | 8-byte
amoand.w	A	d,t,0(s)	dref | 4-byte
amoand.w.aq	A	d,t,0(s)	dref | 4-byte

Instruction	Extension	Parameters	Flags
amoand.w.aqrl	A	d,t,0(s)	dref | 4-byte
amoand.w.rl	A	d,t,0(s)	dref | 4-byte
amomax.d	A	d,t,0(s)	dref | 8-byte
amomax.d.aq	A	d,t,0(s)	dref | 8-byte
amomax.d.aqrl	A	d,t,0(s)	dref | 8-byte
amomax.d.rl	A	d,t,0(s)	dref | 8-byte
amomax.w	A	d,t,0(s)	dref | 4-byte
amomax.w.aq	A	d,t,0(s)	dref | 4-byte
amomax.w.aqrl	A	d,t,0(s)	dref | 4-byte
amomax.w.rl	A	d,t,0(s)	dref | 4-byte
amomaxu.d	A	d,t,0(s)	dref | 8-byte
amomaxu.d.aq	A	d,t,0(s)	dref | 8-byte
amomaxu.d.aqrl	A	d,t,0(s)	dref | 8-byte
amomaxu.d.rl	A	d,t,0(s)	dref | 8-byte
amomaxu.w	A	d,t,0(s)	dref | 4-byte
amomaxu.w.aq	A	d,t,0(s)	dref | 4-byte
amomaxu.w.aqrl	A	d,t,0(s)	dref | 4-byte
amomaxu.w.rl	A	d,t,0(s)	dref | 4-byte
amomin.d	A	d,t,0(s)	dref | 8-byte
amomin.d.aq	A	d,t,0(s)	dref | 8-byte
amomin.d.aqrl	A	d,t,0(s)	dref | 8-byte
amomin.d.rl	A	d,t,0(s)	dref | 8-byte
amomin.w	A	d,t,0(s)	dref | 4-byte
amomin.w.aq	A	d,t,0(s)	dref | 4-byte
amomin.w.aqrl	A	d,t,0(s)	dref | 4-byte
amomin.w.rl	A	d,t,0(s)	dref | 4-byte
amominu.d	A	d,t,0(s)	dref | 8-byte
amominu.d.aq	A	d,t,0(s)	dref | 8-byte
amominu.d.aqrl	A	d,t,0(s)	dref | 8-byte
amominu.d.rl	A	d,t,0(s)	dref | 8-byte
amominu.w	A	d,t,0(s)	dref | 4-byte
amominu.w.aq	A	d,t,0(s)	dref | 4-byte
amominu.w.aqrl	A	d,t,0(s)	dref | 4-byte
amominu.w.rl	A	d,t,0(s)	dref | 4-byte
amoor.d	A	d,t,0(s)	dref | 8-byte
amoor.d.aq	A	d,t,0(s)	dref | 8-byte
amoor.d.aqrl	A	d,t,0(s)	dref | 8-byte
amoor.d.rl	A	d,t,0(s)	dref | 8-byte
amoor.w	A	d,t,0(s)	dref | 4-byte
amoor.w.aq	A	d,t,0(s)	dref | 4-byte
amoor.w.aqrl	A	d,t,0(s)	dref | 4-byte
amoor.w.rl	A	d,t,0(s)	dref | 4-byte
amoswap.d	A	d,t,0(s)	dref | 8-byte
amoswap.d.aq	A	d,t,0(s)	dref | 8-byte
amoswap.d.aqrl	A	d,t,0(s)	dref | 8-byte
amoswap.d.rl	A	d,t,0(s)	dref | 8-byte
amoswap.w	A	d,t,0(s)	dref | 4-byte
amoswap.w.aq	A	d,t,0(s)	dref | 4-byte
amoswap.w.aqrl	A	d,t,0(s)	dref | 4-byte
amoswap.w.rl	A	d,t,0(s)	dref | 4-byte
amoxor.d	A	d,t,0(s)	dref | 8-byte
amoxor.d.aq	A	d,t,0(s)	dref | 8-byte
amoxor.d.aqrl	A	d,t,0(s)	dref | 8-byte
amoxor.d.rl	A	d,t,0(s)	dref | 8-byte

Instruction	Extension	Parameters	Flags
amoxor.w	A	d,t,0(s)	dref | 4-byte
amoxor.w.aq	A	d,t,0(s)	dref | 4-byte
amoxor.w.aqrl	A	d,t,0(s)	dref | 4-byte
amoxor.w.rl	A	d,t,0(s)	dref | 4-byte
and	C	Cs,Ct,Cw	alias
and	C	Cs,Cw,Co	alias
and	C	Cs,Cw,Ct	alias
and	I	d,s,j	alias
and	I	d,s,t	–
andi	C	Cs,Cw,Co	alias
andi	I	d,s,j	–
andn	Zbb | Zbkb	d,s,t	–
auipc	I	d,u	–
bclr	Zbs	d,s,>	alias
bclr	Zbs	d,s,t	–
bclri	Zbs	d,s,>	–
beq	C	Cs,Cz,Cp	alias | condbranch
beq	I	s,t,p	condbranch
beqz	C	Cs,Cp	alias | condbranch
beqz	I	s,p	alias | condbranch
bext	Zbs	d,s,>	alias
bext	Zbs	d,s,t	–
bexti	Zbs	d,s,>	–
bge	I	s,t,p	condbranch
bgeu	I	s,t,p	condbranch
bgez	I	s,p	alias | condbranch
bgt	I	t,s,p	alias | condbranch
bgtu	I	t,s,p	alias | condbranch
bgtz	I	t,p	alias | condbranch
binv	Zbs	d,s,>	alias
binv	Zbs	d,s,t	–
binvi	Zbs	d,s,>	–
ble	I	t,s,p	alias | condbranch
bleu	I	t,s,p	alias | condbranch
blez	I	t,p	alias | condbranch
blt	I	s,t,p	condbranch
bltu	I	s,t,p	condbranch
bltz	I	s,p	alias | condbranch
bne	C	Cs,Cz,Cp	alias | condbranch
bne	I	s,t,p	condbranch
bnez	C	Cs,Cp	alias | condbranch
bnez	I	s,p	alias | condbranch
brev8	Zbkb	d,s	–
brev8	Zbkb	d,s	–
bset	Zbs	d,s,>	alias
bset	Zbs	d,s,t	–
bseti	Zbs	d,s,>	–
c.add	C	d,CV	–
c.addi	C	d,Co	–
c.addi16sp	C	Cc,CL	–
c.addi4spn	C	Ct,Cc,CK	–
c.addiw	C	d,Co	–
c.addw	C	Cs,Ct	–

Instruction	Extension	Parameters	Flags
c.and	C	Cs,Ct	-
c.andi	C	Cs,Co	-
c.beqz	C	Cs,Cp	condbranch
c.bnez	C	Cs,Cp	condbranch
c.ebreak	C	--	-
c.fld	D and C	CD,Cl(Cs)	dref | 8-byte
c.fldsp	D and C	D,Cn(Cc)	dref | 8-byte
c.flw	F and C	CD,Ck(Cs)	dref | 4-byte
c.flwsp	F and C	D,Cm(Cc)	dref | 4-byte
c.fsd	D and C	CD,Cl(Cs)	dref | 8-byte
c.fsdsp	D and C	CT,CN(Cc)	dref | 8-byte
c.fsw	F and C	CD,Ck(Cs)	dref | 4-byte
c.fswsp	F and C	CT,CM(Cc)	dref | 4-byte
c.j	C	Ca	branch
c.jal	C	Ca	jsr
c.jalr	C	d	jsr
c.jr	C	d	branch
c.ld	C	Ct,Cl(Cs)	dref | 8-byte
c.ldsp	C	d,Cn(Cc)	dref | 8-byte
c.li	C	d,Co	-
c.lui	C	d,Cu	-
c.lw	C	Ct,Ck(Cs)	dref | 4-byte
c.lwsp	C	d,Cm(Cc)	-
c.mv	C	d,CV	-
c.nop	C	--	alias
c.nop	C	Cj	alias
c.or	C	Cs,Ct	-
c.sd	C	Ct,Cl(Cs)	dref | 8-byte
c.sdsp	C	CV,CN(Cc)	dref | 8-byte
c.slli	C	d,C>	-
c.slli64	C	d	-
c.srai	C	Cs,C>	-
c.srai64	C	Cs	-
c.srli	C	Cs,C>	-
c.srli64	C	Cs	-
c.sub	C	Cs,Ct	-
c.subw	C	Cs,Ct	-
c.sw	C	Ct,Ck(Cs)	dref | 4-byte
c.swsp	C	CV,CM(Cc)	dref | 4-byte
c.unimp	C	--	0xffffU,-
c.xor	C	Cs,Ct	-
call	I	c	macro
call	I	d,c	macro
cbo.clean	Zicbom	0(s)	-
cbo.flush	Zicbom	0(s)	-
cbo.inval	Zicbom	0(s)	-
cbo.zero	Zicboz	0(s)	-
clmul	Zbc | Zbkc	d,s,t	-
clmulh	Zbc | Zbkc	d,s,t	-
clmulr	Zbc	d,s,t	-
clz	Zbb	d,s	-
clzw	Zbb	d,s	-

Instruction	Extension	Parameters	Flags
cpop	Zbb	d,s	–
cpopw	Zbb	d,s	–
csrc	Zicsr	E,Z	alias
csrc	Zicsr	E,s	alias
csrci	Zicsr	E,Z	alias
csrr	Zicsr	d,E	alias
csrrc	Zicsr	d,E,Z	alias
csrrc	Zicsr	d,E,s	–
csrrci	Zicsr	d,E,Z	–
csrrs	Zicsr	d,E,Z	alias
csrrs	Zicsr	d,E,s	–
csrrsi	Zicsr	d,E,Z	–
csrrw	Zicsr	d,E,Z	alias
csrrw	Zicsr	d,E,s	–
csrrwi	Zicsr	d,E,Z	–
csrs	Zicsr	E,Z	alias
csrs	Zicsr	E,s	alias
csrsi	Zicsr	E,Z	alias
csrw	Zicsr	E,Z	alias
csrw	Zicsr	E,s	alias
csrwi	Zicsr	E,Z	alias
ctz	Zbb	d,s	–
ctzw	Zbb	d,s	–
div	M	d,s,t	–
divu	M	d,s,t	–
divuw	M	d,s,t	–
divw	M	d,s,t	–
dret	I	--	–
ebreak	C	--	alias
ebreak	I	--	–
ecall	I	--	–
fabs.d	D_inx	D,U	alias
fabs.h	Zfh_inx	D,U	alias
fabs.q	q_inx	D,U	alias
fabs.s	f_inx	D,U	alias
fadd.d	D_inx	D,S,T	–
fadd.d	D_inx	D,S,T,m	–
fadd.h	Zfh_inx	D,S,T	–
fadd.h	Zfh_inx	D,S,T,m	–
fadd.q	q_inx	D,S,T	–
fadd.q	q_inx	D,S,T,m	–
fadd.s	f_inx	D,S,T	–
fadd.s	f_inx	D,S,T,m	–
fclass.d	D_inx	d,S	–
fclass.h	Zfh_inx	d,S	–
fclass.q	q_inx	d,S	–
fclass.s	f_inx	d,S	–
fcvt.d.h	Zfhmin and D_NX	D,S	–
fcvt.d.l	D_inx	D,s	–
fcvt.d.l	D_inx	D,s,m	–
fcvt.d.lu	D_inx	D,s	–
fcvt.d.lu	D_inx	D,s,m	–
fcvt.d.q	q_inx	D,S	–

Instruction	Extension	Parameters	Flags
fcvt.d.q	q_inx	D,S,m	–
fcvt.d.s	D_inx	D,S	–
fcvt.d.w	D_inx	D,s	–
fcvt.d.wu	D_inx	D,s	–
fcvt.h.d	Zfhmin and D_inx	D,S	–
fcvt.h.d	Zfhmin and D_inx	D,S,m	–
fcvt.h.l	Zfh_inx	D,s	–
fcvt.h.l	Zfh_inx	D,s,m	–
fcvt.h.lu	Zfh_inx	D,s	–
fcvt.h.lu	Zfh_inx	D,s,m	–
fcvt.h.q	Zfhmin and_q_inx	D,S	–
fcvt.h.q	Zfhmin and_q_inx	D,S,m	–
fcvt.h.s	Zfhmin_INX	D,S	–
fcvt.h.s	Zfhmin_INX	D,S,m	–
fcvt.h.w	Zfh_inx	D,s	–
fcvt.h.w	Zfh_inx	D,s,m	–
fcvt.h.wu	Zfh_inx	D,s	–
fcvt.h.wu	Zfh_inx	D,s,m	–
fcvt.l.d	D_inx	d,S	–
fcvt.l.d	D_inx	d,S,m	–
fcvt.l.h	Zfh_inx	d,S	–
fcvt.l.h	Zfh_inx	d,S,m	–
fcvt.l.q	q_inx	d,S	–
fcvt.l.q	q_inx	d,S,m	–
fcvt.l.s	f_inx	d,S	–
fcvt.l.s	f_inx	d,S,m	–
fcvt.lu.d	D_inx	d,S	–
fcvt.lu.d	D_inx	d,S,m	–
fcvt.lu.h	Zfh_inx	d,S	–
fcvt.lu.h	Zfh_inx	d,S,m	–
fcvt.lu.q	q_inx	d,S	–
fcvt.lu.q	q_inx	d,S,m	–
fcvt.lu.s	f_inx	d,S	–
fcvt.lu.s	f_inx	d,S,m	–
fcvt.q.d	q_inx	D,S	–
fcvt.q.h	Zfhmin and_q_inx	D,S	–
fcvt.q.l	q_inx	D,s	–
fcvt.q.l	q_inx	D,s,m	–
fcvt.q.lu	q_inx	D,s	–
fcvt.q.lu	q_inx	D,s,m	–
fcvt.q.s	q_inx	D,S	–
fcvt.q.w	q_inx	D,s	–
fcvt.q.wu	q_inx	D,s	–
fcvt.s.d	D_inx	D,S	–
fcvt.s.d	D_inx	D,S,m	–
fcvt.s.h	Zfhmin_INX	D,S	–
fcvt.s.l	f_inx	D,s	–
fcvt.s.l	f_inx	D,s,m	–
fcvt.s.lu	f_inx	D,s	–

Instruction	Extension	Parameters	Flags
fcvt.s.lu	f_inx	D,s,m	–
fcvt.s.q	q_inx	D,S	–
fcvt.s.q	q_inx	D,S,m	–
fcvt.s.w	f_inx	D,s	–
fcvt.s.w	f_inx	D,s,m	–
fcvt.s.wu	f_inx	D,s	–
fcvt.s.wu	f_inx	D,s,m	–
fcvt.w.d	D_inx	d,S	–
fcvt.w.d	D_inx	d,S,m	–
fcvt.w.h	Zfh_inx	d,S	–
fcvt.w.h	Zfh_inx	d,S,m	–
fcvt.w.q	q_inx	d,S	–
fcvt.w.q	q_inx	d,S,m	–
fcvt.w.s	f_inx	d,S	–
fcvt.w.s	f_inx	d,S,m	–
fcvt.wu.d	D_inx	d,S	–
fcvt.wu.d	D_inx	d,S,m	–
fcvt.wu.h	Zfh_inx	d,S	–
fcvt.wu.h	Zfh_inx	d,S,m	–
fcvt.wu.q	q_inx	d,S	–
fcvt.wu.q	q_inx	d,S,m	–
fcvt.wu.s	f_inx	d,S	–
fcvt.wu.s	f_inx	d,S,m	–
fdiv.d	D_inx	D,S,T	–
fdiv.d	D_inx	D,S,T,m	–
fdiv.h	Zfh_inx	D,S,T	–
fdiv.h	Zfh_inx	D,S,T,m	–
fdiv.q	q_inx	D,S,T	–
fdiv.q	q_inx	D,S,T,m	–
fdiv.s	f_inx	D,S,T	–
fdiv.s	f_inx	D,S,T,m	–
fence	I	--	alias
fence	I	P,Q	–
fence.i	Zifencei	--	–
fence.tso	I	--	–
feq.d	D_inx	d,S,T	–
feq.h	Zfh_inx	d,S,T	–
feq.q	q_inx	d,S,T	–
feq.s	f_inx	d,S,T	–
fge.d	D_inx	d,T,S	–
fge.h	Zfh_inx	d,T,S	–
fge.q	q_inx	d,T,S	–
fge.s	f_inx	d,T,S	–
fgt.d	D_inx	d,T,S	–
fgt.h	Zfh_inx	d,T,S	–
fgt.q	q_inx	d,T,S	–
fgt.s	f_inx	d,T,S	–
fld	D	D,A,s	macro
fld	D	D,o(s)	dref | 8-byte
fld	D and C	CD,Cl(Cs)	alias | dref|8-byte
fld	D and C	D,Cn(Cc)	alias | dref|8-byte
fle.d	D_inx	d,S,T	–
fle.h	Zfh_inx	d,S,T	–
fle.q	q_inx	d,S,T	–

Instruction	Extension	Parameters	Flags
fle.s	f_inx	d,S,T	-
flh	Zfhmin	D,A,s	macro
flh	Zfhmin	D,o(s)	dref | 2-byte
flq	Q	D,A,s	macro
flq	Q	D,o(s)	dref | 16-byte
flt.d	D_inx	d,S,T	-
flt.h	Zfh_inx	d,S,T	-
flt.q	q_inx	d,S,T	-
flt.s	f_inx	d,S,T	-
flw	F	D,A,s	macro
flw	F	D,o(s)	dref | 4-byte
flw	F and C	CD,Ck(Cs)	alias | dref|4-byte
flw	F and C	D,Cm(Cc)	alias | dref|4-byte
fmadd.d	D_inx	D,S,T,R	-
fmadd.d	D_inx	D,S,T,R,m	-
fmadd.h	Zfh_inx	D,S,T,R	-
fmadd.h	Zfh_inx	D,S,T,R,m	-
fmadd.q	q_inx	D,S,T,R	-
fmadd.q	q_inx	D,S,T,R,m	-
fmadd.s	f_inx	D,S,T,R	-
fmadd.s	f_inx	D,S,T,R,m	-
fmax.d	D_inx	D,S,T	-
fmax.h	Zfh_inx	D,S,T	-
fmax.q	q_inx	D,S,T	-
fmax.s	f_inx	D,S,T	-
fmin.d	D_inx	D,S,T	-
fmin.h	Zfh_inx	D,S,T	-
fmin.q	q_inx	D,S,T	-
fmin.s	f_inx	D,S,T	-
fmsub.d	D_inx	D,S,T,R	-
fmsub.d	D_inx	D,S,T,R,m	-
fmsub.h	Zfh_inx	D,S,T,R	-
fmsub.h	Zfh_inx	D,S,T,R,m	-
fmsub.q	q_inx	D,S,T,R	-
fmsub.q	q_inx	D,S,T,R,m	-
fmsub.s	f_inx	D,S,T,R	-
fmsub.s	f_inx	D,S,T,R,m	-
fmul.d	D_inx	D,S,T	-
fmul.d	D_inx	D,S,T,m	-
fmul.h	Zfh_inx	D,S,T	-
fmul.h	Zfh_inx	D,S,T,m	-
fmul.q	q_inx	D,S,T	-
fmul.q	q_inx	D,S,T,m	-
fmul.s	f_inx	D,S,T	-
fmul.s	f_inx	D,S,T,m	-
fmv.d	D_inx	D,U	alias
fmv.d.x	D	D,s	-
fmv.h	Zfh_inx	D,U	alias
fmv.h.x	Zfhmin	D,s	-
fmv.q	q_inx	D,U	alias
fmv.s	f_inx	D,U	alias
fmv.s.x	F	D,s	-
fmv.w.x	F	D,s	-
fmv.x.d	D	d,S	-

Instruction	Extension	Parameters	Flags
fmv.x.h	Zfhmin	d,S	−
fmv.x.s	F	d,S	−
fmv.x.w	F	d,S	−
fneg.d	D_inx	D,U	alias
fneg.h	Zfh_inx	D,U	alias
fneg.q	q_inx	D,U	alias
fneg.s	f_inx	D,U	alias
fnmadd.d	D_inx	D,S,T,R	−
fnmadd.d	D_inx	D,S,T,R,m	−
fnmadd.h	Zfh_inx	D,S,T,R	−
fnmadd.h	Zfh_inx	D,S,T,R,m	−
fnmadd.q	q_inx	D,S,T,R	−
fnmadd.q	q_inx	D,S,T,R,m	−
fnmadd.s	f_inx	D,S,T,R	−
fnmadd.s	f_inx	D,S,T,R,m	−
fnmsub.d	D_inx	D,S,T,R	−
fnmsub.d	D_inx	D,S,T,R,m	−
fnmsub.h	Zfh_inx	D,S,T,R	−
fnmsub.h	Zfh_inx	D,S,T,R,m	−
fnmsub.q	q_inx	D,S,T,R	−
fnmsub.q	q_inx	D,S,T,R,m	−
fnmsub.s	f_inx	D,S,T,R	−
fnmsub.s	f_inx	D,S,T,R,m	−
frcsr	f_inx	d	alias
frflags	f_inx	d	alias
frrm	f_inx	d	alias
frsr	f_inx	d	alias
fscsr	f_inx	d,s	alias
fscsr	f_inx	s	alias
fsd	D	T,A,s	macro
fsd	D	T,q(s)	dref | 8-byte
fsd	D and C	CD,Cl(Cs)	alias | dref|8-byte
fsd	D and C	CT,CN(Cc)	alias | dref|8-byte
fsflags	f_inx	d,s	alias
fsflags	f_inx	s	alias
fsflagsi	f_inx	Z	alias
fsflagsi	f_inx	d,Z	alias
fsgnj.d	D_inx	D,S,T	−
fsgnj.h	Zfh_inx	D,S,T	−
fsgnj.q	q_inx	D,S,T	−
fsgnj.s	f_inx	D,S,T	−
fsgnjn.d	D_inx	D,S,T	−
fsgnjn.h	Zfh_inx	D,S,T	−
fsgnjn.q	q_inx	D,S,T	−
fsgnjn.s	f_inx	D,S,T	−
fsgnjx.d	D_inx	D,S,T	−
fsgnjx.h	Zfh_inx	D,S,T	−
fsgnjx.q	q_inx	D,S,T	−
fsgnjx.s	f_inx	D,S,T	−
fsh	Zfhmin	T,A,s	macro
fsh	Zfhmin	T,q(s)	dref | 2-byte
fsq	Q	T,A,s	macro
fsq	Q	T,q(s)	dref | 16-byte
fsqrt.d	D_inx	D,S	−

Instruction	Extension	Parameters	Flags
fsqrt.d	D_inx	D,S,m	–
fsqrt.h	Zfh_inx	D,S	–
fsqrt.h	Zfh_inx	D,S,m	–
fsqrt.q	q_inx	D,S	–
fsqrt.q	q_inx	D,S,m	–
fsqrt.s	f_inx	D,S	–
fsqrt.s	f_inx	D,S,m	–
fsrm	f_inx	d,s	alias
fsrm	f_inx	s	alias
fsrmi	f_inx	Z	alias
fsrmi	f_inx	d,Z	alias
fssr	f_inx	d,s	alias
fssr	f_inx	s	alias
fsub.d	D_inx	D,S,T	–
fsub.d	D_inx	D,S,T,m	–
fsub.h	Zfh_inx	D,S,T	–
fsub.h	Zfh_inx	D,S,T,m	–
fsub.q	q_inx	D,S,T	–
fsub.q	q_inx	D,S,T,m	–
fsub.s	f_inx	D,S,T	–
fsub.s	f_inx	D,S,T,m	–
fsw	F	T,A,s	macro
fsw	F	T,q(s)	dref | 4-byte
fsw	F and C	CD,Ck(Cs)	alias | dref|4-byte
fsw	F and C	CT,CM(Cc)	alias | dref|4-byte
hfence.gvma	H	--	alias
hfence.gvma	H	s	alias
hfence.gvma	H	s,t	–
hfence.vvma	H	--	alias
hfence.vvma	H	s	alias
hfence.vvma	H	s,t	–
hinval.gvma	SVINVAL	s,t	–
hinval.vvma	SVINVAL	s,t	–
hlv.b	H	d,0(s)	dref | 1-byte
hlv.bu	H	d,0(s)	dref | 1-byte
hlv.d	H	d,0(s)	dref | 8-byte
hlv.h	H	d,0(s)	dref | 2-byte
hlv.hu	H	d,0(s)	dref | 2-byte
hlv.w	H	d,0(s)	dref | 4-byte
hlv.wu	H	d,0(s)	dref | 4-byte
hlvx.hu	H	d,0(s)	dref | 2-byte
hlvx.wu	H	d,0(s)	dref | 4-byte
hret	I	--	–
hsv.b	H	t,0(s)	dref | 1-byte
hsv.d	H	t,0(s)	dref | 8-byte
hsv.h	H	t,0(s)	dref | 2-byte
hsv.w	H	t,0(s)	dref | 4-byte
j	C	Ca	alias | branch
j	I	a	alias | branch
jal	I	a	alias | jsr
jal	I	d,a	jsr
jal	C	Ca	alias | jsr
jalr	C	d	alias | jsr
jalr	I	d,o(s)	jsr

Instruction	Extension	Parameters	Flags
jalr	I	d,s	alias | jsr
jalr	I	d,s,j	jsr
jalr	I	o(s)	alias | jsr
jalr	I	s	alias | jsr
jalr	I	s,j	alias | jsr
jr	C	d	alias | branch
jr	I	o(s)	alias | branch
jr	I	s	alias | branch
jr	I	s,j	alias | branch
jump	I	c,s	macro
la	I	d,B	macro
la.tls.gd	I	d,A	macro
la.tls.ie	I	d,A	macro
lb	I	d,A	macro
lb	I	d,o(s)	dref | 1-byte
lbu	I	d,A	macro
lbu	I	d,o(s)	dref | 1-byte
ld	C	Ct,Cl(Cs)	alias | dref|8-byte
ld	C	d,Cn(Cc)	alias | dref|8-byte
ld	I	d,A	macro
ld	I	d,o(s)	dref | 8-byte
lh	I	d,A	macro
lh	I	d,o(s)	dref | 2-byte
lhu	I	d,A	macro
lhu	I	d,o(s)	dref | 2-byte
li	C	d,Co	alias
li	C	d,Cv	alias
li	I	d,I	macro
li	I	d,j	alias
lla	I	d,B	macro
lr.d	A	d,0(s)	dref | 8-byte
lr.d.aq	A	d,0(s)	dref | 8-byte
lr.d.aqrl	A	d,0(s)	dref | 8-byte
lr.d.rl	A	d,0(s)	dref | 8-byte
lr.w	A	d,0(s)	dref | 4-byte
lr.w.aq	A	d,0(s)	dref | 4-byte
lr.w.aqrl	A	d,0(s)	dref | 4-byte
lr.w.rl	A	d,0(s)	dref | 4-byte
lui	C	d,Cu	
lui	I	d,u	
lw	C	Ct,Ck(Cs)	alias | dref|4-byte
lw	C	d,Cm(Cc)	alias | dref|4-byte
lw	I	d,A	macro
lw	I	d,o(s)	dref | 4-byte
lwu	I	d,A	macro
lwu	I	d,o(s)	dref | 4-byte
max	Zbb	d,s,t	–
maxu	Zbb	d,s,t	–
min	Zbb	d,s,t	–
minu	Zbb	d,s,t	–
move	C	d,CV	alias
move	I	d,s	alias
mret	I	--	–
mul	Zmmul	d,s,t	–

Instruction	Extension	Parameters	Flags
mulh	Zmmul	d,s,t	–
mulhsu	Zmmul	d,s,t	–
mulhu	Zmmul	d,s,t	–
mulw	Zmmul	d,s,t	–
mv	C	d,CV	alias
mv	I	d,s	alias
neg	I	d,t	alias
negw	I	d,t	alias
nop	C	--	alias
nop	I	--	alias
not	I	d,s	alias
or	C	Cs,Ct,Cw	alias
or	C	Cs,Cw,Ct	alias
or	I	d,s,j	alias
or	I	d,s,t	–
orc.b	Zbb	d,s	–
ori	I	d,s,j	–
orn	Zbb | Zbkb	d,s,t	–
pack	Zbkb	d,s,t	crypto
packh	Zbkb	d,s,t	crypto
packw	Zbkb	d,s,t	crypto
pause	Zihintpause	--	–
prefetch.i	Zicbop	Wif(s)	–
prefetch.r	Zicbop	Wif(s)	–
prefetch.w	Zicbop	Wif(s)	–
rdcycle	I	d	alias
rdcycleh	I	d	alias
rdinstret	I	d	alias
rdinstreth	I	d	alias
rdtime	I	d	alias
rdtimeh	I	d	alias
rem	M	d,s,t	–
remu	M	d,s,t	–
remuw	M	d,s,t	–
remw	M	d,s,t	–
ret	C	--	alias | branch
ret	I	--	alias | branch
rev8	Zbb | Zbkb	d,s	–
rev8	Zbb | Zbkb	d,s	–
rol	Zbb | Zbkb	d,s,t	–
rolw	Zbb | Zbkb	d,s,t	–
ror	Zbb | Zbkb	d,s,>	alias
ror	Zbb | Zbkb	d,s,t	–
rori	Zbb | Zbkb	d,s,>	–
roriw	Zbb | Zbkb	d,s,<	–

Instruction	Extension	Parameters	Flags
rorw	Zbb | Zbkb	d,s,<	alias
rorw	Zbb | Zbkb	d,s,t	–
sb	I	t,A,s	macro
sb	I	t,q(s)	dref | 1-byte
sbreak	C	--	alias
sbreak	I	--	alias
sc.d	A	d,t,0(s)	dref | 8-byte
sc.d.aq	A	d,t,0(s)	dref | 8-byte
sc.d.aqrl	A	d,t,0(s)	dref | 8-byte
sc.d.rl	A	d,t,0(s)	dref | 8-byte
sc.w	A	d,t,0(s)	dref | 4-byte
sc.w.aq	A	d,t,0(s)	dref | 4-byte
sc.w.aqrl	A	d,t,0(s)	dref | 4-byte
sc.w.rl	A	d,t,0(s)	dref | 4-byte
scall	I	--	–
sd	C	CV,CN(Cc)	alias | dref|8-byte
sd	C	Ct,Cl(Cs)	alias | dref|8-byte
sd	I	t,A,s	macro
sd	I	t,q(s)	dref | 8-byte
seqz	I	d,s	alias
sext.b	I	d,s	macro
sext.b	Zbb	d,s	–
sext.h	I	d,s	macro
sext.h	Zbb	d,s	–
sext.w	C	d,CU	alias
sext.w	I	d,s	alias
sfence.inval.ir	SVINVAL	--	–
sfence.vm	I	--	–
sfence.vm	I	s	–
sfence.vma	I	--	alias
sfence.vma	I	s	alias
sfence.vma	I	s,t	–
sfence.w.inval	SVINVAL	--	–
sgt	I	d,t,s	alias
sgtu	I	d,t,s	alias
sgtz	I	d,t	alias
sh	I	t,A,s	macro
sh	I	t,q(s)	dref | 2-byte
sh1add	Zba	d,s,t	–
sh1add.uw	Zba	d,s,t	–
sh2add	Zba	d,s,t	–
sh2add.uw	Zba	d,s,t	–
sh3add	Zba	d,s,t	–
sh3add.uw	Zba	d,s,t	–
sha256sig0	Zknh	d,s	crypto
sha256sig1	Zknh	d,s	crypto
sha256sum0	Zknh	d,s	crypto
sha256sum1	Zknh	d,s	crypto
sha512sig0	Zknh	d,s	crypto
sha512sig0h	Zknh	d,s,t	crypto
sha512sig0l	Zknh	d,s,t	crypto
sha512sig1	Zknh	d,s	crypto

Instruction	Extension	Parameters	Flags
sha512sig1h	Zknh	d,s,t	crypto
sha512sig1l	Zknh	d,s,t	crypto
sha512sum0	Zknh	d,s	crypto
sha512sum0r	Zknh	d,s,t	crypto
sha512sum1	Zknh	d,s	crypto
sha512sum1r	Zknh	d,s,t	crypto
sinval.vma	SVINVAL	s,t	crypto
sll	C	d,CU,C>	alias
sll	I	d,s,>	alias
sll	I	d,s,t	–
slli	C	d,CU,C>	alias
slli	I	d,s,>	–
slli.uw	Zba	d,s,>	–
slliw	I	d,s,<	–
sllw	I	d,s,<	alias
sllw	I	d,s,t	–
slt	I	d,s,j	alias
slt	I	d,s,t	–
slti	I	d,s,j	–
sltiu	I	d,s,j	–
sltu	I	d,s,j	alias
sltu	I	d,s,t	–
sltz	I	d,s	alias
sm3p0	Zksh	d,s	–
sm3p1	Zksh	d,s	–
sm4ed	Zksed	d,s,t,y	–
sm4ks	Zksed	d,s,t,y	–
snez	I	d,t	alias
sra	C	Cs,Cw,C>	alias
sra	I	d,s,>	alias
sra	I	d,s,t	–
srai	C	Cs,Cw,C>	alias
srai	I	d,s,>	–
sraiw	I	d,s,<	–
sraw	I	d,s,<	alias
sraw	I	d,s,t	–
sret	I	--	–
srl	C	Cs,Cw,C>	alias
srl	I	d,s,>	alias
srl	I	d,s,t	–
srli	C	Cs,Cw,C>	alias
srli	I	d,s,>	–
srliw	I	d,s,<	–
srlw	I	d,s,<	alias
srlw	I	d,s,t	–
sub	C	Cs,Cw,Ct	alias
sub	I	d,s,t	–
subw	C	Cs,Cw,Ct	alias
subw	I	d,s,t	–
sw	C	CV,CM(Cc)	alias | dref|4-byte
sw	C	Ct,Ck(Cs)	alias | dref|4-byte
sw	I	t,A,s	macro
sw	I	t,q(s)	dref | 4-byte
tail	I	c	macro

Instruction	Extension	Parameters	Flags
th.addsl	th-ba	d,s,t,Xu2@25	–
th.dcache.call	th-cmo	--	–
th.dcache.ciall	th-cmo	--	–
th.dcache.cipa	th-cmo	s	–
th.dcache.cisw	th-cmo	s	–
th.dcache.civa	th-cmo	s	–
th.dcache.cpa	th-cmo	s	–
th.dcache.cpal1	th-cmo	s	–
th.dcache.csw	th-cmo	s	–
th.dcache.cva	th-cmo	s	–
th.dcache.cval1	th-cmo	s	–
th.dcache.iall	th-cmo	--	–
th.dcache.ipa	th-cmo	s	–
th.dcache.isw	th-cmo	s	–
th.dcache.iva	th-cmo	s	–
th.ext	th-bb	d,s,Xu6@26,Xu6@2a0	–
th.extu	th-bb	d,s,Xu6@26,Xu6@20	–
th.ff0	th-bb	d,s	–
th.ff1	th-bb	d,s	–
th.flrd	th-fmmemidx	D,s,t,Xu2@25	–
th.flrw	th-fmmemidx	D,s,t,Xu2@25	–
th.flurd	th-fmmemidx	D,s,t,Xu2@25	–
th.flurw	th-fmmemidx	D,s,t,Xu2@25	–
th.fmv.hw.x	th-fmv	d,S	–
th.fmv.x.hw	th-fmv	d,S	–
th.fsrd	th-fmmemidx	D,s,t,Xu2@25	–
th.fsrw	th-fmmemidx	D,s,t,Xu2@25	–
th.fsurd	th-fmmemidx	D,s,t,Xu2@25	–
th.fsurw	th-fmmemidx	D,s,t,Xu2@25	–
th.icache.iall	th-cmo	--	–
th.icache.ialls	th-cmo	--	–
th.icache.ipa	th-cmo	s	–
th.icache.iva	th-cmo	s	–
th.ipop	th-int	--	–
th.ipush	th-int	--	–
th.l2cache.call	th-cmo	--	–
th.l2cache.cialt	th-cmo	--	–
th.l2cache.iall	th-cmo	--	–
th.lbia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lbib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lbuia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lbuib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.ldd	th-mempair	d,t,(s),Xu2@25,Xl4	–
th.ldia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.ldib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lhia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lhib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lhuia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lhuib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lrb	th-memidx	d,s,t,Xu2@25	–
th.lrbu	th-memidx	d,s,t,Xu2@25	–
th.lrd	th-memidx	d,s,t,Xu2@25	–
th.lrh	th-memidx	d,s,t,Xu2@25	–
th.lrhu	th-memidx	d,s,t,Xu2@25	–

Instruction	Extension	Parameters	Flags
th.lrw	th-memidx	d,s,t,Xu2@25	–
th.lrwu	th-memidx	d,s,t,Xu2@25	–
th.lurb	th-memidx	d,s,t,Xu2@25	–
th.lurbu	th-memidx	d,s,t,Xu2@25	–
th.lurd	th-memidx	d,s,t,Xu2@25	–
th.lurh	th-memidx	d,s,t,Xu2@25	–
th.lurhu	th-memidx	d,s,t,Xu2@25	–
th.lurw	th-memidx	d,s,t,Xu2@25	–
th.lurwu	th-memidx	d,s,t,Xu2@25	–
th.lwd	th-mempair	d,t,(s),Xu2@25,X13	–
th.lwia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lwib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lwud	th-mempair	d,t,(s),Xu2@25,X13	–
th.lwuia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.lwuib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.mula	th-mac	d,s,t	–
th.mulah	th-mac	d,s,t	–
th.mulaw	th-mac	d,s,t	–
th.muls	th-mac	d,s,t	–
th.mulsh	th-mac	d,s,t	–
th.mulsw	th-mac	d,s,t	–
th.mveqz	th-condmov	d,s,t	–
th.mvnez	th-condmov	d,s,t	–
th.rev	th-bb	d,s	–
th.revw	th-bb	d,s	–
th.sbia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.sbib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.sdd	th-mempair	d,t,(s),Xu2@25,X14	–
th.sdia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.sdib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.sfence.vmas	th-sync	s,t	–
th.shia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.shib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.srb	th-memidx	d,s,t,Xu2@25	–
th.srd	th-memidx	d,s,t,Xu2@25	–
th.srh	th-memidx	d,s,t,Xu2@25	–
th.srri	th-bb	d,s,Xu6@20	–
th.srriw	th-bb	d,s,Xu5@20	–
th.srw	th-memidx	d,s,t,Xu2@25	–
th.surb	th-memidx	d,s,t,Xu2@25	–
th.surd	th-memidx	d,s,t,Xu2@25	–
th.surh	th-memidx	d,s,t,Xu2@25	–
th.surw	th-memidx	d,s,t,Xu2@25	–
th.swd	th-mempair	d,t,(s),Xu2@25,X13	–
th.swia	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.swib	th-memidx	d,(s),Xs5@20,Xu2@25	–
th.sync	th-sync	--	–
th.sync.i	th-sync	--	–
th.sync.is	th-sync	--	–
th.sync.s	th-sync	--	–
th.tst	th-bs	d,s,Xu6@20	–
th.tstnbz	th-bb	d,s	–
unimp	C	--	alias
unimp	I	--	x0

Instruction	Extension	Parameters	Flags
unzip	Zbkb	d,s	–
uret	I	--	–
vaadd.vv	V	Vd,Vt,VsVm	–
vaadd.vx	V	Vd,Vt,sVm	–
vaaddu.vv	V	Vd,Vt,VsVm	–
vaaddu.vx	V	Vd,Vt,sVm	–
vadc.vim	V	Vd,Vt,Vi,V0	–
vadc.vvm	V	Vd,Vt,Vs,V0	–
vadc.vxm	V	Vd,Vt,s,V0	–
vadd.vi	V	Vd,Vt,ViVm	–
vadd.vv	V	Vd,Vt,VsVm	–
vadd.vx	V	Vd,Vt,sVm	–
vand.vi	V	Vd,Vt,ViVm	–
vand.vv	V	Vd,Vt,VsVm	–
vand.vx	V	Vd,Vt,sVm	–
vasub.vv	V	Vd,Vt,VsVm	–
vasub.vx	V	Vd,Vt,sVm	–
vasubu.vv	V	Vd,Vt,VsVm	–
vasubu.vx	V	Vd,Vt,sVm	–
vcompress.vm	V	Vd,Vt,Vs	–
vcpop.m	V	d,VtVm	–
vdiv.vv	V	Vd,Vt,VsVm	–
vdiv.vx	V	Vd,Vt,sVm	–
vdivu.vv	V	Vd,Vt,VsVm	–
vdivu.vx	V	Vd,Vt,sVm	–
vfabs.v	Zvef	Vd,VuVm	alias
vfadd.vf	Zvef	Vd,Vt,SVm	–
vfadd.vv	Zvef	Vd,Vt,VsVm	–
vfclass.v	Zvef	Vd,VtVm	–
vfcvt.f.x.v	Zvef	Vd,VtVm	–
vfcvt.f.xu.v	Zvef	Vd,VtVm	–
vfcvt.rtz.x.f.v	Zvef	Vd,VtVm	–
vfcvt.rtz.xu.f.v	Zvef	Vd,VtVm	–
vfcvt.x.f.v	Zvef	Vd,VtVm	–
vfcvt.xu.f.v	Zvef	Vd,VtVm	–
vfdiv.vf	Zvef	Vd,Vt,SVm	–
vfdiv.vv	Zvef	Vd,Vt,VsVm	–
vfirst.m	V	d,VtVm	–
vfmacc.vf	Zvef	Vd,S,VtVm	–
vfmacc.vv	Zvef	Vd,Vs,VtVm	–
vfmadd.vf	Zvef	Vd,S,VtVm	–
vfmadd.vv	Zvef	Vd,Vs,VtVm	–
vfmax.vf	Zvef	Vd,Vt,SVm	–
vfmax.vv	Zvef	Vd,Vt,VsVm	–
vfmerge.vfm	Zvef	Vd,Vt,S,V0	–
vfmin.vf	Zvef	Vd,Vt,SVm	–
vfmin.vv	Zvef	Vd,Vt,VsVm	–
vfmsac.vf	Zvef	Vd,S,VtVm	–
vfmsac.vv	Zvef	Vd,Vs,VtVm	–
vfmsub.vf	Zvef	Vd,S,VtVm	–
vfmsub.vv	Zvef	Vd,Vs,VtVm	–
vfmul.vf	Zvef	Vd,Vt,SVm	–
vfmul.vv	Zvef	Vd,Vt,VsVm	–
vfmv.f.s	Zvef	D,Vt	–

Instruction	Extension	Parameters	Flags
vfmv.s.f	Zvef	Vd,S	–
vfmv.v.f	Zvef	Vd,S	–
vfncvt.f.f.w	Zvef	Vd,VtVm	–
vfncvt.f.x.w	Zvef	Vd,VtVm	–
vfncvt.f.xu.w	Zvef	Vd,VtVm	–
vfncvt.rod.f.f.w	Zvef	Vd,VtVm	–
vfncvt.rtz.x.f.w	Zvef	Vd,VtVm	–
vfncvt.rtz.xu.f.w	Zvef	Vd,VtVm	–
vfncvt.x.f.w	Zvef	Vd,VtVm	–
vfncvt.xu.f.w	Zvef	Vd,VtVm	–
vfneg.v	Zvef	Vd,VuVm	alias
vfnmacc.vf	Zvef	Vd,S,VtVm	–
vfnmacc.vv	Zvef	Vd,Vs,VtVm	–
vfnmadd.vf	Zvef	Vd,S,VtVm	–
vfnmadd.vv	Zvef	Vd,Vs,VtVm	–
vfnmsac.vf	Zvef	Vd,S,VtVm	–
vfnmsac.vv	Zvef	Vd,Vs,VtVm	–
vfnmsub.vf	Zvef	Vd,S,VtVm	–
vfnmsub.vv	Zvef	Vd,Vs,VtVm	–
vfrdiv.vf	Zvef	Vd,Vt,SVm	–
vfrec7.v	Zvef	Vd,VtVm	–
vfrece7.v	Zvef	Vd,VtVm	–
vfredmax.vs	Zvef	Vd,Vt,VsVm	–
vfredmin.vs	Zvef	Vd,Vt,VsVm	–
vfredosum.vs	Zvef	Vd,Vt,VsVm	–
vfredsum.vs	Zvef	Vd,Vt,VsVm	alias
vfredusum.vs	Zvef	Vd,Vt,VsVm	–
vfrsqrt7.v	Zvef	Vd,VtVm	–
vfrsqrte7.v	Zvef	Vd,VtVm	–
vfrsub.vf	Zvef	Vd,Vt,SVm	–
vfsgnj.vf	Zvef	Vd,Vt,SVm	–
vfsgnj.vv	Zvef	Vd,Vt,VsVm	–
vfsgnjn.vf	Zvef	Vd,Vt,SVm	–
vfsgnjn.vv	Zvef	Vd,Vt,VsVm	–
vfsgnjx.vf	Zvef	Vd,Vt,SVm	–
vfsgnjx.vv	Zvef	Vd,Vt,VsVm	–
vfslide1down.vf	Zvef	Vd,Vt,SVm	–
vfslide1up.vf	Zvef	Vd,Vt,SVm	–
vfsqrt.v	Zvef	Vd,VtVm	–
vfsub.vf	Zvef	Vd,Vt,SVm	–
vfsub.vv	Zvef	Vd,Vt,VsVm	–
vfwadd.vf	Zvef	Vd,Vt,SVm	–
vfwadd.vv	Zvef	Vd,Vt,VsVm	–
vfwadd.wf	Zvef	Vd,Vt,SVm	–
vfwadd.wv	Zvef	Vd,Vt,VsVm	–
vfwcvt.f.f.v	Zvef	Vd,VtVm	–
vfwcvt.f.x.v	Zvef	Vd,VtVm	–
vfwcvt.f.xu.v	Zvef	Vd,VtVm	–
vfwcvt.rtz.x.f.v	Zvef	Vd,VtVm	–
vfwcvt.rtz.xu.f.v	Zvef	Vd,VtVm	–
vfwcvt.x.f.v	Zvef	Vd,VtVm	–
vfwcvt.xu.f.v	Zvef	Vd,VtVm	–
vfwmacc.vf	Zvef	Vd,S,VtVm	–
vfwmacc.vv	Zvef	Vd,Vs,VtVm	–

Instruction	Extension	Parameters	Flags
vfwmsac.vf	Zvef	Vd,S,VtVm	–
vfwmsac.vv	Zvef	Vd,Vs,VtVm	–
vfwmul.vf	Zvef	Vd,Vt,SVm	–
vfwmul.vv	Zvef	Vd,Vt,VsVm	–
vfwnmacc.vf	Zvef	Vd,S,VtVm	–
vfwnmacc.vv	Zvef	Vd,Vs,VtVm	–
vfwnmsac.vf	Zvef	Vd,S,VtVm	–
vfwnmsac.vv	Zvef	Vd,Vs,VtVm	–
vfwredosum.vs	Zvef	Vd,Vt,VsVm	–
vfwredsum.vs	Zvef	Vd,Vt,VsVm	alias
vfwredusum.vs	Zvef	Vd,Vt,VsVm	–
vfwsub.vf	Zvef	Vd,Vt,SVm	–
vfwsub.vv	Zvef	Vd,Vt,VsVm	–
vfwsub.wf	Zvef	Vd,Vt,SVm	–
vfwsub.wv	Zvef	Vd,Vt,VsVm	–
vid.v	V	VdVm	–
viota.m	V	Vd,VtVm	–
vl1r.v	V	Vd,0(s)	dref | alias
vl1re16.v	V	Vd,0(s)	dref
vl1re32.v	V	Vd,0(s)	dref
vl1re64.v	V	Vd,0(s)	dref | eew64
vl1re8.v	V	Vd,0(s)	dref
vl2r.v	V	Vd,0(s)	dref | alias
vl2re16.v	V	Vd,0(s)	dref
vl2re32.v	V	Vd,0(s)	dref
vl2re64.v	V	Vd,0(s)	dref | eew64
vl2re8.v	V	Vd,0(s)	dref
vl4r.v	V	Vd,0(s)	dref | alias
vl4re16.v	V	Vd,0(s)	dref
vl4re32.v	V	Vd,0(s)	dref
vl4re64.v	V	Vd,0(s)	dref | eew64
vl4re8.v	V	Vd,0(s)	dref
vl8r.v	V	Vd,0(s)	dref | alias
vl8re16.v	V	Vd,0(s)	dref
vl8re32.v	V	Vd,0(s)	dref
vl8re64.v	V	Vd,0(s)	dref | eew64
vl8re8.v	V	Vd,0(s)	dref
vle1.v	V	Vd,0(s)	dref | alias
vle16.v	V	Vd,0(s)Vm	dref
vle16ff.v	V	Vd,0(s)Vm	dref
vle32.v	V	Vd,0(s)Vm	dref
vle32ff.v	V	Vd,0(s)Vm	dref
vle64.v	V	Vd,0(s)Vm	dref | eew64
vle64ff.v	V	Vd,0(s)Vm	dref | eew64
vle8.v	V	Vd,0(s)Vm	dref
vle8ff.v	V	Vd,0(s)Vm	dref
vlm.v	V	Vd,0(s)	dref
vloxei16.v	V	Vd,0(s),VtVm	dref
vloxei32.v	V	Vd,0(s),VtVm	dref
vloxei64.v	V	Vd,0(s),VtVm	dref | eew64
vloxei8.v	V	Vd,0(s),VtVm	dref
vloxseg2ei16.v	V	Vd,0(s),VtVm	dref
vloxseg2ei32.v	V	Vd,0(s),VtVm	dref
vloxseg2ei64.v	V	Vd,0(s),VtVm	dref | eew64

Instruction	Extension	Parameters	Flags
vloxseg2ei8.v	V	Vd,0(s),VtVm	dref
vloxseg3ei16.v	V	Vd,0(s),VtVm	dref
vloxseg3ei32.v	V	Vd,0(s),VtVm	dref
vloxseg3ei64.v	V	Vd,0(s),VtVm	dref | eew64
vloxseg3ei8.v	V	Vd,0(s),VtVm	dref
vloxseg4ei16.v	V	Vd,0(s),VtVm	dref
vloxseg4ei32.v	V	Vd,0(s),VtVm	dref
vloxseg4ei64.v	V	Vd,0(s),VtVm	dref | eew64
vloxseg4ei8.v	V	Vd,0(s),VtVm	dref
vloxseg5ei16.v	V	Vd,0(s),VtVm	dref
vloxseg5ei32.v	V	Vd,0(s),VtVm	dref
vloxseg5ei64.v	V	Vd,0(s),VtVm	dref | eew64
vloxseg5ei8.v	V	Vd,0(s),VtVm	dref
vloxseg6ei16.v	V	Vd,0(s),VtVm	dref
vloxseg6ei32.v	V	Vd,0(s),VtVm	dref
vloxseg6ei64.v	V	Vd,0(s),VtVm	dref | eew64
vloxseg6ei8.v	V	Vd,0(s),VtVm	dref
vloxseg7ei16.v	V	Vd,0(s),VtVm	dref
vloxseg7ei32.v	V	Vd,0(s),VtVm	dref
vloxseg7ei64.v	V	Vd,0(s),VtVm	dref | eew64
vloxseg7ei8.v	V	Vd,0(s),VtVm	dref
vloxseg8ei16.v	V	Vd,0(s),VtVm	dref
vloxseg8ei32.v	V	Vd,0(s),VtVm	dref
vloxseg8ei64.v	V	Vd,0(s),VtVm	dref | eew64
vloxseg8ei8.v	V	Vd,0(s),VtVm	dref
vlse16.v	V	Vd,0(s),tVm	dref
vlse32.v	V	Vd,0(s),tVm	dref
vlse64.v	V	Vd,0(s),tVm	dref | eew64
vlse8.v	V	Vd,0(s),tVm	dref
vlseg2e16.v	V	Vd,0(s)Vm	dref
vlseg2e16ff.v	V	Vd,0(s)Vm	dref
vlseg2e32.v	V	Vd,0(s)Vm	dref
vlseg2e32ff.v	V	Vd,0(s)Vm	dref
vlseg2e64.v	V	Vd,0(s)Vm	dref | eew64
vlseg2e64ff.v	V	Vd,0(s)Vm	dref | eew64
vlseg2e8.v	V	Vd,0(s)Vm	dref
vlseg2e8ff.v	V	Vd,0(s)Vm	dref
vlseg3e16.v	V	Vd,0(s)Vm	dref
vlseg3e16ff.v	V	Vd,0(s)Vm	dref
vlseg3e32.v	V	Vd,0(s)Vm	dref
vlseg3e32ff.v	V	Vd,0(s)Vm	dref
vlseg3e64.v	V	Vd,0(s)Vm	dref | eew64
vlseg3e64ff.v	V	Vd,0(s)Vm	dref | eew64
vlseg3e8.v	V	Vd,0(s)Vm	dref
vlseg3e8ff.v	V	Vd,0(s)Vm	dref
vlseg4e16.v	V	Vd,0(s)Vm	dref
vlseg4e16ff.v	V	Vd,0(s)Vm	dref
vlseg4e32.v	V	Vd,0(s)Vm	dref
vlseg4e32ff.v	V	Vd,0(s)Vm	dref
vlseg4e64.v	V	Vd,0(s)Vm	dref | eew64
vlseg4e64ff.v	V	Vd,0(s)Vm	dref | eew64
vlseg4e8.v	V	Vd,0(s)Vm	dref
vlseg4e8ff.v	V	Vd,0(s)Vm	dref
vlseg5e16.v	V	Vd,0(s)Vm	dref

Instruction	Extension	Parameters	Flags
vlseg5e16ff.v	V	Vd,0(s)Vm	dref
vlseg5e32.v	V	Vd,0(s)Vm	dref
vlseg5e32ff.v	V	Vd,0(s)Vm	dref
vlseg5e64.v	V	Vd,0(s)Vm	dref | eew64
vlseg5e64ff.v	V	Vd,0(s)Vm	dref | eew64
vlseg5e8.v	V	Vd,0(s)Vm	dref
vlseg5e8ff.v	V	Vd,0(s)Vm	dref
vlseg6e16.v	V	Vd,0(s)Vm	dref
vlseg6e16ff.v	V	Vd,0(s)Vm	dref
vlseg6e32.v	V	Vd,0(s)Vm	dref
vlseg6e32ff.v	V	Vd,0(s)Vm	dref
vlseg6e64.v	V	Vd,0(s)Vm	dref | eew64
vlseg6e64ff.v	V	Vd,0(s)Vm	dref | eew64
vlseg6e8.v	V	Vd,0(s)Vm	dref
vlseg6e8ff.v	V	Vd,0(s)Vm	dref
vlseg7e16.v	V	Vd,0(s)Vm	dref
vlseg7e16ff.v	V	Vd,0(s)Vm	dref
vlseg7e32.v	V	Vd,0(s)Vm	dref
vlseg7e32ff.v	V	Vd,0(s)Vm	dref
vlseg7e64.v	V	Vd,0(s)Vm	dref | eew64
vlseg7e64ff.v	V	Vd,0(s)Vm	dref | eew64
vlseg7e8.v	V	Vd,0(s)Vm	dref
vlseg7e8ff.v	V	Vd,0(s)Vm	dref
vlseg8e16.v	V	Vd,0(s)Vm	dref
vlseg8e16ff.v	V	Vd,0(s)Vm	dref
vlseg8e32.v	V	Vd,0(s)Vm	dref
vlseg8e32ff.v	V	Vd,0(s)Vm	dref
vlseg8e64.v	V	Vd,0(s)Vm	dref | eew64
vlseg8e64ff.v	V	Vd,0(s)Vm	dref | eew64
vlseg8e8.v	V	Vd,0(s)Vm	dref
vlseg8e8ff.v	V	Vd,0(s)Vm	dref
vlsseg2e16.v	V	Vd,0(s),tVm	dref
vlsseg2e32.v	V	Vd,0(s),tVm	dref
vlsseg2e64.v	V	Vd,0(s),tVm	dref | eew64
vlsseg2e8.v	V	Vd,0(s),tVm	dref
vlsseg3e16.v	V	Vd,0(s),tVm	dref
vlsseg3e32.v	V	Vd,0(s),tVm	dref
vlsseg3e64.v	V	Vd,0(s),tVm	dref | eew64
vlsseg3e8.v	V	Vd,0(s),tVm	dref
vlsseg4e16.v	V	Vd,0(s),tVm	dref
vlsseg4e32.v	V	Vd,0(s),tVm	dref
vlsseg4e64.v	V	Vd,0(s),tVm	dref | eew64
vlsseg4e8.v	V	Vd,0(s),tVm	dref
vlsseg5e16.v	V	Vd,0(s),tVm	dref
vlsseg5e32.v	V	Vd,0(s),tVm	dref
vlsseg5e64.v	V	Vd,0(s),tVm	dref | eew64
vlsseg5e8.v	V	Vd,0(s),tVm	dref
vlsseg6e16.v	V	Vd,0(s),tVm	dref
vlsseg6e32.v	V	Vd,0(s),tVm	dref
vlsseg6e64.v	V	Vd,0(s),tVm	dref | eew64
vlsseg6e8.v	V	Vd,0(s),tVm	dref
vlsseg7e16.v	V	Vd,0(s),tVm	dref
vlsseg7e32.v	V	Vd,0(s),tVm	dref
vlsseg7e64.v	V	Vd,0(s),tVm	dref | eew64

Instruction	Extension	Parameters	Flags
vlsseg7e8.v	V	Vd,0(s),tVm	dref
vlsseg8e16.v	V	Vd,0(s),tVm	dref
vlsseg8e32.v	V	Vd,0(s),tVm	dref
vlsseg8e64.v	V	Vd,0(s),tVm	dref | eew64
vlsseg8e8.v	V	Vd,0(s),tVm	dref
vluxei16.v	V	Vd,0(s),VtVm	dref
vluxei32.v	V	Vd,0(s),VtVm	dref
vluxei64.v	V	Vd,0(s),VtVm	dref | eew64
vluxei8.v	V	Vd,0(s),VtVm	dref
vluxseg2ei16.v	V	Vd,0(s),VtVm	dref
vluxseg2ei32.v	V	Vd,0(s),VtVm	dref
vluxseg2ei64.v	V	Vd,0(s),VtVm	dref | eew64
vluxseg2ei8.v	V	Vd,0(s),VtVm	dref
vluxseg3ei16.v	V	Vd,0(s),VtVm	dref
vluxseg3ei32.v	V	Vd,0(s),VtVm	dref
vluxseg3ei64.v	V	Vd,0(s),VtVm	dref | eew64
vluxseg3ei8.v	V	Vd,0(s),VtVm	dref
vluxseg4ei16.v	V	Vd,0(s),VtVm	dref
vluxseg4ei32.v	V	Vd,0(s),VtVm	dref
vluxseg4ei64.v	V	Vd,0(s),VtVm	dref | eew64
vluxseg4ei8.v	V	Vd,0(s),VtVm	dref
vluxseg5ei16.v	V	Vd,0(s),VtVm	dref
vluxseg5ei32.v	V	Vd,0(s),VtVm	dref
vluxseg5ei64.v	V	Vd,0(s),VtVm	dref | eew64
vluxseg5ei8.v	V	Vd,0(s),VtVm	dref
vluxseg6ei16.v	V	Vd,0(s),VtVm	dref
vluxseg6ei32.v	V	Vd,0(s),VtVm	dref
vluxseg6ei64.v	V	Vd,0(s),VtVm	dref | eew64
vluxseg6ei8.v	V	Vd,0(s),VtVm	dref
vluxseg7ei16.v	V	Vd,0(s),VtVm	dref
vluxseg7ei32.v	V	Vd,0(s),VtVm	dref
vluxseg7ei64.v	V	Vd,0(s),VtVm	dref | eew64
vluxseg7ei8.v	V	Vd,0(s),VtVm	dref
vluxseg8ei16.v	V	Vd,0(s),VtVm	dref
vluxseg8ei32.v	V	Vd,0(s),VtVm	dref
vluxseg8ei64.v	V	Vd,0(s),VtVm	dref | eew64
vluxseg8ei8.v	V	Vd,0(s),VtVm	dref
vmacc.vv	V	Vd,Vs,VtVm	–
vmacc.vx	V	Vd,s,VtVm	–
vmadc.vi	V	Vd,Vt,Vi	–
vmadc.vim	V	Vd,Vt,Vi,V0	–
vmadc.vv	V	Vd,Vt,Vs	–
vmadc.vvm	V	Vd,Vt,Vs,V0	–
vmadc.vx	V	Vd,Vt,s	–
vmadc.vxm	V	Vd,Vt,s,V0	–
vmadd.vv	V	Vd,Vs,VtVm	–
vmadd.vx	V	Vd,s,VtVm	–
vmand.mm	V	Vd,Vt,Vs	–
vmandn.mm	V	Vd,Vt,Vs	–
vmandnot.mm	V	Vd,Vt,Vs	alias
vmax.vv	V	Vd,Vt,VsVm	–
vmax.vx	V	Vd,Vt,sVm	–
vmaxu.vv	V	Vd,Vt,VsVm	–
vmaxu.vx	V	Vd,Vt,sVm	–

Instruction	Extension	Parameters	Flags
vmclr.m	V	Vv	alias
vmcpy.m	V	Vd,Vu	alias
vmerge.vim	V	Vd,Vt,Vi,V0	–
vmerge.vvm	V	Vd,Vt,Vs,V0	–
vmerge.vxm	V	Vd,Vt,s,V0	–
vmfeq.vf	Zvef	Vd,Vt,SVm	–
vmfeq.vv	Zvef	Vd,Vt,VsVm	–
vmfge.vf	Zvef	Vd,Vt,SVm	–
vmfge.vv	Zvef	Vd,Vs,VtVm	alias
vmfgt.vf	Zvef	Vd,Vt,SVm	–
vmfgt.vv	Zvef	Vd,Vs,VtVm	alias
vmfle.vf	Zvef	Vd,Vt,SVm	–
vmfle.vv	Zvef	Vd,Vt,VsVm	–
vmflt.vf	Zvef	Vd,Vt,SVm	–
vmflt.vv	Zvef	Vd,Vt,VsVm	–
vmfne.vf	Zvef	Vd,Vt,SVm	–
vmfne.vv	Zvef	Vd,Vt,VsVm	–
vmin.vv	V	Vd,Vt,VsVm	–
vmin.vx	V	Vd,Vt,sVm	–
vminu.vv	V	Vd,Vt,VsVm	–
vminu.vx	V	Vd,Vt,sVm	–
vmmv.m	V	Vd,Vu	alias
vmnand.mm	V	Vd,Vt,Vs	–
vmnor.mm	V	Vd,Vt,Vs	–
vmnot.m	V	Vd,Vu	alias
vmor.mm	V	Vd,Vt,Vs	–
vmorn.mm	V	Vd,Vt,Vs	–
vmornot.mm	V	Vd,Vt,Vs	alias
vmsbc.vv	V	Vd,Vt,Vs	–
vmsbc.vvm	V	Vd,Vt,Vs,V0	–
vmsbc.vx	V	Vd,Vt,s	–
vmsbc.vxm	V	Vd,Vt,s,V0	–
vmsbf.m	V	Vd,VtVm	–
vmseq.vi	V	Vd,Vt,ViVm	–
vmseq.vv	V	Vd,Vt,VsVm	–
vmseq.vx	V	Vd,Vt,sVm	–
vmset.m	V	Vv	alias
vmsge.vi	V	Vd,Vt,VkVm	alias
vmsge.vv	V	Vd,Vs,VtVm	alias
vmsge.vx	V	Vd,Vt,s,VM,VT	macro
vmsge.vx	V	Vd,Vt,sVm	macro
vmsgeu.vi	V	Vd,Vt,VkVm	alias
vmsgeu.vi	V	Vd,Vu,0Vm	alias
vmsgeu.vv	V	Vd,Vs,VtVm	alias
vmsgeu.vx	V	Vd,Vt,s,VM,VT	macro
vmsgeu.vx	V	Vd,Vt,sVm	macro
vmsgt.vi	V	Vd,Vt,ViVm	–
vmsgt.vv	V	Vd,Vs,VtVm	alias
vmsgt.vx	V	Vd,Vt,sVm	–
vmsgtu.vi	V	Vd,Vt,ViVm	–
vmsgtu.vv	V	Vd,Vs,VtVm	alias
vmsgtu.vx	V	Vd,Vt,sVm	–
vmsif.m	V	Vd,VtVm	–
vmsle.vi	V	Vd,Vt,ViVm	–

Instruction	Extension	Parameters	Flags
vmsle.vv	V	Vd,Vt,VsVm	–
vmsle.vx	V	Vd,Vt,sVm	–
vmsleu.vi	V	Vd,Vt,ViVm	–
vmsleu.vv	V	Vd,Vt,VsVm	–
vmsleu.vx	V	Vd,Vt,sVm	–
vmslt.vi	V	Vd,Vt,VkVm	alias
vmslt.vv	V	Vd,Vt,VsVm	–
vmslt.vx	V	Vd,Vt,sVm	–
vmsltu.vi	V	Vd,Vt,VkVm	alias
vmsltu.vi	V	Vd,Vu,0Vm	alias
vmsltu.vv	V	Vd,Vt,VsVm	–
vmsltu.vx	V	Vd,Vt,sVm	–
vmsne.vi	V	Vd,Vt,ViVm	–
vmsne.vv	V	Vd,Vt,VsVm	–
vmsne.vx	V	Vd,Vt,sVm	–
vmsof.m	V	Vd,VtVm	–
vmul.vv	V	Vd,Vt,VsVm	–
vmul.vx	V	Vd,Vt,sVm	–
vmulh.vv	V	Vd,Vt,VsVm	–
vmulh.vx	V	Vd,Vt,sVm	–
vmulhsu.vv	V	Vd,Vt,VsVm	–
vmulhsu.vx	V	Vd,Vt,sVm	–
vmulhu.vv	V	Vd,Vt,VsVm	–
vmulhu.vx	V	Vd,Vt,sVm	–
vmv.s.x	V	Vd,s	–
vmv.v.i	V	Vd,Vi	–
vmv.v.v	V	Vd,Vs	–
vmv.v.x	V	Vd,s	–
vmv.x.s	V	d,Vt	–
vmv1r.v	V	Vd,Vt	–
vmv2r.v	V	Vd,Vt	–
vmv4r.v	V	Vd,Vt	–
vmv8r.v	V	Vd,Vt	–
vmxnor.mm	V	Vd,Vt,Vs	–
vmxor.mm	V	Vd,Vt,Vs	–
vnclip.wi	V	Vd,Vt,VjVm	–
vnclip.wv	V	Vd,Vt,VsVm	–
vnclip.wx	V	Vd,Vt,sVm	–
vnclipu.wi	V	Vd,Vt,VjVm	–
vnclipu.wv	V	Vd,Vt,VsVm	–
vnclipu.wx	V	Vd,Vt,sVm	–
vncvt.x.x.w	V	Vd,VtVm	alias
vneg.v	V	Vd,VtVm	alias
vnmsac.vv	V	Vd,Vs,VtVm	–
vnmsac.vx	V	Vd,s,VtVm	–
vnmsub.vv	V	Vd,Vs,VtVm	–
vnmsub.vx	V	Vd,s,VtVm	–
vnot.v	V	Vd,VtVm	alias
vnsra.wi	V	Vd,Vt,VjVm	–
vnsra.wv	V	Vd,Vt,VsVm	–
vnsra.wx	V	Vd,Vt,sVm	–
vnsrl.wi	V	Vd,Vt,VjVm	–
vnsrl.wv	V	Vd,Vt,VsVm	–
vnsrl.wx	V	Vd,Vt,sVm	–

Instruction	Extension	Parameters	Flags
vor.vi	V	Vd,Vt,ViVm	–
vor.vv	V	Vd,Vt,VsVm	–
vor.vx	V	Vd,Vt,sVm	–
vpopc.m	V	d,VtVm	alias
vredand.vs	V	Vd,Vt,VsVm	–
vredmax.vs	V	Vd,Vt,VsVm	–
vredmaxu.vs	V	Vd,Vt,VsVm	–
vredmin.vs	V	Vd,Vt,VsVm	–
vredminu.vs	V	Vd,Vt,VsVm	–
vredor.vs	V	Vd,Vt,VsVm	–
vredsum.vs	V	Vd,Vt,VsVm	–
vredxor.vs	V	Vd,Vt,VsVm	–
vrem.vv	V	Vd,Vt,VsVm	–
vrem.vx	V	Vd,Vt,sVm	–
vremu.vv	V	Vd,Vt,VsVm	–
vremu.vx	V	Vd,Vt,sVm	–
vrgather.vi	V	Vd,Vt,VjVm	–
vrgather.vv	V	Vd,Vt,VsVm	–
vrgather.vx	V	Vd,Vt,sVm	–
vrgatherei16.vv	V	Vd,Vt,VsVm	–
vrsub.vi	V	Vd,Vt,ViVm	–
vrsub.vx	V	Vd,Vt,sVm	–
vs1r.v	V	Vd,0(s)	dref
vs2r.v	V	Vd,0(s)	dref
vs4r.v	V	Vd,0(s)	dref
vs8r.v	V	Vd,0(s)	dref
vsadd.vi	V	Vd,Vt,ViVm	–
vsadd.vv	V	Vd,Vt,VsVm	–
vsadd.vx	V	Vd,Vt,sVm	–
vsaddu.vi	V	Vd,Vt,ViVm	–
vsaddu.vv	V	Vd,Vt,VsVm	–
vsaddu.vx	V	Vd,Vt,sVm	–
vsbc.vvm	V	Vd,Vt,Vs,V0	–
vsbc.vxm	V	Vd,Vt,s,V0	–
vse1.v	V	Vd,0(s)	dref | alias
vse16.v	V	Vd,0(s)Vm	dref
vse32.v	V	Vd,0(s)Vm	dref
vse64.v	V	Vd,0(s)Vm	dref | eew64
vse8.v	V	Vd,0(s)Vm	dref
vsetivli	V	d,Z,Vb	–
vsetvl	V	d,s,t	–
vsetvli	V	d,s,Vc	–
vsext.vf2	V	Vd,VtVm	–
vsext.vf4	V	Vd,VtVm	–
vsext.vf8	V	Vd,VtVm	–
vslide1down.vx	V	Vd,Vt,sVm	–
vslide1up.vx	V	Vd,Vt,sVm	–
vslidedown.vi	V	Vd,Vt,VjVm	–
vslidedown.vx	V	Vd,Vt,sVm	–
vslideup.vi	V	Vd,Vt,VjVm	–
vslideup.vx	V	Vd,Vt,sVm	–
vsll.vi	V	Vd,Vt,VjVm	–
vsll.vv	V	Vd,Vt,VsVm	–
vsll.vx	V	Vd,Vt,sVm	–

Instruction	Extension	Parameters	Flags
vsm.v	V	Vd,0(s)	dref
vsmul.vv	V	Vd,Vt,VsVm	−
vsmul.vx	V	Vd,Vt,sVm	−
vsoxei16.v	V	Vd,0(s),VtVm	dref
vsoxei32.v	V	Vd,0(s),VtVm	dref
vsoxei64.v	V	Vd,0(s),VtVm	dref | eew64
vsoxei8.v	V	Vd,0(s),VtVm	dref
vsoxseg2ei16.v	V	Vd,0(s),VtVm	dref
vsoxseg2ei32.v	V	Vd,0(s),VtVm	dref
vsoxseg2ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsoxseg2ei8.v	V	Vd,0(s),VtVm	dref
vsoxseg3ei16.v	V	Vd,0(s),VtVm	dref
vsoxseg3ei32.v	V	Vd,0(s),VtVm	dref
vsoxseg3ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsoxseg3ei8.v	V	Vd,0(s),VtVm	dref
vsoxseg4ei16.v	V	Vd,0(s),VtVm	dref
vsoxseg4ei32.v	V	Vd,0(s),VtVm	dref
vsoxseg4ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsoxseg4ei8.v	V	Vd,0(s),VtVm	dref
vsoxseg5ei16.v	V	Vd,0(s),VtVm	dref
vsoxseg5ei32.v	V	Vd,0(s),VtVm	dref
vsoxseg5ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsoxseg5ei8.v	V	Vd,0(s),VtVm	dref
vsoxseg6ei16.v	V	Vd,0(s),VtVm	dref
vsoxseg6ei32.v	V	Vd,0(s),VtVm	dref
vsoxseg6ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsoxseg6ei8.v	V	Vd,0(s),VtVm	dref
vsoxseg7ei16.v	V	Vd,0(s),VtVm	dref
vsoxseg7ei32.v	V	Vd,0(s),VtVm	dref
vsoxseg7ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsoxseg7ei8.v	V	Vd,0(s),VtVm	dref
vsoxseg8ei16.v	V	Vd,0(s),VtVm	dref
vsoxseg8ei32.v	V	Vd,0(s),VtVm	dref
vsoxseg8ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsoxseg8ei8.v	V	Vd,0(s),VtVm	dref
vsra.vi	V	Vd,Vt,VjVm	−
vsra.vv	V	Vd,Vt,VsVm	−
vsra.vx	V	Vd,Vt,sVm	−
vsrl.vi	V	Vd,Vt,VjVm	−
vsrl.vv	V	Vd,Vt,VsVm	−
vsrl.vx	V	Vd,Vt,sVm	−
vsse16.v	V	Vd,0(s),tVm	dref
vsse32.v	V	Vd,0(s),tVm	dref
vsse64.v	V	Vd,0(s),tVm	dref | eew64
vsse8.v	V	Vd,0(s),tVm	dref
vsseg2e16.v	V	Vd,0(s)Vm	dref
vsseg2e32.v	V	Vd,0(s)Vm	dref
vsseg2e64.v	V	Vd,0(s)Vm	dref | eew64
vsseg2e8.v	V	Vd,0(s)Vm	dref
vsseg3e16.v	V	Vd,0(s)Vm	dref
vsseg3e32.v	V	Vd,0(s)Vm	dref
vsseg3e64.v	V	Vd,0(s)Vm	dref | eew64
vsseg3e8.v	V	Vd,0(s)Vm	dref
vsseg4e16.v	V	Vd,0(s)Vm	dref

Instruction	Extension	Parameters	Flags
vsseg4e32.v	V	Vd,0(s)Vm	dref
vsseg4e64.v	V	Vd,0(s)Vm	dref | eew64
vsseg4e8.v	V	Vd,0(s)Vm	dref
vsseg5e16.v	V	Vd,0(s)Vm	dref
vsseg5e32.v	V	Vd,0(s)Vm	dref
vsseg5e64.v	V	Vd,0(s)Vm	dref | eew64
vsseg5e8.v	V	Vd,0(s)Vm	dref
vsseg6e16.v	V	Vd,0(s)Vm	dref
vsseg6e32.v	V	Vd,0(s)Vm	dref
vsseg6e64.v	V	Vd,0(s)Vm	dref | eew64
vsseg6e8.v	V	Vd,0(s)Vm	dref
vsseg7e16.v	V	Vd,0(s)Vm	dref
vsseg7e32.v	V	Vd,0(s)Vm	dref
vsseg7e64.v	V	Vd,0(s)Vm	dref | eew64
vsseg7e8.v	V	Vd,0(s)Vm	dref
vsseg8e16.v	V	Vd,0(s)Vm	dref
vsseg8e32.v	V	Vd,0(s)Vm	dref
vsseg8e64.v	V	Vd,0(s)Vm	dref | eew64
vsseg8e8.v	V	Vd,0(s)Vm	dref
vssra.vi	V	Vd,Vt,VjVm	−
vssra.vv	V	Vd,Vt,VsVm	−
vssra.vx	V	Vd,Vt,sVm	−
vssrl.vi	V	Vd,Vt,VjVm	−
vssrl.vv	V	Vd,Vt,VsVm	−
vssrl.vx	V	Vd,Vt,sVm	−
vssseg2e16.v	V	Vd,0(s),tVm	dref
vssseg2e32.v	V	Vd,0(s),tVm	dref
vssseg2e64.v	V	Vd,0(s),tVm	dref | eew64
vssseg2e8.v	V	Vd,0(s),tVm	dref
vssseg3e16.v	V	Vd,0(s),tVm	dref
vssseg3e32.v	V	Vd,0(s),tVm	dref
vssseg3e64.v	V	Vd,0(s),tVm	dref | eew64
vssseg3e8.v	V	Vd,0(s),tVm	dref
vssseg4e16.v	V	Vd,0(s),tVm	dref
vssseg4e32.v	V	Vd,0(s),tVm	dref
vssseg4e64.v	V	Vd,0(s),tVm	dref | eew64
vssseg4e8.v	V	Vd,0(s),tVm	dref
vssseg5e16.v	V	Vd,0(s),tVm	dref
vssseg5e32.v	V	Vd,0(s),tVm	dref
vssseg5e64.v	V	Vd,0(s),tVm	dref | eew64
vssseg5e8.v	V	Vd,0(s),tVm	dref
vssseg6e16.v	V	Vd,0(s),tVm	dref
vssseg6e32.v	V	Vd,0(s),tVm	dref
vssseg6e64.v	V	Vd,0(s),tVm	dref | eew64
vssseg6e8.v	V	Vd,0(s),tVm	dref
vssseg7e16.v	V	Vd,0(s),tVm	dref
vssseg7e32.v	V	Vd,0(s),tVm	dref
vssseg7e64.v	V	Vd,0(s),tVm	dref | eew64
vssseg7e8.v	V	Vd,0(s),tVm	dref
vssseg8e16.v	V	Vd,0(s),tVm	dref
vssseg8e32.v	V	Vd,0(s),tVm	dref
vssseg8e64.v	V	Vd,0(s),tVm	dref | eew64
vssseg8e8.v	V	Vd,0(s),tVm	dref
vssub.vv	V	Vd,Vt,VsVm	−

Instruction	Extension	Parameters	Flags
vssub.vx	V	Vd,Vt,sVm	−
vssubu.vv	V	Vd,Vt,VsVm	−
vssubu.vx	V	Vd,Vt,sVm	−
vsub.vv	V	Vd,Vt,VsVm	−
vsub.vx	V	Vd,Vt,sVm	−
vsuxei16.v	V	Vd,0(s),VtVm	dref
vsuxei32.v	V	Vd,0(s),VtVm	dref
vsuxei64.v	V	Vd,0(s),VtVm	dref | eew64
vsuxei8.v	V	Vd,0(s),VtVm	dref
vsuxseg2ei16.v	V	Vd,0(s),VtVm	dref
vsuxseg2ei32.v	V	Vd,0(s),VtVm	dref
vsuxseg2ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsuxseg2ei8.v	V	Vd,0(s),VtVm	dref
vsuxseg3ei16.v	V	Vd,0(s),VtVm	dref
vsuxseg3ei32.v	V	Vd,0(s),VtVm	dref
vsuxseg3ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsuxseg3ei8.v	V	Vd,0(s),VtVm	dref
vsuxseg4ei16.v	V	Vd,0(s),VtVm	dref
vsuxseg4ei32.v	V	Vd,0(s),VtVm	dref
vsuxseg4ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsuxseg4ei8.v	V	Vd,0(s),VtVm	dref
vsuxseg5ei16.v	V	Vd,0(s),VtVm	dref
vsuxseg5ei32.v	V	Vd,0(s),VtVm	dref
vsuxseg5ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsuxseg5ei8.v	V	Vd,0(s),VtVm	dref
vsuxseg6ei16.v	V	Vd,0(s),VtVm	dref
vsuxseg6ei32.v	V	Vd,0(s),VtVm	dref
vsuxseg6ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsuxseg6ei8.v	V	Vd,0(s),VtVm	dref
vsuxseg7ei16.v	V	Vd,0(s),VtVm	dref
vsuxseg7ei32.v	V	Vd,0(s),VtVm	dref
vsuxseg7ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsuxseg7ei8.v	V	Vd,0(s),VtVm	dref
vsuxseg8ei16.v	V	Vd,0(s),VtVm	dref
vsuxseg8ei32.v	V	Vd,0(s),VtVm	dref
vsuxseg8ei64.v	V	Vd,0(s),VtVm	dref | eew64
vsuxseg8ei8.v	V	Vd,0(s),VtVm	dref
vt.maskc	ventana condops	d,s,t	−
vt.maskcn	ventana condops	d,s,t	−
vwadd.vv	V	Vd,Vt,VsVm	−
vwadd.vx	V	Vd,Vt,sVm	−
vwadd.wv	V	Vd,Vt,VsVm	−
vwadd.wx	V	Vd,Vt,sVm	−
vwaddu.vv	V	Vd,Vt,VsVm	−
vwaddu.vx	V	Vd,Vt,sVm	−
vwaddu.wv	V	Vd,Vt,VsVm	−
vwaddu.wx	V	Vd,Vt,sVm	−
vwcvt.x.x.v	V	Vd,VtVm	alias
vwcvtu.x.x.v	V	Vd,VtVm	alias
vwmacc.vv	V	Vd,Vs,VtVm	−
vwmacc.vx	V	Vd,s,VtVm	−
vwmaccsu.vv	V	Vd,Vs,VtVm	−

Instruction	Extension	Parameters	Flags
vwmaccsu.vx	V	Vd,s,VtVm	–
vwmaccu.vv	V	Vd,Vs,VtVm	–
vwmaccu.vx	V	Vd,s,VtVm	–
vwmaccus.vx	V	Vd,s,VtVm	–
vwmul.vv	V	Vd,Vt,VsVm	–
vwmul.vx	V	Vd,Vt,sVm	–
vwmulsu.vv	V	Vd,Vt,VsVm	–
vwmulsu.vx	V	Vd,Vt,sVm	–
vwmulu.vv	V	Vd,Vt,VsVm	–
vwmulu.vx	V	Vd,Vt,sVm	–
vwredsum.vs	V	Vd,Vt,VsVm	–
vwredsumu.vs	V	Vd,Vt,VsVm	–
vwsub.vv	V	Vd,Vt,VsVm	–
vwsub.vx	V	Vd,Vt,sVm	–
vwsub.wv	V	Vd,Vt,VsVm	–
vwsub.wx	V	Vd,Vt,sVm	–
vwsubu.vv	V	Vd,Vt,VsVm	–
vwsubu.vx	V	Vd,Vt,sVm	–
vwsubu.wv	V	Vd,Vt,VsVm	–
vwsubu.wx	V	Vd,Vt,sVm	–
vxor.vi	V	Vd,Vt,ViVm	–
vxor.vv	V	Vd,Vt,VsVm	–
vxor.vx	V	Vd,Vt,sVm	–
vzext.vf2	V	Vd,VtVm	–
vzext.vf4	V	Vd,VtVm	–
vzext.vf8	V	Vd,VtVm	–
wfi	I	--	–
wrs.nto	Zawrs	--	–
wrs.sto	Zawrs	--	–
xnor	Zbb | Zbkb	d,s,t	–
xor	C	Cs,Ct,Cw	alias
xor	C	Cs,Cw,Ct	alias
xor	I	d,s,j	alias
xor	I	d,s,t	–
xori	I	d,s,j	–
xperm4	Zbkx	d,s,t	–
xperm8	Zbkx	d,s,t	–
zext.b	I	d,s	alias
zext.h	I	d,s	macro
zext.h	Zbb	d,s	–
zext.h	Zbb	d,s	–
zext.w	I	d,s	macro
zext.w	Zba	d,s	alias
zip	Zbkb	d,s	–

1.16 Further reading

1. The latest release of the Instruction Set Architecture (ISA):
   github/riscv

2. Linux Standard Base Core Specification, Generic Part. This is the official documentation for the ELF file format, for the debug frame machinery, etc. Download it from linuxfoundation.org.

3. The MaskRay blog. This is a very interesting blog full of references to low level stuff both for ARM, RISCV and other stuff, even windows. See: maskray.me

4. Improving DWARF. This is a very good and very readable critique of DWARF tables, presenting a DWARF table verifier, and, in general, a new perspective in debug tables and stack unwinding. (Slides) inria/france

5. This is the research article for the above slides. acm.org

6. A complete description of riscv relocations: sifive-blog

7. The specifications of the "Zbb" (bit manipulation) extension. github/riscv/bitmanip

8. The official specifications for the assembler: github/riscv/asm

9. Specifications for the Thead processor. www.t-head.cn

1.17   Answers to all exercises

Exercise 1: *The instruction* `bgt` *is an alias. How would you build it from the other instructions?*
 Answer:
Just invert the arguments. `bgt rs1,rs2,label` → `blt rs2,rs1,label`

Exercise 2: *Write a small program that uses a conditional branch.*
 Answer:

```
1    .globl main
2    main:
3    addi sp,sp,-16      // Build a stack frame
4    sd ra,8(sp)
5    sd s0,0(sp)
6    li t1,1             // t1 ← 1
7    li t2,2             // t2 ← 2
8    bgt t1,t2,.L1       // is t1 bigger than t2 ?
9    la a0,.LC1          // We did not branch. Load string address of LC1
10   j .L2               // Branch to call instruction
11   .L1:
12   la a0,.LC2          // We did branch. Load LC2 string
13   .L2:
14   call printf         // Do the call
15   ld ra,8(sp)         // Restore stack frame
16   ld s0,(sp)
17   jr ra               // Return
18   .LC1:
19   .string "Branch not taken\n"
20   .p2align 2
21   .LC2:
22   .string "Branch taken\n"
```

Executing:

```
1    star64:~/tiny-asm$ ./asm -o bgt.o bgt.s
2    star64:~/tiny-asm$ gcc bgt.o
3    star64:~/tiny-asm$ ./a.out
4    Branch not taken
```

Exercise 3: *Disassemble the program. What you see instead of bgt?*
 Answer:

```
14: 0063c863          blt t2,t1,24 <.L1>
```

The assembler changed source and destination, using `blt`.
Exercise 4: *Change bgt into blt line 8. Does the output change?*
 Answer:

```
star64:~/tiny-asm$ ./a.out
Branch taken
start64:~/tiny-asm$
```

Exercise 5: *How is the change achieved? Look at the source asm.c.*
 Answer:
Looking at the opcode table we have:

```
{"bgt",0,INSN_CLASS_I,"t,s,p",MATCH_BLT,MASK_BLT,match_opcode,INSN_ALIAS|
    INSN_CONDBRANCH}
{"blt",0,INSN_CLASS_I,"s,t,p",MATCH_BLT,MASK_BLT,match_opcode,INSN_CONDBRANCH}
```

We can see that in the case of `bgt`, the instruction is marked as an *alias*. We see also that the match and the mask are identical of `bgt` and `blt`. The essential difference is in the argument string: bgt has `"t,s,p"`, and in the mask of `blt` we have `"s,t,p"`.

Exercise 6: *Use the XOR instruction to invert all bits in an integer register*
 Answer:
Since $1 \oplus 1$ is zero, and $0 \oplus 1$ is 1, it suffices to have a right hand side of all ones (the number -1) and we are all set.

The instruction `not` (invert all bits) is `xor rs1,-1`. You can see this in the opcode table:

```
{"not",0,INSN_CLASS_I,"d,s",MATCH_XORI|MASK_IMM,MASK_XORI|MASK_IMM,match_opcode,
    INSN_ALIAS},
```

It has the `INSN_ALIAS` bit set, and the match is `MATCH_XORI`.

Exercise 7: *Write a program in assembler to print these 3 counters.*
 Answer:

```
1    .globl main
2    main:              Use of this name allows us to use C runtime
3    addi sp,sp,-16     Make room to establish a stack frame
4    sd ra,8(sp)        Save return address
5    sd s0,0(sp)        Save old stack frame
6    addi s0,sp,16      Establish a new frame. This is not actually needed.
7    rdtime a1          Read time into a1, that is the second argument
8    lla a0,.LC0        to printf. The first is the LC0 string
9    call printf        Let printf do the job
10   rdcycle a1         The same thing for the cycles. Use LC1.
11   lla a0,.LC1        String into a0
12   call printf
13   rdinstret a1       And the same for instructions returned. Use LC2
14   lla a0,.LC2        String into a0
15   call printf
16   ld ra,8(sp)        Restore return address
17   ld s0,0(sp)        Restore old frame pointer
18   jr ra              Return to the startup code
19   .LC0:
20   .string "Time=0x%x\n"
21   .LC1:
22   .string "Cycles=0x%x\n"
23   .LC2:
24   .string "Instructions executed=%d\n"
```

Executing this yields:

```
Time=133663617138
Cycles=164249353990
Instructions executed=85266785454
```

Exercise 8: *Try to verify that time corresponds to a time measure*
 Answer:
One way to do that is to call our program, then do something, then call it again. This should be a measure of how much time this "do something" takes. If we repeat that, we should arrive at similar results.

We will use "uptime" a utility that prints the time since startup.

```
star64:~/tiny-asm$ ./a.out;uptime;./a.out
Time=139156277825
Cycles=81886690065
Instructions executed=52342547651
```

```
15:22:18 up  9:39,  2 users,  load average: 0.00, 0.00, 0.00
Time=139156334252
Cycles=81891385713
Instructions executed=52345512910

star64:~/tiny-asm$ ./a.out;uptime;./a.out
Time=139235738528
Cycles=81923228282
Instructions executed=52376705932
15:22:38 up  9:40,  2 users,  load average: 0.00, 0.00, 0.00
Time=139235795755
Cycles=81927968778
Instructions executed=52379666534
star64:~/tiny-asm$
```

Exercise 9: *Use the "max" instruction to calculate the absolute value of a signed integer.*
 Answer:
Use:

```
    neg rd,rs1
    max rd,rs1,rd
```

Exercise 10: *Write an assembler program to show the CSR flags*
 Answer:

```
1       .globl main
2       main:
3       addi sp,sp,-16      Establish stack frame
4       sd ra,8(sp)         Save return address
5       sd s0,0(sp)         Save frame pointer
6       frcsr a1            Read the control register into the fisrt arg (a1)
7       lla  a0,.LC0        Read the string into the first argument (a0)
8       call printf         Call printf
9       ld ra,8(sp)         Restore return address
10      ld s0,0(sp)         Restore frame pointer
11      add sp,sp,16        Destroy stack frame
12      ret                 Bye bye
13      .LC0:
14      .string "CSR= 0x%lx\n"
```

This whole things makes just `printf("CSR=0x%x\n",csr);` But... wait, there is a bug.

```
$ asm -o rcsr.o rcsr.s
$ gcc rcsr.o
$ ./a.out
CSR= 0x0
$ ???
```

Well, of course. There wasn't any motives to set any of those flags above, and the rounding mode is zero (RNE). To see that we are really reading the csr let's provoke a division by zero, so at least we have something in there. We add following lines:

```
5       sd s0,0(sp)         Save frame pointer
6       li t1,12            Put 12 in register t1
7       fcvt.d.l f20,t1     Convert it to 12.0 in register f20
8       fcvt.d.l f21,x0     Put zero into register f21
9       fdiv.d f10,f20,f21  Divide 12.0/0.0
```

```
10    frcsr a1           Read the control register into the fisrt arg (a1)
11    etc                Rest is the same
```

Now, it should show the Division by zero bit as ON.

```
$ asm -o rcsr.o rcsr.s
$ gcc rcsr.o
$ ./a.out
CSR= 0x8
$
```

That was it!

Exercise 11: *Write a subroutine that returns the flags of the CSR*
 Answer:

```
1    # C Interface: int readcsr(void);
2    .globl readcsr
3    readcsr:
4    frcsr a0          Read the control and status register into a0
5    andi a0,a0,15     Select the 4 lower bits and leave result in a0 (ABI)
6    jr ra             return
```

Here it is not necessary to build a stack frame since we do not make any calls, and we do not use any local variables. We build our result in the established register for returning results (a0).

Exercise 12: *Mismatch between source and disassembly. Explain*
 Answer:

The explanation: actually, the instruction at address 4 is `addi a0,0`, what is actually a mov instruction. But the zero is not zero, since it is just a placeholder for a relocation. Looking at the relocations we see:

```
1    Disassembly of section .text:
2
3    0000000000000000 <main>:
4    0:  00000517          auipc  a0,0x0
5    0: R_RISCV_PCREL_HI20 .L2
6    0: R_RISCV_RELAX  *ABS*
7    4:  00050513          mv  a0,a0
8    4: R_RISCV_PCREL_LO12_I .L0
9    4: R_RISCV_RELAX  *ABS*
10   8:  00008067          ret
```

We see that a relocation points to address 4, indicating to put the lower 12 bits of the main address into an immediate and add them to a0.

Index