

tiny-asm: an assembler for riscv

jacob navia

Contents

| | |
|---------------------------------------------|----|
| Contents | 3 |
| 1 The RISC-V assembler | 5 |
| 1.1 Introduction | 5 |
| 1.2 Building tiny-asm | 6 |
| 1.3 Overview | 6 |
| 1.4 Instruction formats and encoding | 11 |
| 1.5 The instruction formats | 12 |
| 1.5.1 The "R" format | 12 |
| Software handling | 13 |
| 1.5.2 The "I" format | 13 |
| Software handling | 13 |
| 1.5.3 The "U" format | 14 |
| Software handling | 15 |
| 1.5.4 The "S" format | 15 |
| 1.5.5 The "B" format | 16 |
| 1.5.6 The "J" format | 17 |
| 1.6 The compressed instructions | 17 |
| 1.6.1 The compressed register (CR) format | 18 |
| 1.6.2 The compressed immediate (CI) format | 19 |
| 1.6.3 The stack relative store (CSS) format | 20 |
| 1.6.4 The wide immediate (CIW) format | 20 |
| 1.6.5 The compressed load (CL) format | 21 |
| 1.7 The opcode table | 22 |

1 The RISC-V assembler

1.1 Introduction

The tiny assembler is a "digest" of the GNU `gas` assembler. I have extracted from the 1.3Gb of `binutils-gdb` source code¹ two files: `asm.c` and `asm.h`.

There are two goals here:

1. To produce a small and fast assembler to be used as a compiler assembler. The elimination of features proceeds according to this goal: assemble machine generated output, without consideration for any human user, since all input to the assembler is supposed to be machine generated.
2. To produce a minimal set of sources that is *easy to read and understand* so that people can hack away without a lengthy learning curve. This documentation also, contributes to this objective.

In this version of the tiny-assembler there isn't:

- No input pre-processing. No include files, nor any fancy macro processing.
- No fancy error messages, messages will be emitted only in english. If you want other language error output you are welcome to do it yourself. The rationale behind this is obviously that a high level language user, programming in C++ or C, will be completely unable to understand the assembler messages even if they are translated into his/her native language.
- This assembler is geared to the riscv CPU. All support for any other machine has been dropped, specially support for machines that have ceased to exist for more than 20 years: the Motorola 68000 family, the Sparc, the Z80, etc. I think that even `gas` could drop support for those machines also.
- The code has been cleaned up from all cruft like this:

```
/* The magic number BSD_FILL_SIZE_CROCK_4 is from BSD 4.2 VAX
 * flavoured AS. The following bizarre behaviour is to be
 * compatible with above. I guess they tried to take up to 8
 * bytes from a 4-byte expression and they forgot to sign
 * extend. */
#define BSD_FILL_SIZE_CROCK_4 (4)
```

So, we are still in 2023 keeping bug compatibility with an assembler for a machine that ceased production in 2000?

¹I have just done a `du -b ./binutils-gdb` Probably is a bit less since I didn't do an extensive search for only `.c` and `.h` files.

- All the indirection through macros that are expanded into members of function tables that makes the code impossible to follow are eliminated. Now, if you see code like `foo(bar)`; it is highly likely that you are calling function `foo` with argument `bar`...
- All libraries are eliminated. Tiny-asm doesn't use `BFD` nor `libiberty` nor `libopcodes`. The only library used is `zlib`.
- There are only two files: `asm.c` and `asm.h`. No other include files are there, as far as I remember, excepting system includes like `stdio.h` of course.

1.2 Building tiny-asm

The build process runs as follows:

1. Download the software from github
2. Build it:

```
$ gcc -o asm asm.c -lz
```

That is it. There is no Makefile but you can write one. I wrote this one:

```
star64:~/tiny-asm$ cat Makefile
asm: asm.o
    gcc -o asm asm.o -g -lz

asm.o: asm.c asm.h
    gcc -W -Wall -Wstrict-prototypes -Wmissing-prototypes\
        -Wshadow -Wwrite-strings -g -c asm.c
clean:
    rm -f asm.o asm
```

The Makefile for `gas` is 2268 lines... an impressive piece of software. However I think that 9 lines is much easier to understand. The user wants to use an assembler, maybe modify it, so there is no point in making him/her try to modify a 2 thousand line Makefile.

1.3 Overview

Like all assemblers, this assembler has a **parser**, where the text of the input file is converted into logical units that represent either instructions for the machine, or for the assembler itself, called *pseudo instructions*, and an **encoder**, where the instruction and its arguments are encoded into a 32 or 16 bit instruction and added to the current fragment. And then there is the object file generation, where the instructions and associated information are packed into the ELF format.

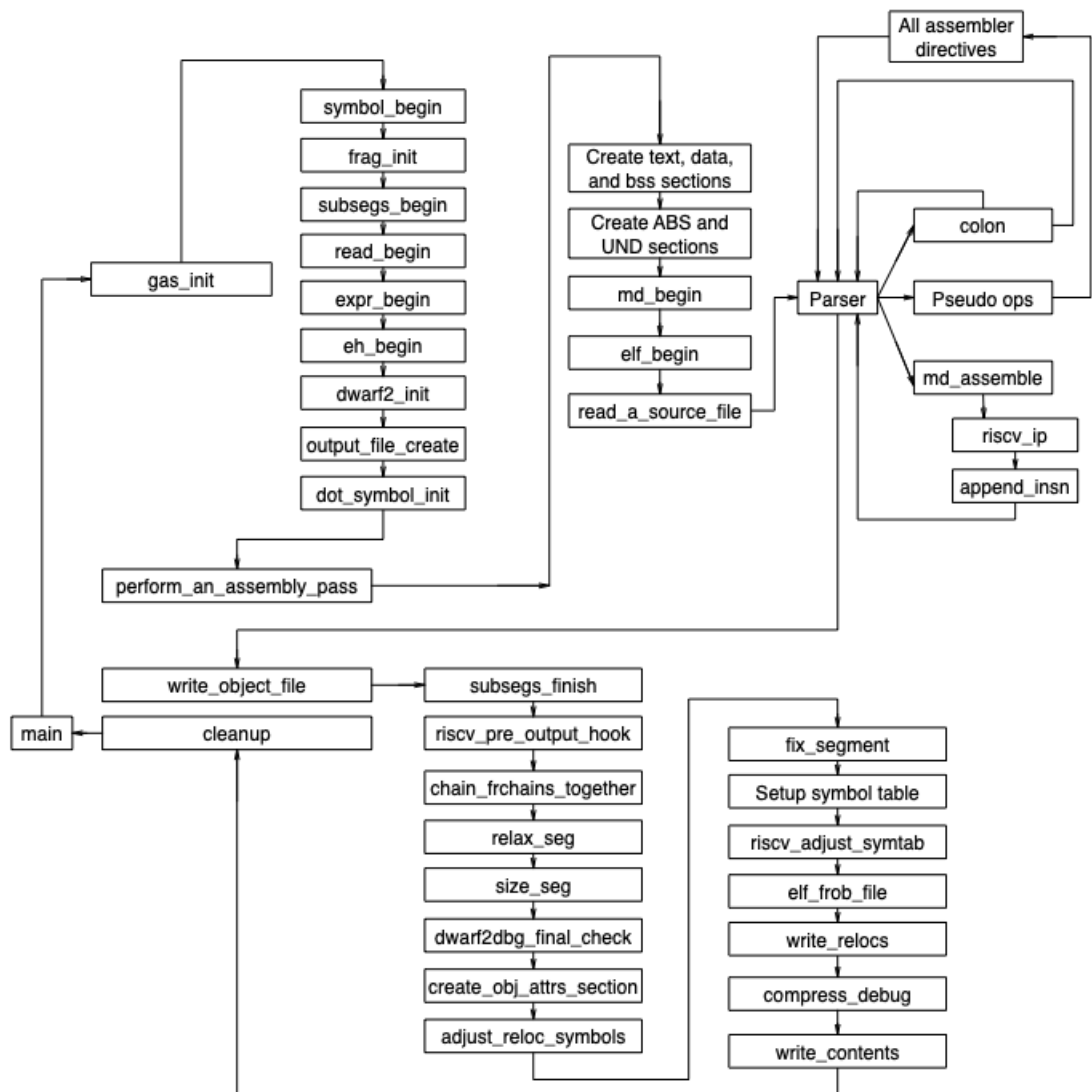


Figure 1.1: Overview of the assembler control flow

In figure 1.1 (page 7) we have these three main parts. Please keep in mind that this is a high level abstraction of the control flow. Obviously, if we would put each statement in the diagram we would have crammed 40 000 lines into a diagram... too much.

We start with `main` that organizes all three parts². It calls the initialization, `gas_init`, that initializes the symbols (`symbol_begin`), the fragments initialization, the sub-segments, etc.

"Fragments" are understood in the assembler as pieces of code already assembled but that can grow, getting new instructions or other data. They are of variable length, and they

²Please be aware that in the diagram there is a direct link between, for instance, the function `dot_symbol_init` and `perform_an_assembly_pass`. This does NOT mean that the first calls the second directly. It means that the flow of the program returns to `gas_init` and then returns to the main function, and it is `main` function that calls `perform_an_assembly_pass`.

That would be quite complicated to draw, however. So, the diagram simplifies this.

will be strung together in a process called "relaxation" at the end of the assembly.

The initialization of the "sub-segments" means the text, data, and bss sections are created. Are "sub-segments" just plain object file sections? Not quite. There are "sections" like the "ABS" (absolute) section or the "UND" (undefined) sections that will never be written out in the object file.

There are other initializations that give us the opportunity of explaining some concepts that will be important later on. The `eh_begin` function, for instance, initializes the "exception handling" stuff. This is a complicated system that allows languages like C++ to walk the stack at run time, searching for a handler that will accept handling the exception that has just occurred.

This process involves an impressive machinery that contains a set of tables that associate addresses in the code to descriptions of the stack contents that allow a debugger or a run-time interpreter to see what functions have in terms of local variables and the space that each stack frame uses in the stack. And even if you are programming in C and you do not have any need for exceptions you will get them anyway since your C code could be called from a C++ program.

Other initializations concerns the start of the `dwarf2` debug information generation. Yes, the assembler can emit debug information for the program it is assembling. This way, the assembly programmer can follow the program line by line. `tiny-asm` has kept this even if it is highly unlikely that the compiler, that emits its own and much richer debug information, will need this.

The initialization of the "dot symbol" needs also some explaining. The current location when assembling a program is called "dot", i.e. a point. This symbol is always associated with the current address following a long assembler tradition that goes back to the start of the micro-computer age.

Eventually we come to the `perform_an_assembly_pass` function. This one continues the initialization process by creating the standard sections of the object file:

- The text section. This is a misnomer since there isn't anything textual inside. It contains the binary codes that will be interpreted by the integrated circuit. This is the most important output of the whole assembly process.
- The data section. This contains the tables, constants, structures and everything that the programmer has defined as static data that will be loaded at the start of the program by the program loader.
- The BSS section that contains nothing. It is just a reserved memory space that will be allocated by the program loader when it loads the program and contains always zeroes at the start.
- There are many other sections in an ELF format file. Let's stop here.

The `md_begin` function reads all the static tables and builds hash tables from the for fast access. The opcodes are stored in hash tables, together with other data like the register names, the Control and Status Registers (CSRs) and what have you. Just read the code for more information.

The `elf_begin` function builds symbols for each section in the object file. This allows to emit relocations or symbol addresses as an offset from the start of the section.

In the diagram below, the functions aren't shown with their actual names but with their functional description. The GAS developers took (as you can see) a lot of effort to choose clear names that describe quite well what each function is doing. Still, I thought that here we will use functional boxes instead of function names, since some of the functions described here do not exist as a separated subroutine but they are just pieces of `read_a_source_file`.

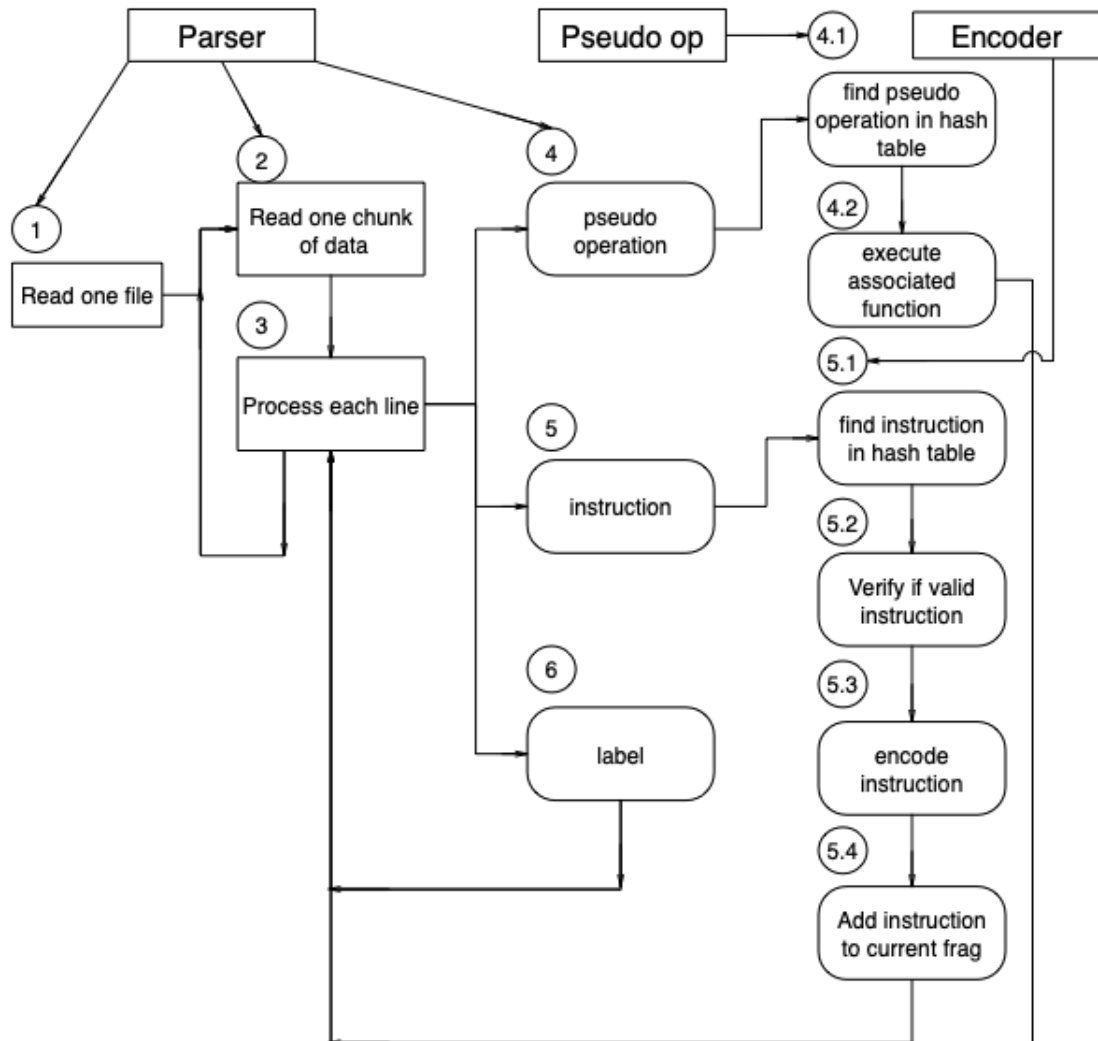



Figure 1.2: A more detailed view of the parser

We assume that the assembler input is a single file containing instructions, data, and assembler directives. In this version of the assembler, parsing is reduced to a bare minimum since we assume that we are assembling compiler output, and all the sophistication that is needed for an assembler adapted to human use is not needed for an assembler that is used to parse machine output.

We start with the function `read_a_source_file` that organizes the parsing and the instruction generation³.

1. Setup. Here, we setup the input file name, in variable `physical_input_file` and we care about writing a file name record if we are emitting debug information.
2. We read a chunk of the input file. Currently, `BUFFER_SIZE` is set to `256 * 1024`, and can be changed just by editing the corresponding line in `asm.h`

³Actually, the initialization phase is executed before, but we will abstract that away for the time being

Figure 1.3: `read_a_source_file` function

3. We start parsing lines. The first thing we read should be a symbol. If it ends with a colon, it is a label definition. We call the corresponding function `colon()` and continue parsing. If it is not finished by a colon, we see if the first letter is a point. If it is, it is an assembler directive. We call the corresponding function stored in the `pseudo_ops` structure (called `pseudo_typeS`) and we go fishing for the next line. If it is not a pseudo-operation, it *must* be a machine instruction. We call the `md_assemble` function.

The `md_assemble` function does basically following things:

1. Test if the instruction is valid using the current set of RISC-V specifications. There are instructions that can be issued only with 64 or even 128 bits, or floating point instructions that depend on floating point being implemented in hardware, etc. RISC-V machines can have a number of extensions implemented, since the basic ISA (Instruction Set Architecture) doesn't even have multiplication or division!

Each "extension" has a letter that characterizes it. For instance, in the machine I am using we have in `/proc/cpuinfo` a line with:

```
isa : rv64imafdc
```

This means that the machine is a risc V 64 bits machine (rv64), with the integer (i), multiplication (m), a (Atomic), f (single precision floating point), d (double precision floating point) and c (Compressed instructions in 16 bits) extensions. The assembler should test if any instruction is legal in the current subset, and reject those that do not comply.

Since we are an assembler for reading compiler output, we just assume the compiler doesn't emit wrong instructions and skip this test.

2. We call the `riscv_ip` function to encode the instruction. Basically it uses the `args` character string to know what arguments are expected. It verifies that those are correct, and inserts all necessary bits at the required positions. We will see later how these formats are defined.
3. If assembly succeeded the new instruction is added to the current fragment.

The `riscv_ip` function is basically a huge switch statement. The function will go through each one of the characters present in the `args` string of the opcode and add the necessary bits to the instruction.

In the tables below you will find the description of the different formats defined for each of the instructions in a riscv machine. These tables will help you understand how `riscv_ip` works.

1.4 Instruction formats and encoding

To understand how the assembler works, it is important to keep in mind how the machine works, the names of its parts, and the intricacies of instruction encoding. Yes, yes, that looks awfully dry and uninteresting. But (to me) it is interesting, and if you do not like to *understand* how things work, please go to tik-tok and play some games...

There are several types of instruction encoding, named **R, I, S, B, U, J**.

- All are 32 bits, like the ARM.
- The first 7 bits are reserved for the opcode (bits 0 to 6).
- The same operand, for instance the source register 1 (sr1) is at the same position, bits 15 to 19.
- All instructions have at least one register operand.
- Since we have 32 registers, all register encoding take 5 bits.

The riscv introduces a more functional naming schema, where registers are assigned usage names, instead of the register numbers. Here is a correspondence table between them:

Table 1.1: RISCv symbolic register names

| Register name | ABI name | Description | Register name | ABI name | Description |
|-------------------|----------|-----------------------------------|---------------|----------|------------------|
| Integer registers | | | | | |
| x0 | zero | Hard-wired zero | x16 | a6 | Seventh argument |
| x1 | ra | Return Address | x17 | a7 | Eighth argument |
| x2 | sp | Stack pointer | x18 | s2 | Saved 2 |
| x3 | gp | Global pointer | x19 | s3 | Saved 3 |
| x4 | tp | Thread Pointer | x20 | s4 | Saved 4 |
| x5 | t0 | Temporary/Alternate link register | x21 | s5 | Saved 5 |
| x6 | t1 | Temporary | x22 | s6 | Saved 6 |
| x7 | t2 | Temporary | x23 | s7 | Saved 7 |
| x8 | fp/s0 | Frame pointer | x24 | s8 | Saved 8 |
| x9 | s1 | Saved 1 | x25 | s9 | Saved 9 |
| x10 | a0 | First argument / Return value | x26 | s10 | Saved 10 |

Table 1.1: RISC-V symbolic register names

| | | | | | |
|--------------------------|---------|-----------------------------------|---------|----------|---------------------|
| x11 | a1 | Second Argument / Return value | x27 | s11 | Saved 11 |
| x12-x15 | a2-a5 | Argument 3-5 | x28-x31 | t3-t6 | Temporary registers |
| Floating point registers | | | | | |
| f0-f7 | ft0-ft7 | Fp temps | f2-f7 | fa2-fa7 | function arguments |
| f8-f9 | fs0-fs1 | Fp saved registers | f18-f27 | fs2-fs11 | saved registers |
| f10-f11 | fa0-fa1 | Fp arguments/return value | f28-f31 | ft8-ft11 | Temporary registers |

The difference between the ABI names and the actual register numbers is due to the fact that the ranges of registers are not contiguous. For instance the range of saved registers has two of them as x8 and x9, then the rest is x18 to x27.

1.5 The instruction formats

Each format is designed to be used by similar type of instructions.

- **R** Register to register ALU instructions.
- **I** Immediate and load.
- **S** Store and comparisons.
- **B** Branch.
- **U J** Jump and jump with link (call) instructions.

The RISC-V manual comments these formats like this

The RISC-V ISA keeps the source (rs1 and rs2) and destination (rd) registers at the same position in all formats to simplify decoding. Except for the 5-bit immediates used in CSR instructions, immediates are always sign-extended, and are generally packed towards the leftmost available bits in the instruction and have been allocated to reduce hardware complexity. In particular, the sign bit for all immediates is always in bit 31 of the instruction to speed sign-extension circuitry. ⁴


1.5.1 The "R" format

This format features 3 registers (destination, source 1 and source 2) and has two fields of 3 and seven bits available for use to customize the opcodes. We use a 32 bit addition as an example of this format: `addw a0,a0,a1`. The addition using ABI names is `addw a0,a0,a1` but using actual register numbers we have `addw x10,x10,x11`. For this instruction the 10 bits of `extra-3` and `extra-7` are empty.

We have then:

- Opcode: 0 1 1 1 0 1 1 \rightarrow 0x3b (59 decimal).
- Destination register: 0 1 0 1 0 \rightarrow 0xA (10 decimal). Register 10 is a0, that contains the first argument and is loaded with the result.
- Source 1: 0 1 0 1 0 \rightarrow 0xA (10 decimal). Register 10 (a0) is the first source.
- Source 2: 0 1 0 1 1 \rightarrow 0xB (11 decimal). Register 11 (a1) is the second source.

⁴RISC-V User level ISA V 2.2 §2.2. They add further down: Decoding register specifiers is usually on the critical paths in implementations, and so the instruction format was chosen to keep all register specifiers at the same position in all formats at the expense of having to move immediate bits across formats

Figure 1.4: **R** Instruction layout

Software handling

We have an instruction with the `args` format of "`Cs,Cw,Ct`" that expects source and destination to be identical (`s` and `w`) followed by a target register in the expected range for compressed registers. All of that is true, and we succeed with a compressed 16 bit instruction.

Obviously this is not what we wanted. We wanted a 32 bit 'R' instruction. To be able to do that, we add the following instruction at the top of our assembler file


```
.option arch -c
```

I.e. we disable all compressed instructions.

We see here that the *order* in the layout of the opcode table is very important. The instructions that are **more** constrained should come first, and the general formats should come last. For instance the compressed instruction should come first, and non-compressed last, since the software stops at the first match.

1.5.2 The "I" format

This format changes the "R" format by merging `src2` with `extra-7` to give a 12 bit field where an immediate integer value can be stored (up to $2^{12} - 1 \rightarrow 4095$ values can be stored).

Figure 1.5: **I** Instruction layout

Software handling

The first instruction that the software tries has its `args` string: "`Cs,Cw,Ct`", we expect a source register in compressed format, i.e. register 8-15, followed by the *same* register. The second condition succeeds, and the software passes to the third argument: we expect a register, and we find the constant 1025. Nope, this instruction is not the one.

The next `addw` instruction to be tested has the string "`Cs,Ct,Cw`", a permutation of the above that fails also, for the same reasons.

More instructions are tried, with strings `d,Cu,Co` that fails, "`d,s,t`" that fails also since we have an immediate constant and not a register in the third position (`'t'` field). At last we arrive at an instruction with `args` field of `d,s,j`", i.e. a sign extended immediate (`'j'`) in the third position. This time the software succeeds and we are done. Accessing the different fields is done with macros. Here is one example of a series of macros that extracts the immediate field of the immediate value in the instruction above

```
#define RV_X(x, s, n)  (((x) >> (s)) & ((1 << (n)) - 1))
```

This macro extracts `<n>` bits from `<x>`, beginning in bit position `<s>`. It has two parts:

1. The left side of the "and" operation that shifts the given number `<s>` bits to the right to bring it to position zero, and
2. An expression that builds a mask of `<n>` 1 bits by shifting a 1 `<n>` positions to the right and subtracting one, what gives a power of two minus 1. A power of two minus 1 is a field full of ON bits in two's complement notation. For instance $1 \ll 3 \rightarrow 8$ (1000). You subtract 1 from that and you obtain 0111 (7), i.e. 3 bits "on", a mask to extract the lower 3 bits from a number.

```
#define RV_IMM_SIGN(x) (-(((x) >> 31) & 1))
```

This macro returns either -1 or 0, depending if the sign of the 32 bit number is negative or positive. Since -1 is 32 bits of "1" bits, it can be used to sign extend a number.

The two macros above are used in these new ones:


```
#define EXTRACT_ITYPE_IMM(x) (RV_X(x,20,12)|(RV_IMM_SIGN(x) << 12))
#define ENCODE_ITYPE_IMM(x) (RV_X(x, 0, 12) << 20)
```

The first macro extracts 12 bits from the given number (`<x>`) and sign-extends its sign. The second extracts the lower 12 bits of the value, and puts them at position 20-31 ⁵

1.5.3 The "U" format

A variant of the **I** format featuring more space for immediate constants is the **U** format, that can hold immediate constants with 20 bits.

Figure 1.6: **U** Instruction layout



The `lui`⁶ instruction loads an unsigned 20 bits immediate stored in the bits 12 to 31 of the instruction into the upper 20 bits of the destination and sets the lower 12 bits to zero.

⁵It is a pity that machines implementing the boolean extension aren't widely available yet. I miss the ARM boolean instructions that will reduce many of those macros to a couple of instructions.

⁶`lui` stands for **load upper immediate**

In C language notation we have: `dst = (imm20 << 12);`. The authors justify these choices with:

In practice, most immediates are either small or require all XLEN bits. We chose an asymmetric immediate split (12 bits in regular instructions plus a special load upper immediate instruction with 20 bits) to increase the opcode space available for regular instructions. Immediates are sign-extended because we did not observe a benefit to using zero-extension for some immediates as in the MIPS ISA and wanted to keep the ISA as simple as possible.⁷

Software handling

Looking up the `args` description for this instruction, we find the character string `"d,Cu"`. This means we should expect a register name, followed by a comma, and an immediate value to be able to use a C (compressed) instruction. But that doesn't work, our constant is beyond bounds of the compressed immediate.

The software continues its search for the correct instruction and we come to the next instruction in the list that has the `args` string `"d,u"`, without any compression requirements. This time a match is found, and necessary bits are inserted as shown in figure 1.6 page 14.

Obviously, loading an immediate constant that will be shifted by 12 bits is seldom used. This is thought for loading the upper 20 bits of an *address*, then adding the lower 12 bits with another instruction. This constant was chosen in this example so that it has a 1 bit at the end of 10 bits, and 1 at the start to be visible in the drawing.


To extract the J type immediate we use the following macro:

```
1  #define EXTRACT_UTYPE_IMM(x) ((RV_X(x, 12, 20) << 12) | (RV_IMM_SIGN(x) << 32))
```

1.5.4 The "S" format

In this format, the `dst` field disappears and its bits are used to hold the lower 4 bits of an immediate value. An instruction that uses this format is the `sd` (store double word) instruction.

Figure 1.7: S Instruction layout



We use the instruction `sd ra,64(sp)` as example. This instruction means: Store the contents of the return address register (`ra`) at the memory address obtained by adding 64 to the contents of the `sp` register. We have here an address that is obtained by adding the contents of a register and a *displacement* that must fit into 12 bit. As you can see here, this is a much easier format than the ARM jungle of different types of offsets where you never really know which one to use. The Risc-V manual specifies that all offsets are signed.⁸

We have then for this instruction:

⁷Riscv ISA Architecture §2.2

⁸They say:

- `src1` is 0 0 0 1 0, or register 2.
- `src2` is 0 0 0 0 1, or register 1.
- The immediate is the concatenation of `imm[4:0]` and `imm[11:5]` i.e; 0 0 0 0 0 0 1 0 0 0 0 0 or 64.


For extracting the immediate from the instruction we use the macro

```
1 #define EXTRACT_STYPE_IMM(x) \
2 (RV_X(x, 7, 5) | (RV_X(x, 25, 7) << 5) | (RV_IMM_SIGN(x) << 12))
```

This macro extracts five bits beginning at position seven, then 7 bits from position 25 upwards, shifted by 5 left, so that they come right after the first five. The whole is sign extended in the same way as explained in section 1.5.2 page 14.

1.5.5 The "B" format

Figure 1.8: **B** Instruction layout



In this format, we have a 13 bit immediate for branches. The immediate represents the amount that will be added to the program counter to reach the specified location, in multiples of 2. Since the lowest bit of the immediate will be always zero, it has been replaced by bit 11 (the twelfth bit) adding one bit to the quantity being written. The range of the branch is $\pm 4K$.

The different conditional branches are specified in the **extra-3** group, with

| extra-3 | Instruction | Description |
|---------|-------------|----------------------------------|
| 0 0 0 | beq | branch if equal |
| 0 0 1 | bne | branch if different |
| 1 0 0 | blt | branch if less than |
| 1 0 1 | bge | branch if greater/equal |
| 1 1 0 | bltu | branch if less than unsigned |
| 1 1 1 | bgeu | branch if greater equal unsigned |

All these instructions share the same opcode: 99. The **extra-3** field is used to extend the opcode for different instructions.

The macro to access the immediate value is way more complicated due to the bit scrambling...


```
1 #define EXTRACT_BTYPE_IMM(x) ((RV_X(x, 8, 4) << 1) | \
2 (RV_X(x, 25, 6) << 5) | (RV_X(x, 7, 1) << 11) | (RV_IMM_SIGN(x) << 12))
```

Except for the 5-bit immediates used in CSR instructions, immediates are always sign-extended, and are generally packed towards the leftmost available bits in the instruction and have been allocated to reduce hardware complexity. In particular, the sign bit for all immediates is always in bit 31 of the instruction to speed sign-extension circuitry.

1.5.6 The "J" format

The only difference between the U and J formats is that the 20-bit immediate is shifted left by 12 bits to form U immediates and by 1 bit to form J immediates. In the "J" format, the immediate represents an offset in pairs of 16 bit instructions from the current PC.

Figure 1.9: J Instruction layout



Why this scrambled layout? The Risc-v manual tells us:

Although more complex implementations might have separate adders for branch and jump calculations and so would not benefit from keeping the location of immediate bits constant across types of instruction, we wanted to reduce the hardware cost of the simplest implementations. By rotating bits in the instruction encoding of B and J immediates instead of using dynamic hardware muxes to multiply the immediate by 2, we reduce instruction signal fanout and immediate mux costs by around a factor of 2. The scrambled immediate encoding will add negligible time to static or ahead-of-time compilation. For dynamic generation of instructions, there is some small additional overhead, but the most common short forward branches have straightforward immediate encodings.

The macro to extract this monster from its hiding place looks like this

```

1 #define EXTRACT_JTYPE_IMM(x) ((RV_X(x, 21, 10) << 1) | (RV_X(x, 20, 1) << 11) | \
2   (RV_X(x, 12, 8) << 12) | (RV_IMM_SIGN(x) << 20))
3 #define ENCODE_JTYPE_IMM(x) ((RV_X(x, 1, 10) << 21) | (RV_X(x, 11, 1) << 20) | \
4   (RV_X(x, 12, 8) << 12) | (RV_X(x, 20, 1) << 31))

```

1.6 The compressed instructions

The Risc-v instructions are normally 32 bits in length. The "C" extension (C for **C**ompressed) encodes certain instructions in 16 bits, what leads to big savings in code size. These instructions aren't enabled by default in the assembler. You can enable them (if your machine actually supports them) with the instruction: `.option arch, +c`. Enabling them or not is not that important, since the linker will replace longer with shorter instruction whenever possible. For instance the jumps can't be really calculated until all the instructions are compressed, what only the linker can know.

The compressed instructions are enabled when one of these conditions is true:

- The compressed 16 bit instructions have the lowest 2 bits of the opcode set to either 00, 01, or 10.
- 32 bits instructions have their lowest two bits set to 11. The following 3 bits should have any value different from 111.
- The 48 bit instructions have their lowest 6 bits set to 011111. (5 bits set)

- 64 bit instructions have the 7 lower bits set to 0111111. (6 bits set)

The criteria for making a compressed instruction are as follows:

- The immediate or the address offset is small.
- One of the registers used is the **zero** register (x0), the return address register or link register **ra** (x1), or the stack pointer **sp** (x2).
- The destination and first source register are the same.
- The registers used belong to the 8 most popular ones, described with 3 bits in the table below⁹.

Table 1.3: Compressed register numbers

| number | 000 | 001 | 010 | 011 | 100 | 101 | 110 | 111 |
|-----------------|-----|-----|-----|-----|-----|-----|-----|-----|
| int reg. number | x8 | x9 | x10 | x11 | x12 | x13 | x14 | x15 |
| ABI name | s0 | s1 | a0 | a1 | a2 | a3 | a4 | a5 |
| FP reg number | f8 | f9 | f10 | f11 | f12 | f13 | f14 | f15 |
| FP ABI name | fs0 | fs1 | fa0 | fa1 | fa2 | fa3 | fa4 | fa5 |


There are nine different compressed instruction layouts.

In the table below the registers that use the 3 bit number are marked with a '.

| Meaning | Code | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 | |
|-------------|------|---------|----|----|-------------|----------|----|----------|---|-----|-----------|-----|------|----|----|----|---|--|
| Register | CR | Extra-4 | | | | dst/src1 | | | | | src2 | | | | | op | | |
| Immediate | CI | Extra-3 | | | S | rd/rs1 | | | | | immediate | | | | | op | | |
| Store local | CSS | Extra-3 | | | imm | | | | | rs2 | | | | | op | | | |
| Wide imm | CIW | Extra-3 | | | imm | | | | | | | rd' | | op | | | | |
| Load | CL | Extra-3 | | | imm | | | rs1' | | | imm | | rd' | | op | | | |
| Store | CS | Extra-3 | | | imm | | | rs1' | | | imm | | rs2' | | op | | | |
| Arithmetic | CA | Extra-6 | | | | | | rd'/rs1' | | | Extra-2 | | rs2' | | op | | | |
| Branch | CB | Extra-3 | | | offset | | | rs1' | | | offset | | | | | op | | |
| Jump | CJ | Extra-3 | | | jump target | | | | | | | | | | | op | | |

1.6.1 The compressed register (CR) format

Figure 1.10: Compressed **CR** Instruction layout



This format accepts instructions where the destination and the first source register are the same. It has four fields, here from right to left, i.e. from bit 0 to 15:


⁹Actually those numbers are just the normal register number modulo 8.

1. OP: Bits 0-1. Value: 2.
2. Src2: Bits 2-6. The second source register. Note that it is specified in 5 bits, like dst/src1, so any register of the set of 32 is possible, except the zero register. In this case it is 13, i.e. register a3 (x13).
3. dst / src1: Bits 7-11. The source 1 and the destination register are the same. Also specified in 5 bits, in this case it is 12: the a2 (x12) register.
4. Extra-4: Bits 12-15. Value: 9. Complements the opcode. This field can have two values that correspond to mv (move) or, in the example, add.

1.6.2 The compressed immediate (CI) format

These instructions perform operations between a register and a small immediate encoded in only 6 bits. The register can't be the zero register, and the immediate can't be zero. There

Figure 1.11: Compressed immediate **CI** Instruction layout



are four instructions that use the compressed immediate format. They differ in the **extra-3** field. From least significant bit to the most significant one we have:

1. OP: Bits 0-1, always with value 1 for the CI format.
2. The immediate field, in bits 2 to 6 that encodes immediate bits 0 to 4. In the example above this is 27, 1 1 0 1 1 in binary.
3. The destination and the source register number over 5 bits. In the example we have 14 since the register a4 has the number 14.
4. The sign of the immediate value in a single bit (index 12th).
5. The **Extra-3** field, that allows for 3 instructions to be distinguished: **addi**, **addiw**, and **addi16sp**. The last one adds a number of 16 bits quantities to the stack and is used to adjust the stack at the prologue or at the epilogue of a function. Since the stack must be aligned to a multiple of 16, there is no need to keep the lower 4 bits. This makes for adjustments of -512 to 496 bytes.

To access the immediate value we use

```

1  #define EXTRACT_CITYPE_IMM(x) (RV_X(x, 2, 5) | (-RV_X(x, 12, 1) << 5))
2  #define ENCODE_CITYPE_IMM(x) ((RV_X(x, 0, 5) << 2) | (RV_X(x, 5, 1) << 12))

```


The first macro uses the same technique for sign extending than our **RV_IMM_SIGN** (see 1.5.2 page 14). We just need another expression since the other was fixed for 32 bits.

1.6.3 The stack relative store (CSS) format

Five instructions use the CSS format:

1. `c.swsp` or store word to an offset from `sp`, the stack pointer.
2. `c.sdsp` or store double word (64 bits) to an offset from `sp`.
3. `c.fswsp` or store single precision (32 bits) to an offset from `sp`.
4. `c.fsdsp` or store double precision (64 bits) to an `sp` offset.

Figure 1.12: Store to stack offset (CSS) instructions layout



In our example instruction we have an `op` field of 2, an `src` field of 16 (10000) and the cryptic "011000" sequence that is translated into 00110 (6 decimal) since the bits are scrambled: they are stored as bits 5 4 3 2 7 6. The macros to access the immediate displacement here are:

```
1 #define EXTRACT_CSSTYPE_IMM(x) (RV_X(x, 7, 6) << 0)
2 #define ENCODE_CSSTYPE_IMM(x) (RV_X(x, 0, 6) << 7)
```

The encoding of instruction `c.swsp` needs only one source register: the source of the 32 bit data to store in memory. Any register will do since we have a register number in 5 bits. The value of the immediate displacement will be added to the stack pointer scaled by 4 to form the effective address. In the example above the 6 binary is scaled to 24. ¹⁰

The argument description string is "`CV,CM(Cc)`": We need a register name (`CV`), followed by a small constant (`CM`) that is a displacement (the parentheses) of the stack pointer (`Cc`). The constant value will be zero extended, since obviously negative offsets for the stack aren't very useful!


The reach of this instruction is $2^7 - 1$ values since we have 7 bits. Scaled by 4, i.e. $127 * 4 \rightarrow 508$.

And... "one more thing" as Steve Jobs liked to say, there is a problem with zero offsets from the stack pointer. Normally a zero offset is omitted, i.e. you do NOT write `sw a6,0(sp)`, you just write `sw a6,(sp)`. The handling of the `CM` directive tests for this with the function `riscv_handle_implicit_zero_offset`.

1.6.4 The wide immediate (CIW) format

This format is used to encode a constant in bits 5 to 12. It is used in the `addi4spn` instruction. The constant encoded in those 8 bits is scaled by 4, i.e. the two lower bits are implicit zeroes. The scaled value will be added to the stack pointer and written to the register whose index is stored in the 3 bits `rd`. This instruction builds then pointers to values stored in the local stack frame.

¹⁰By an unfortunate coincidence the scrambled bits of the constant are 011000, what is 24 in binary. Beware, nothing in this business is simple, and a 24 can be scrambled to 6, then scaled to 24 back again.

Figure 1.13: Store to stack offset (**CIW**) instructions layout

1. The OP field is zero.
2. The destination (**rd'**) is 7, the register number in 3 bits of the a5 register
3. Now, this is more complicated to explain. The poor immediate bits are *scrambled*, i.e. they are **not** in the natural order but in the order: 5, 4, 9, 8, 7, 6, 2, 3. The bits 1 and 0 are implicitly zero. The quantity (128) has a single bit on at the position 7, what in our scrambled layout corresponds to bit 8.¹¹ The Risc-V ISA manual justifies this saying:

The immediate fields are scrambled in the instruction formats instead of in sequential order so that as many bits as possible are in the same position in every instruction, thereby simplifying implementations.¹²

The "simplifying" above refers to hardware simplification.

4. The Extra-3 field is zero.

The macros used to access the immediate are:


```

1  #define EXTRACT_CIWTYPE_ADDI4SPN_IMM(x) ((RV_X(x, 6, 1) << 2) |\
2  (RV_X(x, 5, 1) << 3) | (RV_X(x, 11, 2) << 4) | (RV_X(x, 7, 4) << 6))
3  #define ENCODE_CIWTYPE_ADDI4SPN_IMM(x) ((RV_X(x, 2, 1) << 6) |\
4  (RV_X(x, 3, 1) << 5) | (RV_X(x, 4, 2) << 11) | (RV_X(x, 6, 4) << 7))

```

The argument description string for this instruction is "Ct,Cc,CK"

1.6.5 The compressed load (CL) format

Figure 1.14: Compressed load **CL** Instruction layout

¹¹The number 128 is 1000 0000 in binary. Bit 7 is one. In the scrambled order we have bit 7 in the fourth position of the immediate field, counting from left to right, as shown in the figure 1.13

¹²Risc-V Unprivileged ISA V20191213 §16.2

1. The OP field is zero.
2. The destination register is 5 (a3).¹³
3. This field corresponds to an offset from a register. The constant should be aligned by a multiple of 4, since we are loading 4 bytes. The two lower bits then should be zero and they are implicit, i.e. they are absent from the encoding. The value is split between two bits at positions 5 and 6, and the rest in positions 10, 11, and 12. The two bits in positions 5 and 6 are scrambled, and bit 6 corresponds to bit 2 of the immediate and bit 5 is bit 6 of the immediate value, they are not consecutive.
4. The `rs1'` field contains 1 1 1, what corresponds to x15 (a5).
5. We have in bits 10, 11, and 12 the bits 3, 4, and 5 of the immediate value.
6. The `extra-3` field contains constant 2.

1.7 The opcode table

The full table of opcodes (called `riscv_opcodes`) consists of entries with the following structure:

```
struct riscv_opcode {
    const char *name;
```

The name of the instruction in lower case. This is also the used as the key to the hash table. Several instructions can share the same name, and they are recognized by their different arguments.

```
    unsigned xlen_requirement;
```

The word bit length (32, 64, or 128) that is required to use this instruction. A zero here means no requirement.

```
    enum riscv_insn_class insn_class;
```

The instruction class to which it belongs. For instance the instructions belonging to the basic integer operations are `INSN_CLASS_I` one of the member of the `enum riscv_insn_class`. This was used to decide whether or not this instruction is legal in the current machine architecture context, but this test has been dropped since we assume that the compiler will not generate instructions that are illegal for the target machine.

```
    const char *args;
```

A string describing the arguments for this instruction. This string will be interpreted by the `riscv_ip` function in a rather big set of nested `switch` statements.

```
    insn_t match;
    insn_t mask;
```

The basic opcode for the instruction. When assembling, this opcode is modified by the arguments to produce the actual instruction that is used. If `pinfo` is `INSN_MACRO`, then this is 0. Otherwise the `mask` field is a bit mask used to isolate the relevant portions of the opcode when disassembling. If `pinfo` is `INSN_MACRO` then this field contains the macro identifier, encoded as a member of an anonymous enumeration and casted to an integer.

¹³These values are in table 1.3

```
int (*match_func) (const struct riscv_opcode *op, insn_t word);
```

A function to determine if a word corresponds to this instruction. Usually, this computes `((word & mask) == match)`.

```
unsigned long pinfo;
```

Additional information about the instruction. They are:

| Symbol | Description |
|-----------------|-----------------------------------------------------------------------|
| INSN_ALIAS | Just an alias, for example "mv" for "addi dest,src,zero" |
| INSN_BRANCH | Unconditional branch |
| INSN_CONDBRANCH | Conditional branch |
| INSN_JSR | Jump to a subroutine |
| INSN_DREF | Data reference |
| INSN_V_EEW64 | Instruction allowed only when the machine is a 64 bit machine or more |
| INSN_XX_BYTE | 5 different data size specifiers, for XX=1, 2, 4, 8, or 16 bytes |

```
};
```

The field `args` above needs more explanation. It is a one (or more) letters that represent the type of argument that can be expected in an instruction. This can be a register, a constant within a certain range, or other things. During assembly, the assembler reads and interprets this character string to weed out wrong choices or emit warnings, and to verify that all constraints are met.

The table below should document all the letters used by the `riscv_ip` function. They are listed in the order they appear there; only for the first level. If a letter has a continuation (for instance for the compressed instructions), the secondary switch statement is explained in another table¹⁴.

Table 1.6: Argument descriptions

| Char | Description |
|------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| \0 | End of the argument string. Here are done the final checks, for instance that this instruction corresponds to the bit length of the machine (64 bit instructions can't be done in a 32 bit machine). It checks also if the end of the argument string coincides with the end of the actual arguments present. If everything goes well it sets the errors to zero and branches to the end of the <code>riscv_ip</code> function. |
| C | Compressed format instructions. This leads to a nested switch statement, since all the compressed argument descriptions begin with a C letter. This switch is described in table 1.8 page 25. |
| V | Vector instructions. This leads to a nested switch statement too. |
| , | Synchronization. Arguments are separated by commas. The software tests this and ignores the separators. |
| ()[] | Displacement or index. Same behavior as for commas. |
| < | Shift amount for shifts less than 32. |
| > | Shift amount for 0 to word length - 1. |

¹⁴Nested tables are as difficult to read as nested `switch` statements.

Table 1.6: Argument descriptions

| | |
|----|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Z | CSRRxI Immediate. C ontrol and S tatus R egisters are specified in a different instruction format. For this to work, you have to have access to a CPU with the 'z' extension. |
| E | Control register number. This is used only in privileged instructions. |
| m | Rounding mode. This argument expects a character string representing the rounding mode. It can be one of "rne", "rtz", "rdn", "rup", "rmm", 0, 0, "dyn". See table 1.7 page 25. |
| PQ | Fence predecessor or successor |
| d | Destination register. |
| s | First source register. Also called src1 in the documentation. |
| t | Second source register. The 't' is for target. It is also called src2 in the documentation. |
| r | RS3 |
| D | Floating point destination register |
| S | Floating point source 1. |
| T | Floating point source 2 |
| U | Floating point source 1 and 2 |
| R | Floating point RS3. |
| F | Expects a bit field, that is defined by the following character |
| I | M_LI macro. Immediate value. |
| A | Requests a symbol |
| B | Requests a symbol or a constant. |
| j | Sign extended immediate. |
| q | Expects a register store displacement. |
| o | Expects a load displacement. |
| l | Used for thread local storage. |
| p | PC relative offset |
| 0 | Expects a zero displacement. For instance: <code>lr.w a5,0(sp)</code> . |
| u | Expects a 20 bit immediate |
| a | 20 bit relative offset. |
| c | Call using the global object table |
| O | Opcode field |
| y | bs immediate for branch offsets. |
| Y | rnum immediate |
| z | Expects a zero |
| W | Various operands |
| X | Integer immediate |

Below is the set of rounding modes for the `m` parameter. It has been taken from the Sifive site¹⁵. Edited in May 27th 2020.

Table 1.7: Accepted rounding modes for the 'm' parameter

| Binary Value | Mnemonic | Meaning |
|--------------|----------|---------------------------------|
| 000 | rne | Round to Nearest, ties to Even |
| 001 | rtz | Round towards Zero |
| 010 | rdn | Round Down (towards $-\infty$) |

¹⁵<https://observablehq.com/@nschwass/riscv-f-extension-single-precision-floating-point-instruction>
The URL seems truncated but it is not...

Table 1.7: Accepted rounding modes for the 'm' parameter

| | | |
|-----|-----|-----------------------------------------------------------------------------------------------|
| 011 | rup | Round Up (towards $+\infty$) |
| 100 | rmm | Round to Nearest, ties to Max Magnitude |
| 101 | | Invalid. Reserved for future use. |
| 110 | | Invalid. Reserved for future use. |
| 111 | dyn | In instruction's rm field, selects dynamic rounding mode; In Rounding Mode register, Invalid. |

The C (compressed) instructions are differentiated by the following letters:

Table 1.8: Compressed instruction types

| Char | Description | Char | Description |
|------|-----------------------------------------------|------|----------------------------------------------------------------------------------------------------------------------|
| s | Source register 1 (x8-x15) | w | Source 1 and destination when they are the same |
| t | Source 2 with x8-x15 | x | Source 2 and destination are the same. x8-x15 only. |
| U | Source 1 and destination the same. | v | Source 2 |
| c | Source 1 constrained to be sp | z | Source 2 should be the zero register |
| > | Shift amount between 0 and word length - 1 | 5 | Five bit field |
| 6 | Six bit numeric field | 8 | Eight bit field |
| j | Non-zero immediate | k | Immediate (possibly zero) |
| l | Load immediate (64 bits) | m | Load immediate |
| n | Immediate offset from SP | o | C.addiw, c.li, and c.andi allow zero immediate. C.addi allows zero immediate as hint. Otherwise this is same as 'j'. |
| K | scaled by 4 stack addend | L | Stack offset scaled by 16 |
| M | Scaled by 4 stack displacement(32 bits store) | N | Data reference with offset from stack scaled by 4(64 bits store) |
| u | Immediate for jumps | v | Immediate for jumps |
| S | Floating point source 1 x8-x15 | D | Floating point source 2 x8-x15 |
| T | Floating point source 2 | F | Field of 6, 4, 3, or 2 bits |

This is an example for an instruction entry in the opcodes table:

```
{"addi",0,INSN_CLASS_C,"Ct,Cc,CK",MATCH_C_ADDI4SPN,MASK_C_ADDI4SPN,\
match_c_addi4spn,INSN_ALIAS},
```

After parsing the name of the instruction, the `riscv_ip` function examines entries in the opcode table starting with the first one that has this name. It copies this entry into temporary storage because it will modify it later (using the `create_insn` function).

Then, it uses the letter in the `args` character string to check if there is a match. If there is, it stores immediately the bits into the instruction copy. But, as mentioned above, if there isn't any match, all the work is discarded and `riscv_ip` starts over using a saved pointer to the start of the arguments.

This way it ensures that eventually, the good instruction will be discovered, if at all. It is a slow process, since in many cases 4 other 5 instructions will be parsed and discarded until

the correct one is found. Since the order of the opcodes is crucial the most used instructions can be the last ones to be found, what compounds the problem.

Several solutions can be imagined to speed up things, but the question arises if the speed of the assembler encoding is really the limiting factor for the compilation process. In a very cheap riscv machine assembling a 3.6Mb file takes 1.7 seconds, including the time for i/o from disk.